

Bioinformatics Algorithms

COS-BIOL-530/630

Lecture 10

Days & Times	Room	Meeting Dates
Tu 2:00PM - 3:50PM	Thomas Gosnell Hall (GOS)-2178	01/13/2025 - 04/28/2025
Th 2:00PM - 3:50PM	Thomas Gosnell Hall (GOS)-2178	01/13/2025 - 04/28/2025

Instructor:

Fernando Rodriguez

email: frvsbi@rit.edu

Office: Orange Hall 1311

Gene Prediction

- Lecture 10 -

Announcements

Lecture 10

Lab 10

- Activity 10
- Discussion 10

Quiz 9: Open Friday April 4th 5 pm (**next week!**)

- Lecture/Lab 10 (Gene Prediction)
- Lecture/Lab 11 (High Throughput Sequencing)

Gene Prediction

- Lecture 10 -

Topics:

- Gene recognition in sequences
- CpG islands
- Markov Model
- Hidden Markov Models

Gene Prediction

The problem:

We are given a sequence of **DNA**, and we wish to know which subsequence or concatenation of subsequences constitutes a **gene**.

>chr

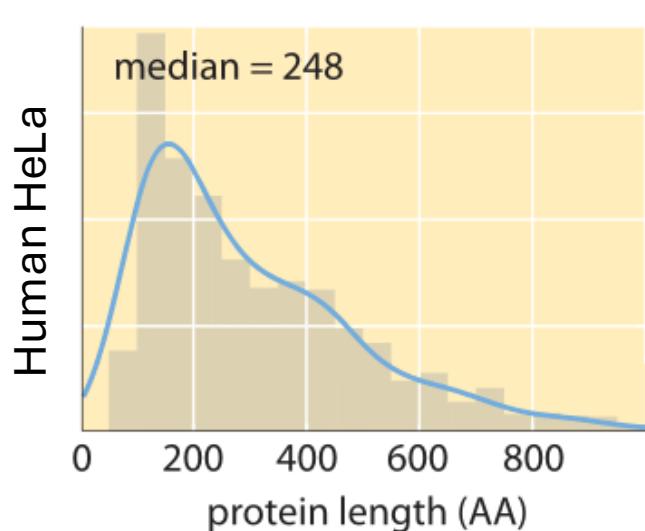
ATCTTTTCGGCTTTTTAGTATCCACAGAGGTTATCGACAACATTTCACATTACCAACCCCTGTGGA
CAAGGTTTTCAACAGGTTGCCGTTGTGGATAAGATTGTGACAACCATTGCAAGCTCTCGTTATT
TTGGTATTATTTGTGTTAACTCTGATTACTAATCCTACCTTCTTATCCACAAAGTGTGGA
TAAGTTGTGGATTGATTCACACAGCTGTGAGAAGGTTGTCCACAAGTTGTGAAATTGTGAAAAGC
TATTATCTACTATATTATGTTCAACATTAAATGTGACGAAATGGTAAGCGCCATTGCTTTTT
TTGTGTTCTATAACAGAGAAAGACGCCATTCTAAGAAAAGGAGGGACGTGCCGGAGATGGAAAATAT
ATTAGACCTGTGGAACCAAGCCCCTGCTCAAATCGAAAAAAAGTTGAGCAAACCGAGTTTGAGACTTGG
ATGAAGTCACCAAAGCCCCTGCAAGGCATACATTAACAATCACGGCTCCAATGAATTGCCA
GAGACTGGCTGGAGTCCAGATACTGCATCTGATTGAGATACTATATGAATTAAACGGGGAGAATT
GAGCATTAAGTTGTCATTCTCAAATCAAGATGTTGAGGACTTTATGCCGAAACCGCAAGTCAAAAAA
GCGGTCAAAGAACGATACATGATTTCTCAAATATGCTCAATCCAAATATACTTTGATACTTTG
TCATCGGATCTGAAACGATTGACATGCTGCTTCCCTCGCAGTAGCGGAAGCGCCCGCGAAAGCTTA
CAACCTTATTATCTATGGGGCGTCGGCTAGGGAAAACACACTTAATGCATGCGATGCCATTAT
GTAATAGATCATAATCCTCTGCAAAGTGGTTATCTGCTTCTGAGAAATTACAAACGAATTCATCA
ACTCTATCCGAGATAATAAAGCCGTGACTTCCGCAATCGCTATCGAAATGTTGATGTGCTTTGATAGA
TGATATTCAATTAGCGGGAAAGAACAAACCCAGGAAGAATTTCATACATTAAACACATTACAC
GAAGAAAGCAAACAAATCGTATTCAGTGACCGGCCAAAGGAAATTCCGACACTTGAAGACAGAT
TGCCTCACGTTTGAATGGGACTTATTACAGATATCACACCGCCTGATCTAGAACGAGAATTGCAAT
TTAAGAAAAAGGCCAAGCAGAGGGCCTCGATATTCCGACGAGGTTATGCTTACATCGCAATCAA
ATCGACAGCAATATCGGGAACTCGAAGGAGCATTAAATCAGAGTTGTCGCTTATTCACTTTAATTAA
AAGATATTAAATGCTGATCTGCCGCTGAGGCCTGAAAGATATTATCCTCTCAAACCGAAAGTCAT
TACGATAAAAGAAATTCAAGAGGGTAGTAGGCCAGCAATTAAATATTAAACTCGAGGATTCAAAGCAAA
AACGGACAAAGTCAGTAGCTTCCCGCTCAAATGCCATGTACTTATCAAGGAAATGACTGATTCT
CTCTCCTAAATCGGTGAAGAGTTGGAGGACGTGATCAGACCGTTATTGCGCATGAAAAAAT
TCAAAATGCTGGCAGATGATGAACAGCTTCAGCAGCATGTAAGAAATTAAAGAACAGCTAAATAG
CAGGACCGGGGATCAATCGGGAAAGTGTGAATAACTTCCGGAAGTCATACACAGTCTGCCACATGT
GATAGGCTGTGTTCTGCTTTTACAACCTATCCACAAATCCACAGGCCCTACTATTACTTACTA
TTTTTATAAATATATTAACATTATCGTTAGGAGGATAAAATGAAATTACGATTAAAAAG
ATCGTCTTGTGAAAGTGTCCAAGATGTATTAAAGCAGTTCATCCAGAACACGATTCCCATTCTGAC
TGGTATTAAATGTTGCATCAGATGATGGAGTATCCTTACAGGGAGTGAACGATATTCTATTGAA
TCCTCATTCAAAGAACAGAGGAGATAAGAAATCGTCACTATTGAAACAGGCCGGAGCATGTTTAC
AGGCTCGCTTTTAGTGAAGATTGTAAGGGGAAATTGCCGATGGCAACTGTAGAAATTGAAGTCCAAAATCA
GTATTGACGATTATCGTTCTGGTAAAGCTGAATTAAATCTAAACGGACTGGATGCTGATGAATATCG
CACTTGCCGAGATTGAAGAGCATATCGCATTGAGATCCAACTGATTGTTAAAAATCTAATCAGAC
AAACAGTATTGAGTGTCCACCTCAGAAACACGCCCTATCTGACAGGTGAAACTGGAAAGTGGAGCA
AAAGTGAATTATTATGCACTGCAACGGATAGCCACCGTCTGCATTAAGAAAGGCGAAACTGATATTCCA
GAAGACAGATCTTATAACGTCGTGATTCCGGGAAAAAGTTAACTGAACTCAGCAAGATTAGTGACA
ACCAAGGAACCTGTAGATATCGTCATCACAGAAACCCAAGTTCTGTTAAAGCAAAAACGTCTGTTCTT
CTCACGGCTTCTGGACGGGAATTATCCAGACACAAACCGCTGATTCCGCAAGACAGCAAAACAGAAATC

Algorithms and Models in Bioinformatics

Model of a gene

A gene is a sequence of nucleotides that encodes a protein sequence

- between 50 and 1000 residues in length
- A gene starts with Methionine
- And ends with a Stop Codon



Source: <https://book.bionumbers.org>

Gene finding algorithm

Search for and identify all sequences that:

- start with ATG
- end with either TAG/TAA/TGA
- between 150 and 3,000 nucleotides in length.

Define the DNA sequence as a string

Create an empty list to store identified gene sequences

Set `minimum_length = 150`
Set `maximum_length = 3000`

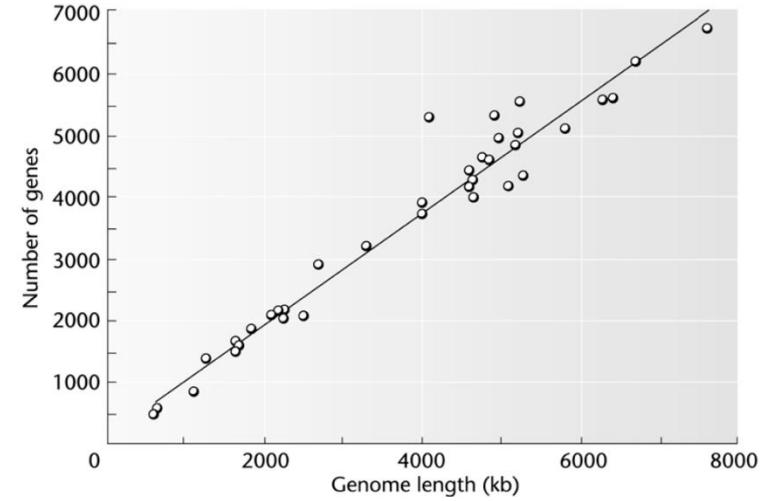
```
For each position i in the DNA sequence from 0 to (length of sequence - 3):
    If the substring from i to i+2 is "ATG": # Check for start codon
        For each position j from i+3 to (length of sequence - 2):
            If the substring from j to j+2 is "TAG" or "TAA" or "TGA": # Check for stop codon
                Calculate gene_length = j + 3 - i # Length of the gene
                If minimum_length <= gene_length <= maximum_length:
                    Extract the substring from i to j+2 as the gene
                    Add the gene sequence to the list
                Break the inner loop # Move to the next start codon
```

Output the list of identified gene sequences

Lecture
01!

Gene Prediction

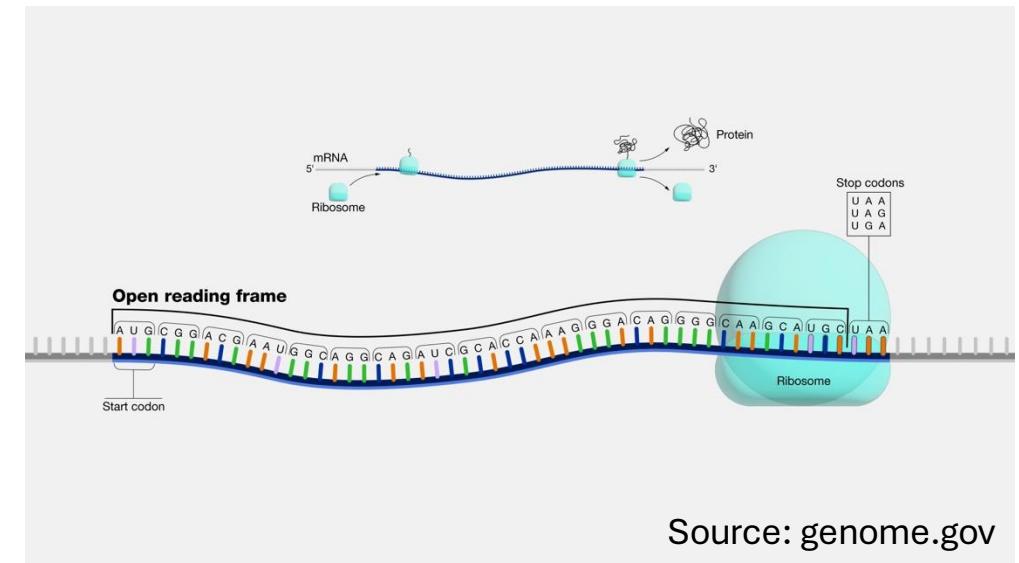
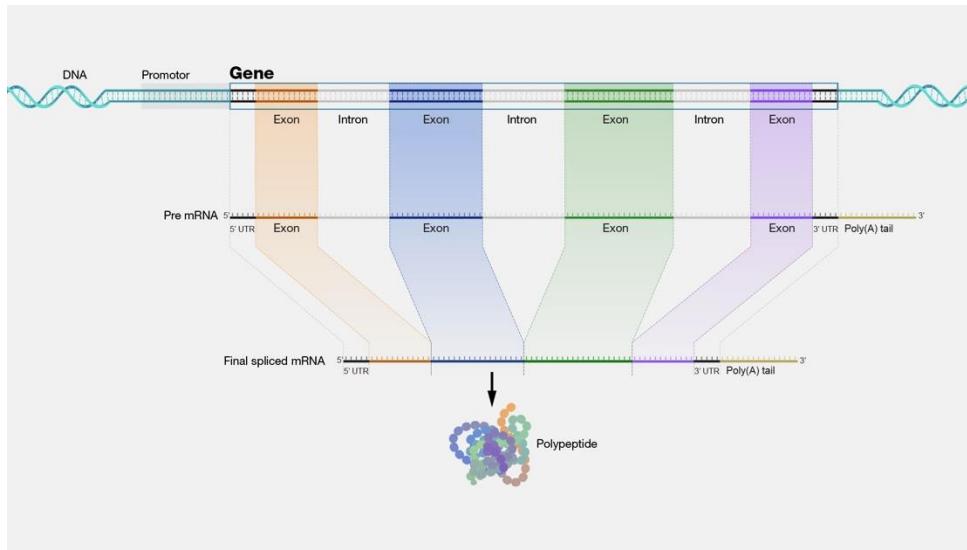
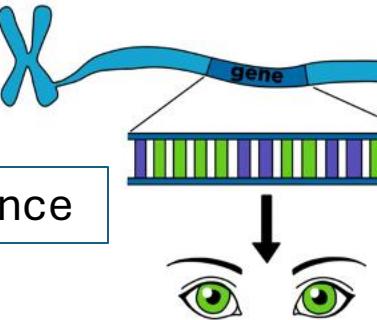
- There is a key difference between prokaryotes and eukaryotes.
 - > 85% of the prokaryotic genome is coding DNA →
 - The portion of eukaryotic genome coding:
 - Fungi: 70% in *Saccharomyces cerevisiae* (yeast)
 - Plants: 20% in *Arabidopsis*
 - Human: 3%
- We need to parse the genomic DNA to identify a complete gene with some knowledge (**a model**).



In prokaryotes most of the space is composed by genes.

Gene vs. Open Reading Frame

The gene is considered the basic unit of inheritance



Source: genome.gov

An **open reading frame (ORF)**, as related to genomics, **is a portion of a DNA sequence that does not include a stop codon** (which functions as a stop signal). A codon is a DNA or RNA sequence of three nucleotides (a trinucleotide) that forms a unit of genomic information encoding a particular amino acid or signaling the termination of protein synthesis (stop codon).

Prokaryotic genes

- A prokaryotic gene typically begins with a start codon (eg. ATG, GTG, TTG).
- Ends with one of the three stops codon (eg. TAG, TAA, or TGA).
- Most of the genes are organized in ***operons***: gene clusters of **more than one ORF** that are under the control of a shared set of regulatory sequences.
 - Promoters
 - Silencers
 - Terminators
 - Operators
- Regulatory sequences constitute 10–15% of the prokaryote genome.
- Promoters are located near the transcription start sites (TSSs) of genes, on the same strand and upstream of the gene or ORF.

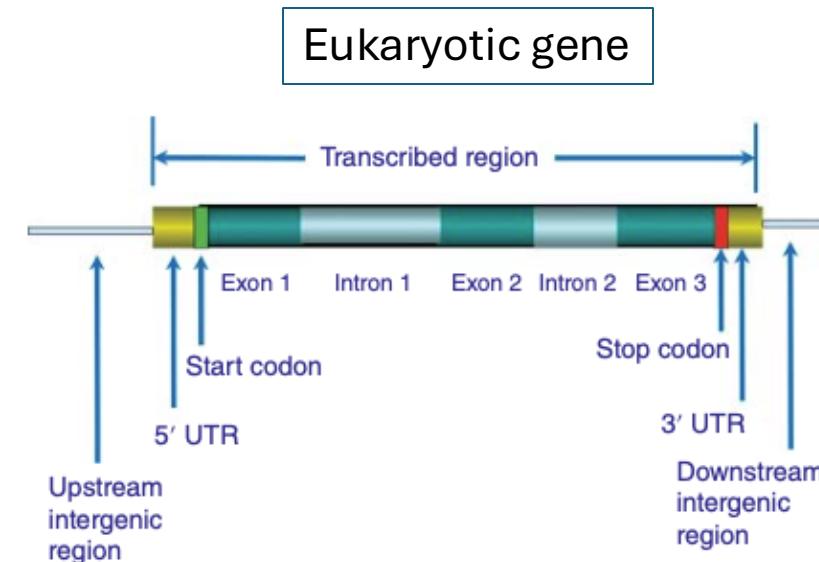


Source: Bioinformatics, Baxevanis

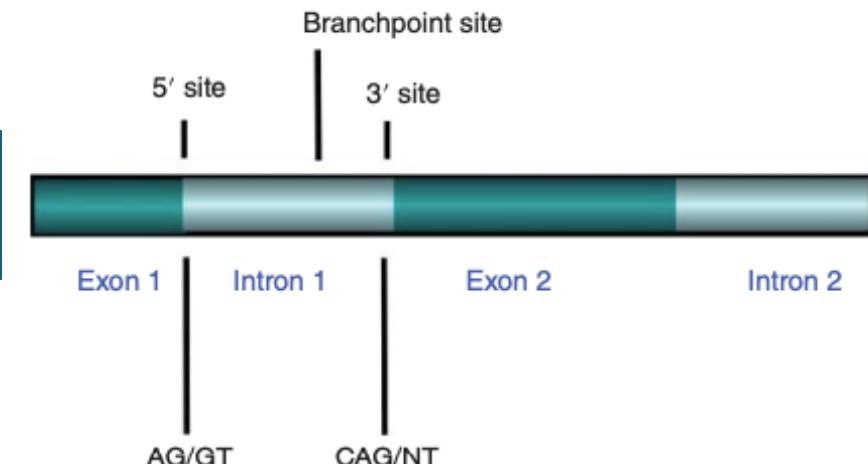
Eukaryotic genes

To find a eukaryotic gene, we must identify 4 signals:

- Start codon
- Stop codon
- Beginning of intron (donor site)
- End of intron (acceptor site)



Splice site region exon/intron

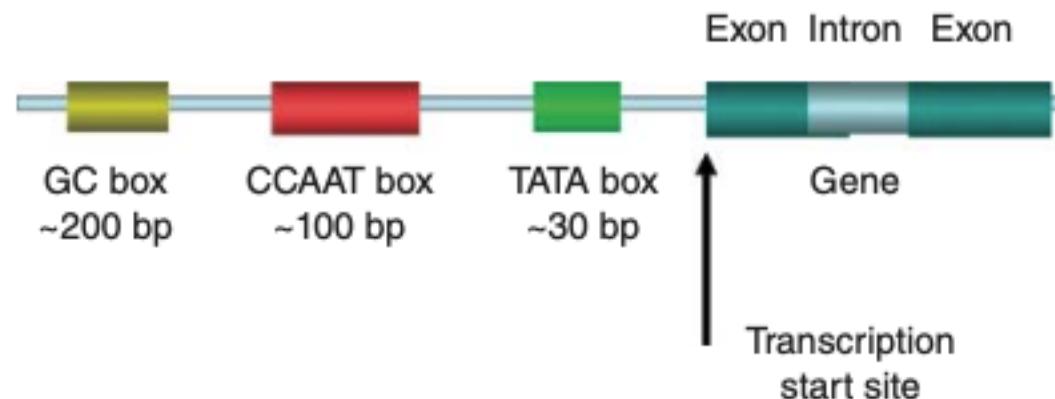


Eukaryotic genes

To find a eukaryotic gene, we must identify 4 signals:

- Start codon
- Stop codon
- Beginning of intron (donor site)
- End of intron (acceptor site)

* It helps to find other signals outside the gene, such as promoters and ribosomal binding sites



Start codon

What defines a start codon?

- Eukaryotes: ATG
- Prokaryotes: ATG, GTG, TTG
- In a random sequence, **the probability of an initiation codon is $(1/4)^3 = 1/64$ (for each).**
- There is usually a characteristic regulatory sequence upstream (promoter).
- However, regulatory element location and sequence are not consistent between species.

Stop codon

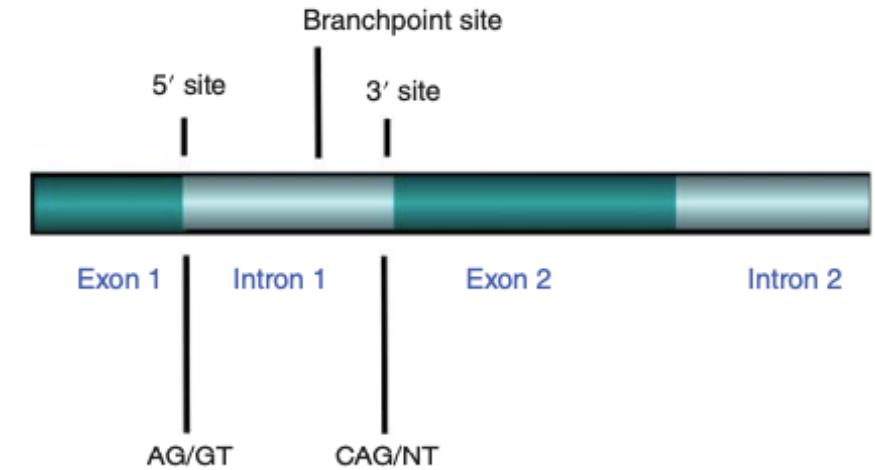
What defines a stop codon?

- TAA, TAG, TGA
- No sequence or regulatory element upstream.
- In a random sequence, **the probability of a stop codon is $(1/4)^3 = 1/64$ (for each).**

Splice sites

What defines a splice site (eukaryotes)?

- The donor is **almost** always GT.
- The acceptor is **almost** always AG.
- There are **certain consistencies** around the splice sites.
- However, the (average) size of introns varies between species.
- Exon base count is not always multiple of 3: introns can split codons!
- The split is not necessary in the same frameshift.



ORFs

What defines an open reading frame?

- Begins with **START codon**, ends with **STOP codon**.
- If we have a DNA sequence (not genome), an ORF must be the longest sequence without a stop codon.
- This is easy! In theory, an ORF is a gene. Indeed, in bacteria, it is a gene.
- The problem:
 - An ORF is ended by a stop codon. How do you know that the stop codon is the **real stop**?
 - Maybe the stop codon is just a random sequence embedded (randomly) in a non-coding area.
 - And the stop codon could come from any of the SIX reading frames (sense and anti-sense, 3 frames each).



```
>chr
ATCTTTTCGGCTTTTAGTATCCACAGAGGTTATCGACAACATTTCACATTACCAACCCGTGGA
CAAGGTTTTCAACAGGTTGCCGTTGGATAAGATTGTGACAACCATTGCAAGCTCGTTATT
TTGGTATTATATTGTGTTTACTCTGATTACTAATCCTACCTTCCTCTTATCCACAAAGTGTGGA
TAAGTTCTGCAATTGAGAGACCTTGCTGAGAACCTTGCGAGAACCTTGCAAAATTGCTGAAAGC
```

Algorithm *Find_Stop_Codons*

Input: DNA_sequence (a string of nucleotides)

Output: A list of positions of stop codons in each reading frame

Define stop_codons as ["TAA", "TAG", "TGA"]

Initialize `stop_positions` as an empty dictionary with keys "Frame 1", "Frame 2", and "Frame 3".

For each frame in {0, 1, 2}: // Three reading frames

Initialize `stop_positions["Frame " + (frame + 1)]` as an empty list

For i from frame to length(DNA_sequence) - 2 step 3:

codon \leftarrow substring of DNA_sequence from i to i+2 (inclusive)

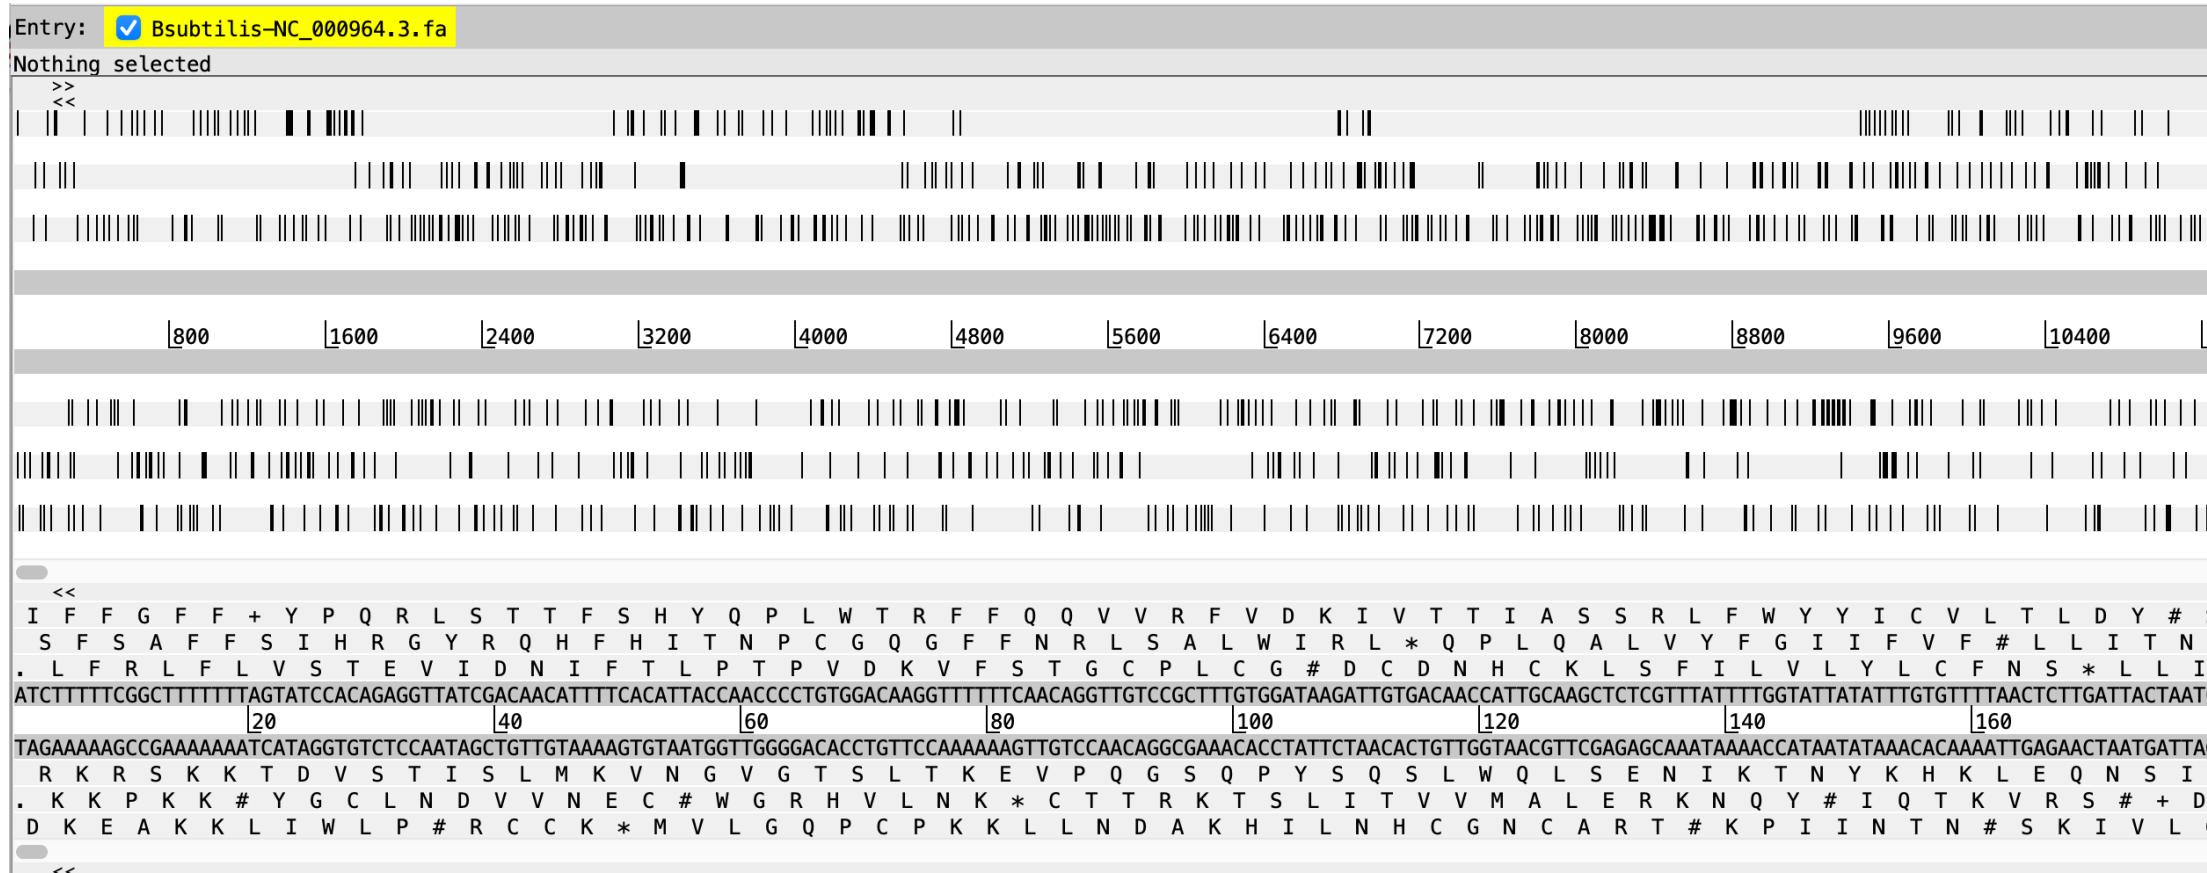
If codon is in stop_codons:

Append i (position of stop codon) to stop_positions["Frame " + (frame + 1)]

Return stop_positions

Algorithm *Find_Stop_Codons*

For each frame



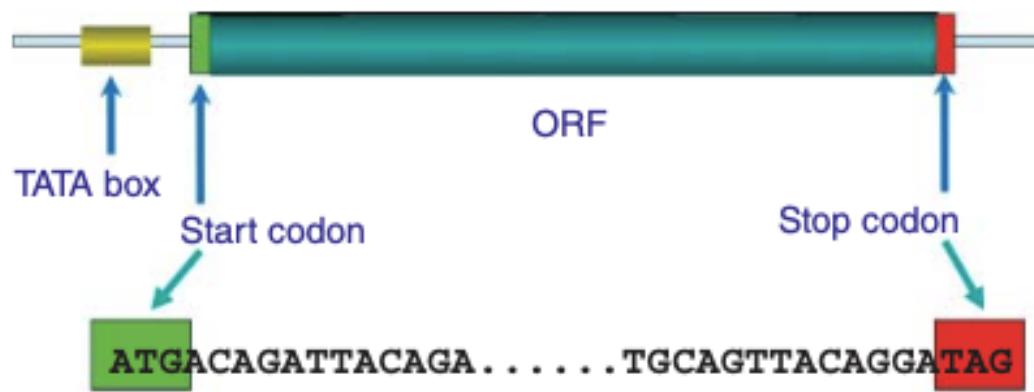
Gene Prediction

The solution:

- In prokaryotes:
 - Find the ORF
- In Eukaryotes:
 - Need to identify more structures (splice sites)

Gene Finding Strategy

in prokaryotes



Gene Finding Strategy

in prokaryotes



- Brute force:
 - Find ATG start codon
 - Longest ORF (>150 bases)
 - Move to the next ATG downstream
 - Repeat the process in the opposite strand
- Find motif signal (matrix profile)
 - TATA - Pribnow box

Gene Finding Models

- Markov Models (MMs)
- Hidden Markov Models (HMMs)

Models and Algorithms used in Computational Biology

A **model** is a parametric explanation of the observations of interest.

Probabilistic models/methods:

- Maximum Likelihood
- Bayesian
- Machine Learning
- Markov Chain Models
- Hidden Markov Models

An **algorithm** is a set of instructions for solving a problem, e.g. , inferring the optimal value of a model's parameter.

Algorithms/Methods:

- Sequence (string): sort/search algorithms
- Optimization algorithms:
 - Linear programming
 - Dynamic programming
 - Greedy algorithms
 - Heuristic methods

“All models are wrong, but some models are useful”

George Box

**Correct versus incorrect
Algorithms**

Lecture
01!

Markov Models (MMs)

A Markov chain, model, or process refers to a series of observations in which the **probability of an observation depends on a number of previous observations**. The number of observations defines the “order” of the chain.

For example, in a first-order Markov model, the probability of an observation depends only on the previous observation. In a Markov chain of order 5, the probability of an observation depends on the five preceding observations.

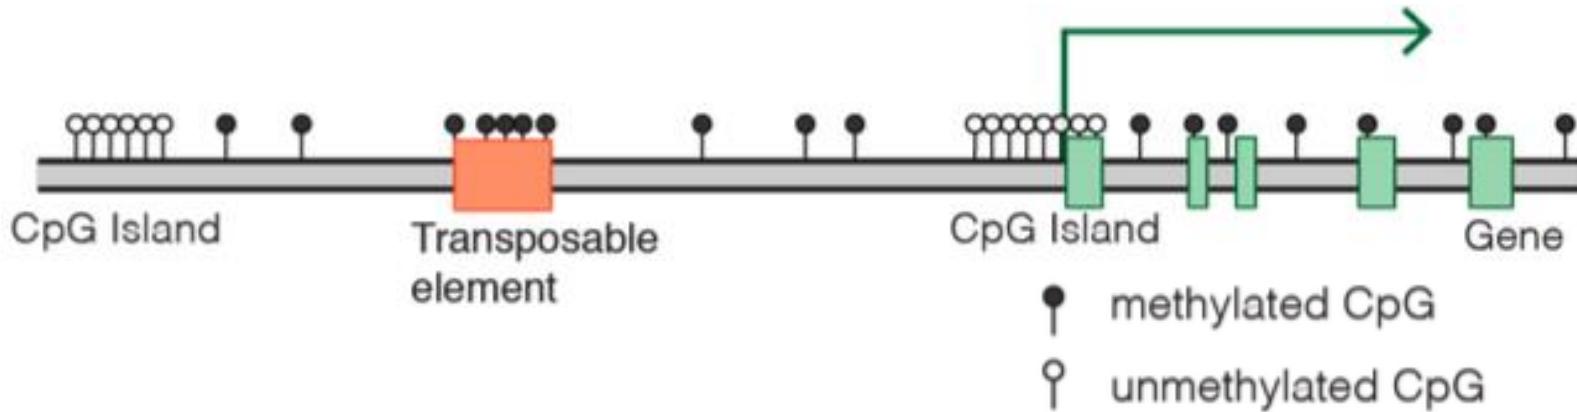
A DNA sequence can be considered to be an example of a Markov model because the likelihood of **observing a particular base at a given position may depend on the bases preceding it**. In particular, in coding regions, it is well known that **the probability of a given base depends on the five preceding bases**, reflecting observed **codon biases** and dependencies between adjacent codons.

Such dependence is **not observed in non-coding regions**. When scanning an anonymous genomic region, one can compute how well the local nucleotide sequence conforms to the fifth-order dependencies observed in coding regions and assign appropriate coding likelihood scores.

CpG islands

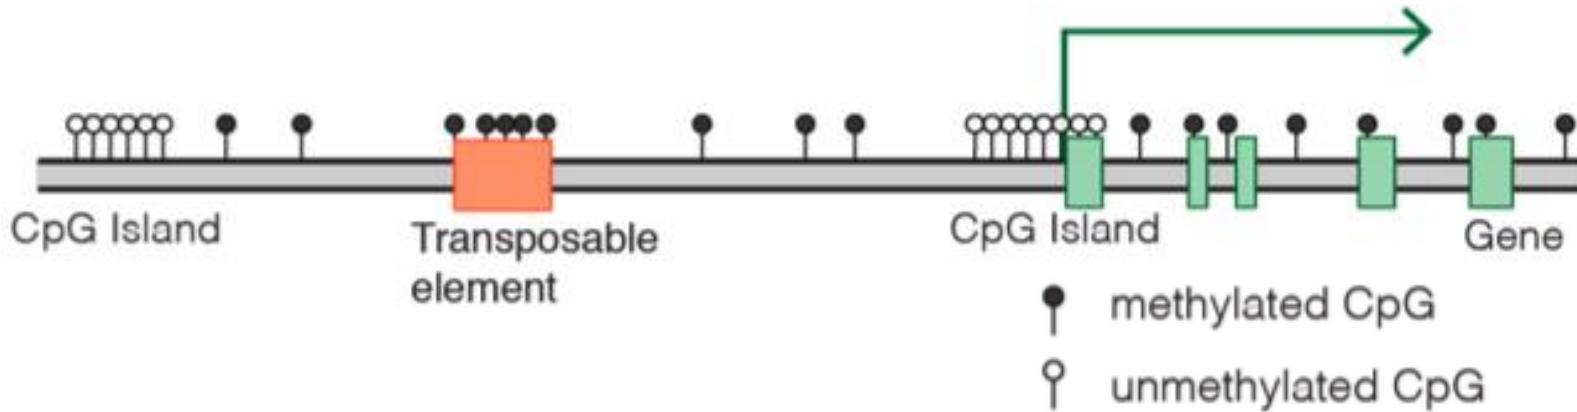
- In the (human) genome, CpG dinucleotides should be rare (probabilities of C and G). **Why?**
- They are rich regions in CG dinucleotides.
- From 100-1000 bases long.
- The cytosine is usually modified by methylation (5mC).
- CpG regions in the genome play an essential role in regulation (suppressing nearby promoters/genes).
- CpG islands have other bases (A and T); they are just rich in CG dinucleotides.
- Question: given a stretch of a genomic region, **how can we say if it comes from a CpG island or not?**

CpG islands



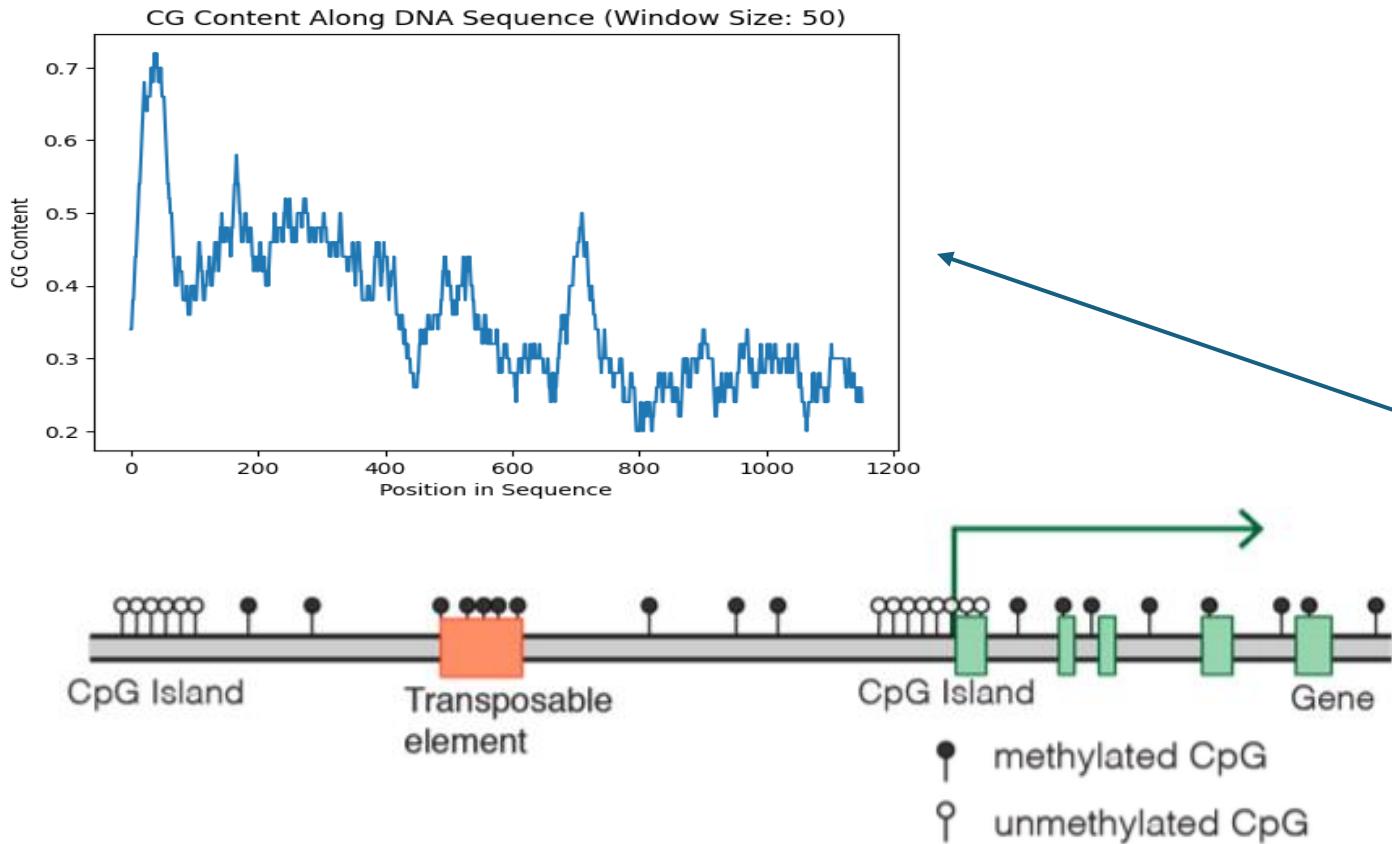
Check
Seq_CG-content
Notebook

CpG islands



Is the sequence
ATCG more likely
to be from a CpG
island?

Finding CpG islands



Problems with sliding window heuristics:

- What is the right window size?
 - Too small: break up real islands
 - Too large: we miss islands
- What cut-off (threshold) should we use?

*Check Seq_CG-content [Notebook](#)

Markov property example

Sequence: ATG ATG

Table of transition probabilities

	A	G	C	T
A				
G				
C				
T				

- If you know you are looking at a sequence of $(ATG)_n$...
- The probability of the next character being a G?
Depending on what character we are looking at:
 - If you are looking at a T: the odds are good that the next is a G
 - If you are looking at a A: the odds are weak that the next is a G
 - If you are looking at a G: the odds are weak that the next is a G

↓ : Low probability

↑ : High probability

Markov property example

Table of transition probabilities

	A	G	C	T
A				
G				
C				
T				

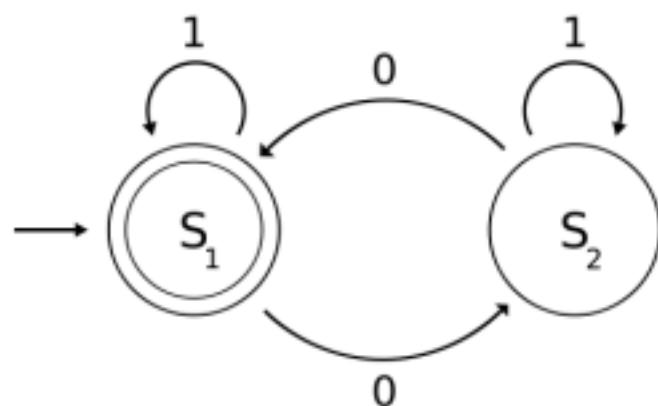
- As a result, we can model such a system with a series of transition probabilities.
- A table representing
 - If we are here...
 - What is the probability of getting there.
 - For all possible scenarios

↓ : Low probability

↑ : High probability

Markov chains

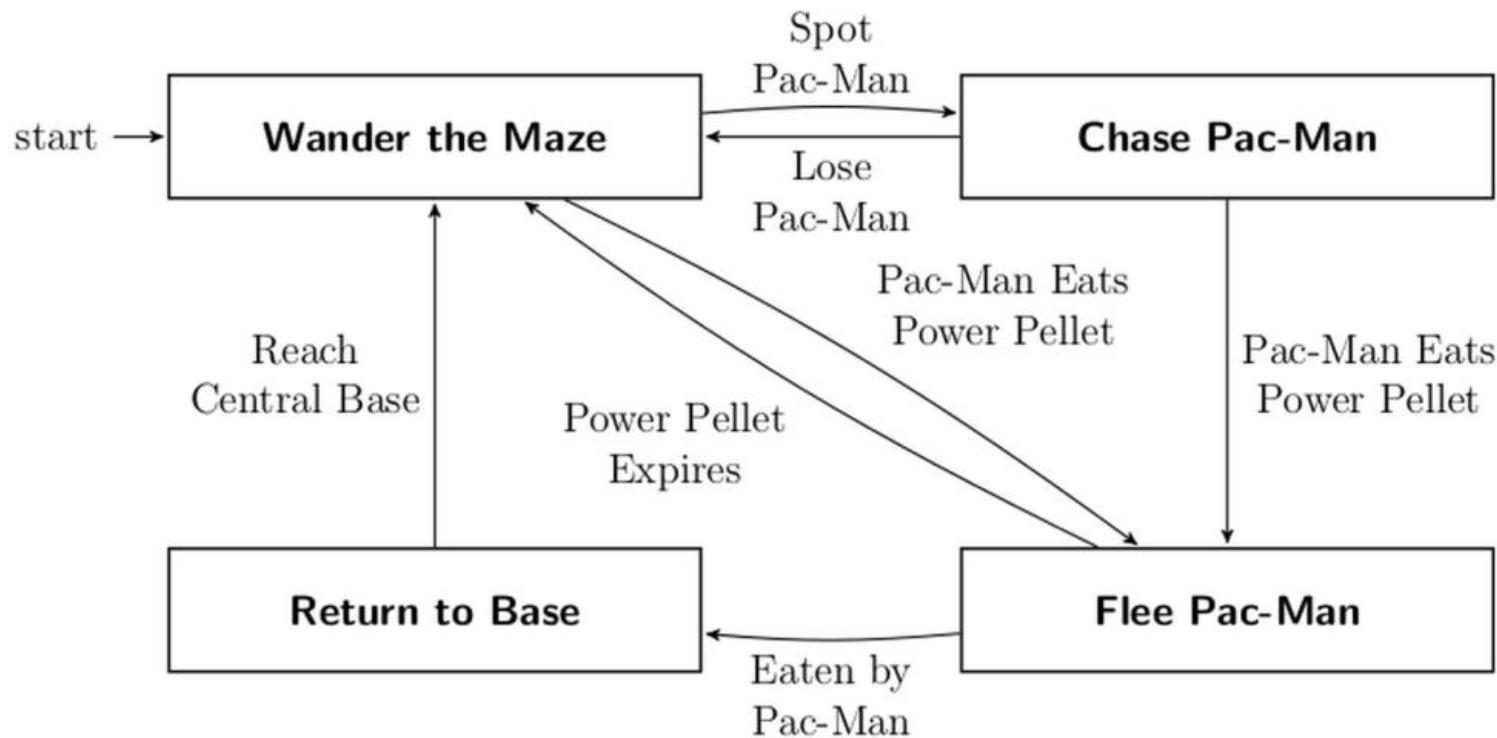
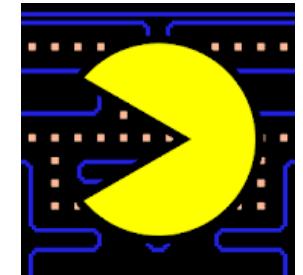
- Markov Chains are implementations of models for systems with the Markov property.
- They are represented as **Deterministic Finite Automaton** (DFA).
- The **edges** (arrows) are represented as probabilities of a transition.
- The **vertices** (nodes) represent states.
- Often the DFA has “begin” and “end” states.



*A DFA requires **O(1)** memory (constant), regardless of the length of the input.*

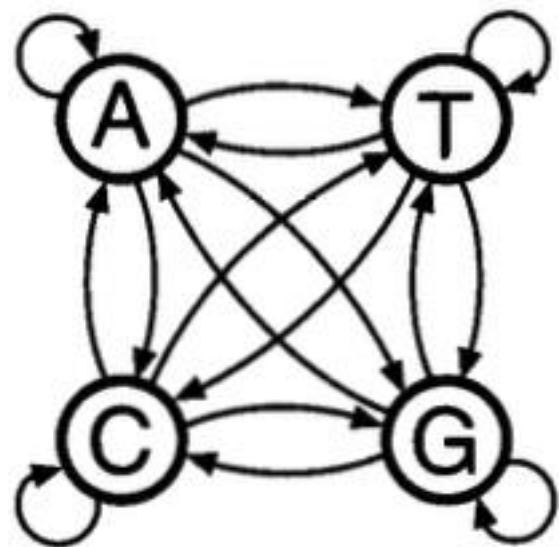


Markov chains



Pac-Man uses a four-state automaton

DNA Markov chains



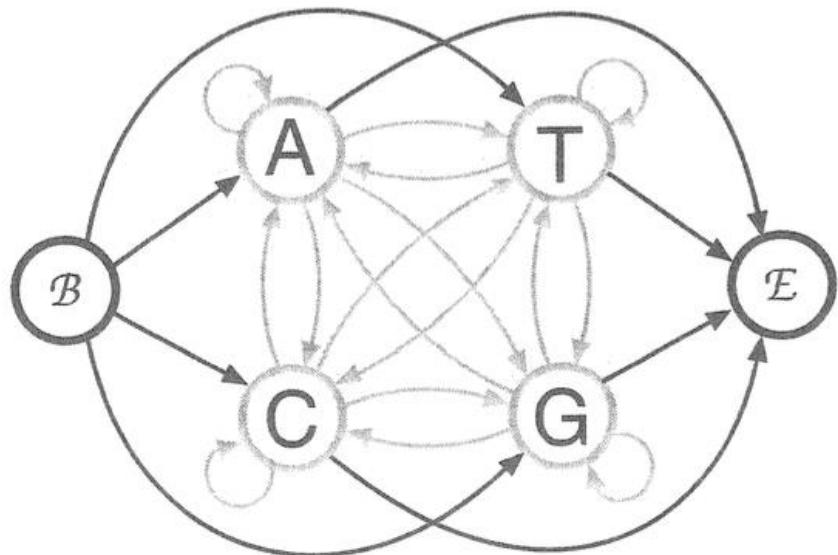
A Markov chain or DNA can be drawn with a state for each of the four letters A, C, G and T.

A probability parameter is associated with each arrow (edges): the probability of a certain residue following another residue.

The probability parameters are called **transition probabilities**.

Source: *Biological sequence analysis*, Durbin

DNA Markov chains



Source: *Biological sequence analysis*, Durbin

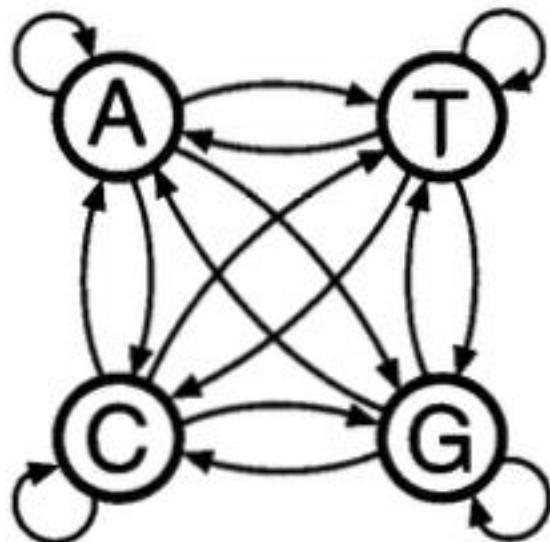
A Markov chain or DNA can be drawn with a state for each of the four letters A, C, G and T.

A probability parameter is associated with each arrow (edges): the probability of a certain residue following another residue.

The probability parameters are called **transition probabilities**.

Begin and end states can be added to the Markov Chain.

DNA Markov chains



L = 4
P (ACGT)
P (AGTC)
...
P (TGCA)



When moving between states, it accumulates the product of probabilities.

$$\sum_x P(x)$$

The **probability for a given model** approaches to zero.

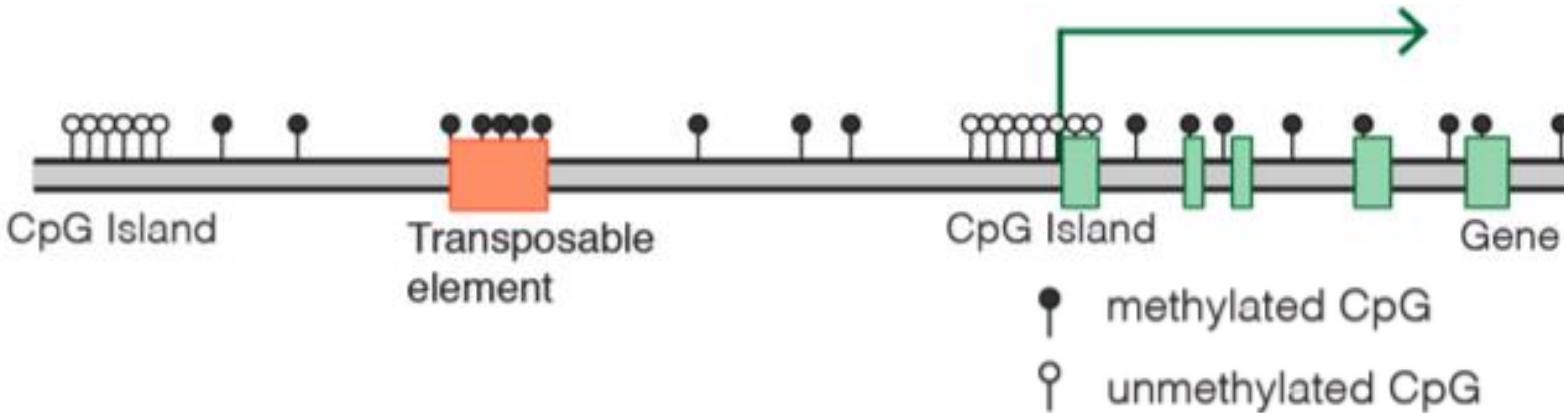
$$P = 0.5 * 0.2 * 0.1 * 0.3 \dots$$

The longer the model runs, the smaller the result.

Each chain represents the probability of following the exact chain/string (eg. sequence nucleotide).

What is the sum of the probabilities of all possible sequences of length L?

CpG islands



We have to train our model for:

- CpG island
- Non-CpG island (or CpG oceans)

Compute the probability of each possible transition.

$$\frac{\sum_{\forall i, \forall j} transition_{island} i \rightarrow j}{number\ of\ all\ possible\ transitions_{island}}$$

CpG islands

Table of transition probabilities

- Rows = from
- Columns = to
- Red = high probability
- Orange = medium
- Yellow = low

Non-CpG island

	A	G	C	T
A	orange	orange	orange	orange
G	orange	orange	orange	orange
C	orange	yellow	orange	orange
T	orange	orange	orange	orange

CpG island

	A	G	C	T
A	orange	red	red	orange
G	orange	red	red	orange
C	orange	red	high	orange
T	orange	red	red	orange

CpG islands

Probabilities

	A	G	C	T
A	0.19	0.27	0.40	0.14
G	0.17	0.33	0.36	0.14
C	0.19	0.36	0.25	0.20
T	0.10	0.34	0.38	0.19

What's the probability of the sequence ATCG in a CpG island?

$x = \text{ATCG}$

$$P(x) = P(x_4|x_3)P(x_3|x_2)P(x_2|x_1)P(x_1)$$

$$P(x) = P(G|C)P(C|T)P(T|A)P(A)$$

$$P(x) = 0.36 * 0.38 * 0.14 * 0.16$$

* $P(A)$ approx. = mean $P(A|X) = 0.16$

Simulating

$$P(C|A) = 0.40$$

Building

$$P(C|A) = \# \text{ times AC occurs} / \# \text{ times AX occurs}$$

CpG islands

Probabilities

Non-CpG island

	A	G	C	T
A	0.34	0.23	0.18	0.25
G	0.30	0.25	0.20	0.25
C	0.38	0.04	0.26	0.33
T	0.22	0.26	0.21	0.31

$x = \text{ATCG}$

$$P(x) = 0.04 * 0.21 * 0.25 * 0.31$$

$$P(x) = 0.000651$$

Which one is more likely?

Which model/scenario
(CpG island or non-CpG
island) is more likely for
the sequence ATCG to
be?

4.7 time more likely

CpG island

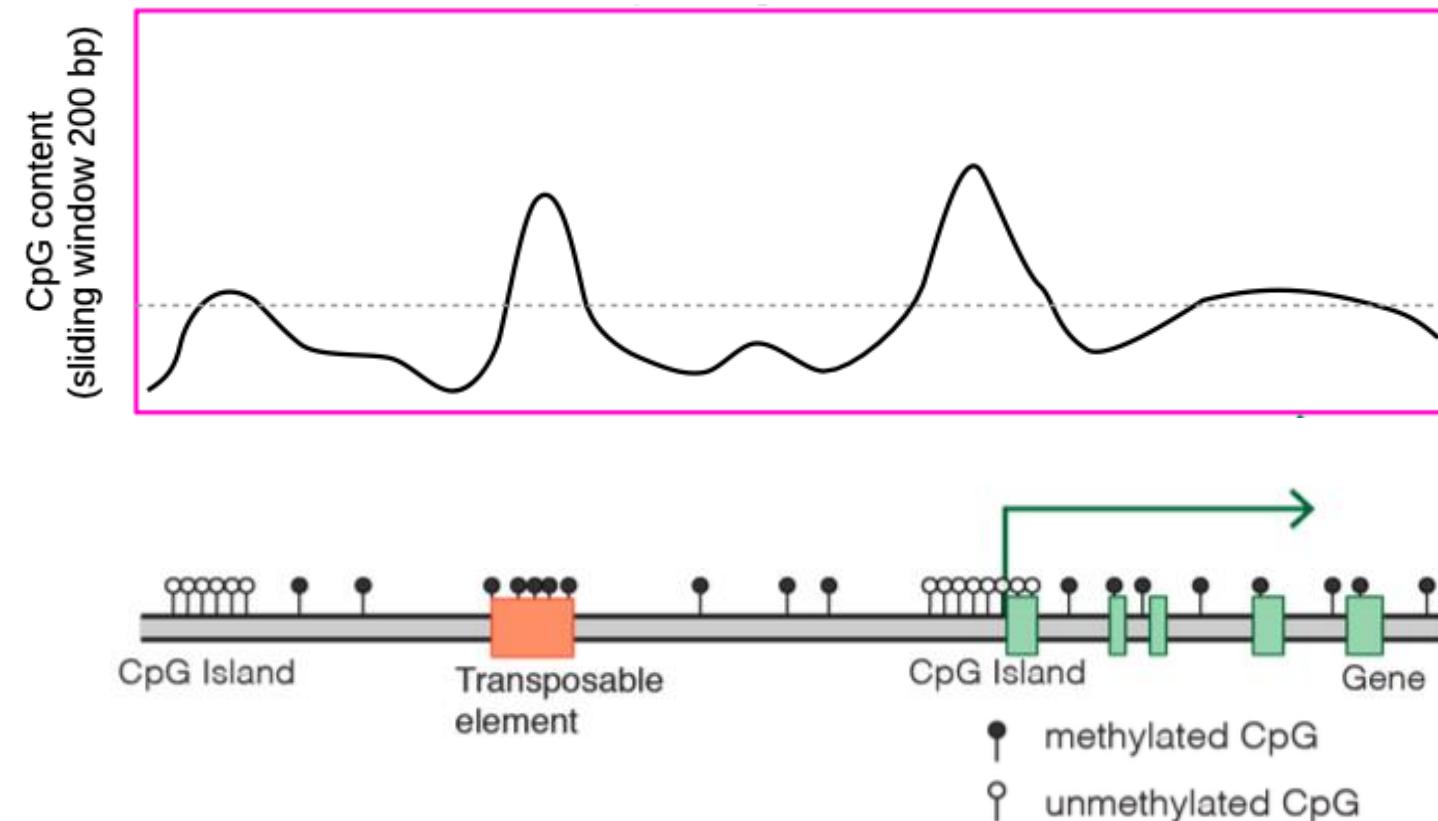
	A	G	C	T
A	0.19	0.27	0.40	0.14
G	0.17	0.33	0.36	0.14
C	0.19	0.36	0.25	0.20
T	0.10	0.34	0.38	0.19

$x = \text{ATCG}$

$$P(x) = 0.36 * 0.38 * 0.14 * 0.16$$

$$P(x) = 0.00306$$

Finding CpG islands



Problems with sliding window heuristics:

- What is the right window size?
 - Too small: break up real islands
 - Too large: we miss islands
- What cut-off (threshold) should we use?

*Check Seq_CG-content [Notebook](#)

Finding CpG islands

Markov Models

Non-CpG island

	A	G	C	T
A	0.34	0.23	0.18	0.25
G	0.30	0.25	0.20	0.25
C	0.38	0.04	0.26	0.33
T	0.22	0.26	0.21	0.31

- In order to train the model, you must know up front whether the training data came from an **island** or an **ocean**.

- The **cut-off problem is solved** with MMs: calculate which model is more likely.

CpG island

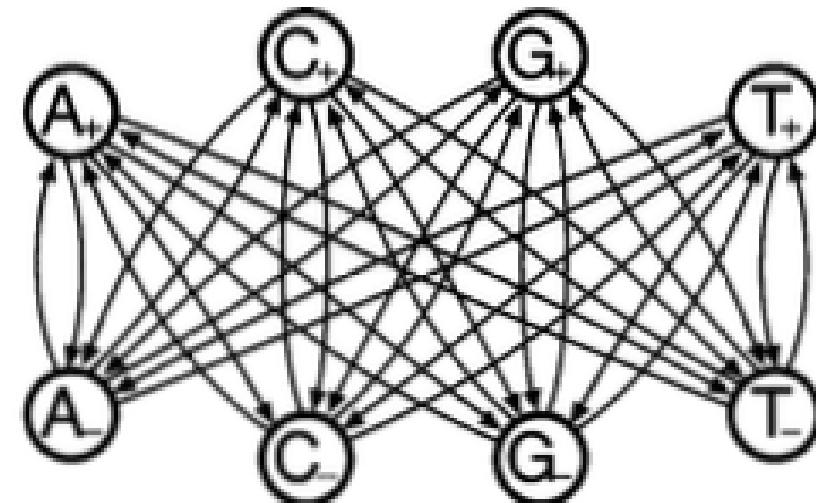
	A	G	C	T
A	0.19	0.27	0.40	0.14
G	0.17	0.33	0.36	0.14
C	0.19	0.36	0.25	0.20
T	0.10	0.34	0.38	0.19

- We can use Markov Chains for discrimination.

- But we still have to deal with size and boundary problems:
*“Distinguishing the **shorelines**. ”*

Hidden Markov Models

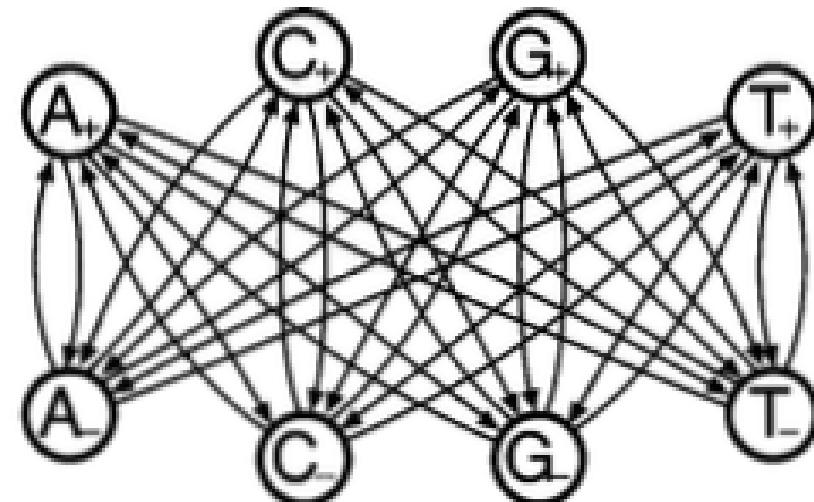
- We can use Hidden Markov Models (HMMs) to identify the boundaries of the CpG island.
- CpG island boundaries are "sharp" but with variable length.
- We can implement a model with an additional set of transition probabilities and changes of states.
- To simulate in one model the “islands” in an “ocean” of non-island genomic sequence, we want to have both Markov chains present in the same model.
- We relabel the states with “+” and “-” symbols.
- Look for the most probable state path.



Source: *Biological sequence analysis*, Durbin

Hidden Markov Models

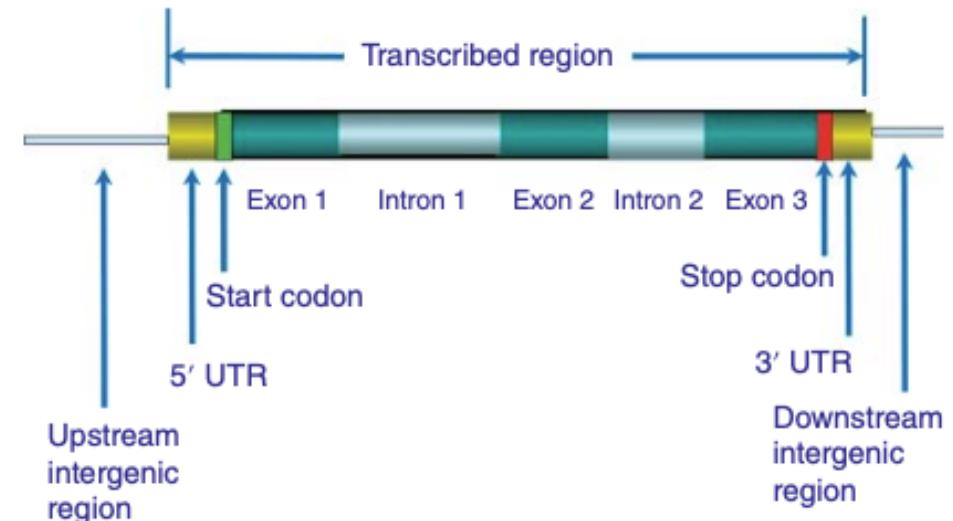
- A set of states (eg. CpG island, CpG ocean).
- A set of symbols to be determine (A, C, G, T)
- A set of **emission probabilities** for each symbol from each stated
- An index (eg. next nucleotide)
- A **transition probability** between successive states.
- **What other states can we look for in a genome?**



Source: *Biological sequence analysis*, Durbin

Hidden Markov Models

- A set of states (eg. CpG island, CpG ocean).
- A set of symbols to be determine (A, C, G, T)
- A set of emission probabilities for each symbol from each stated
- An index (eg. next nucleotide)
- A transition probability between successive states.
- **What other states can we look for in a genome?**
- Hidden Markov Models (HMM) forms the core component of most gene predictors.



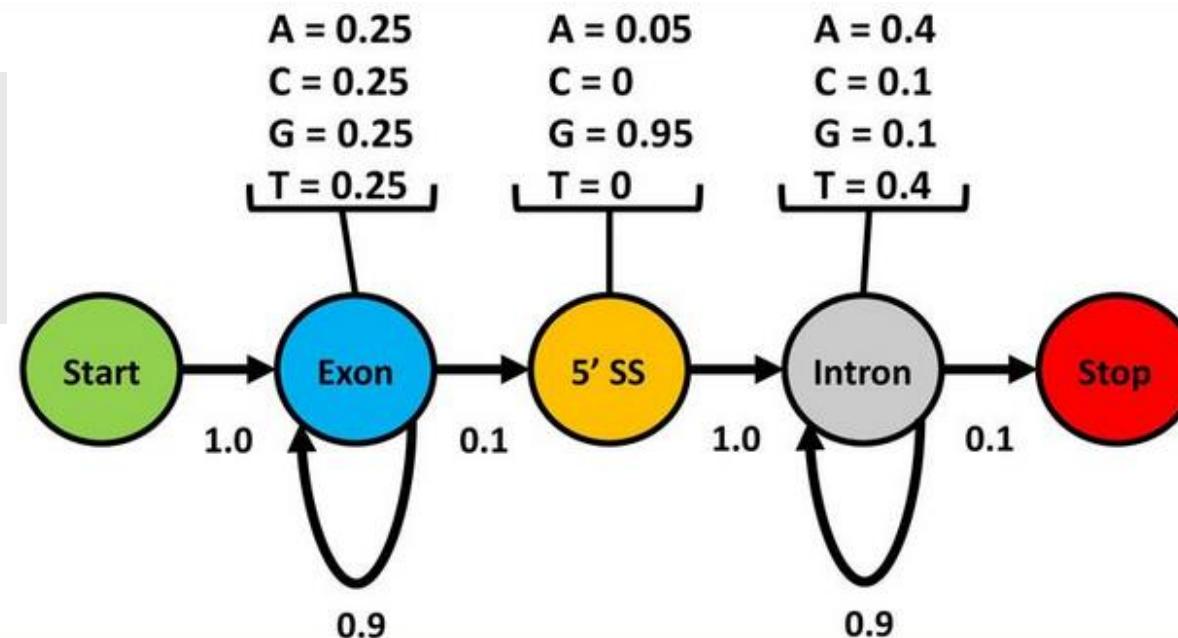
Hidden Markov Models in gene prediction

- Hidden Markov models (HMMs) are used to provide a **statistical representation** of real biological processes.
- HMMs are used in speech recognition, facial recognition, and other applications.
- They have found widespread use in many areas of bioinformatics, including multiple sequence alignment, the characterization and classification of protein families, the comparison of protein structures, and the prediction of gene structure.
- In general, gene-finding methods use a raw nucleotide sequence as their input and, for each position in the sequence, they attempt to predict whether **a given base is most likely found in an intron, an exon, or within an intergenic region**.
- In making these predictions, the algorithm applied (variation of HMMs) must consider what is known about the structure of a gene for that specific genome or taxa.
- Each of the elements – exons, introns, and so forth – are referred to as **states**.

HMM Probabilities

- The probability of switching from one state to another (eg. exon -> intron) is called **transition probability**.
- The probability of observing a nucleotide (A, T, C, G) that is of a certain state (exon, intron, splice junction) is called an **emission probability**. (eg. the probability of observing an adenine in an exon).

The probability of switching from one state type to another (ex. Exon - Intron).



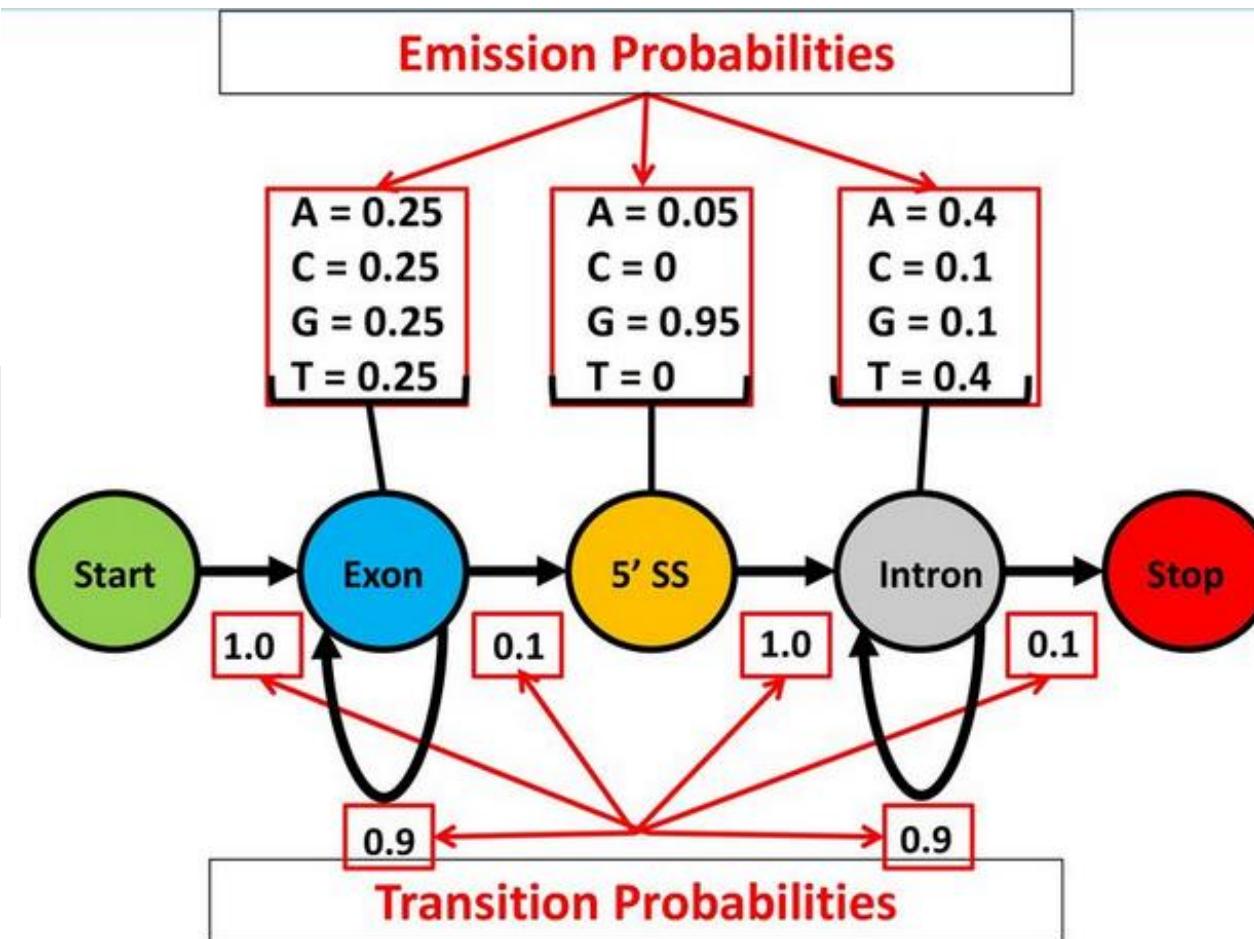
Source: *Introduction to HMMs*, Weisstein

ss

Splice Site

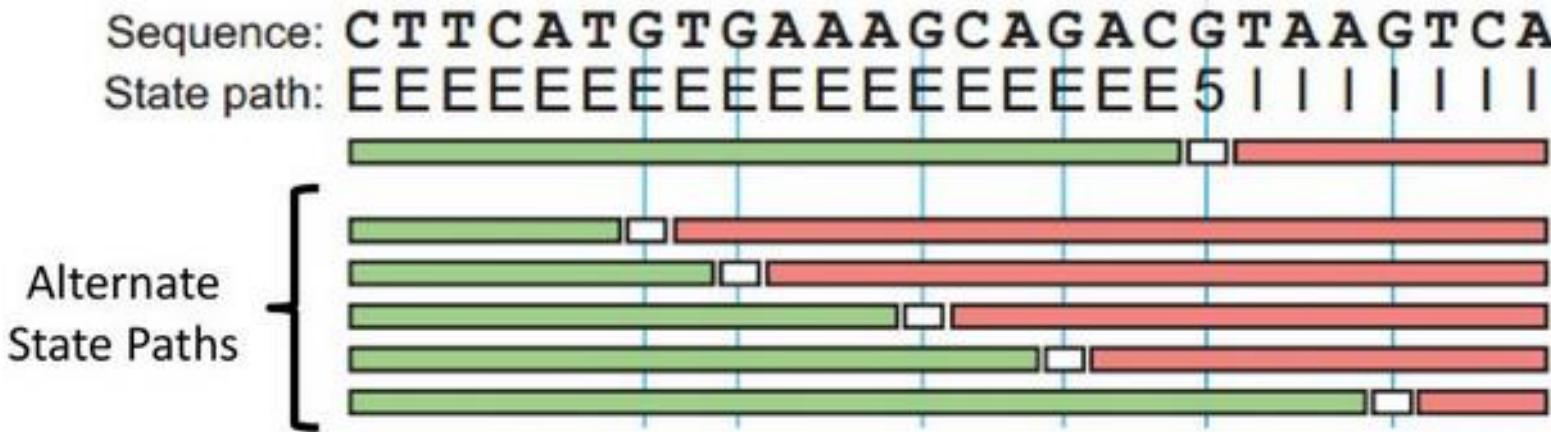
HMM Probabilities

The probability of switching from one state type to another (ex. Exon - Intron).



Source: *Introduction to HMMs*, Weisstein

HMM Probabilities



- A **state path** is the list of states (E: exon; I: intron; 5: 5' splice junction).
- An HMM can produce many state paths for a single sequence.
- We use an algorithm to determine what is the most likely path given the emissions:
 - **Viterbi algorithm:**
 - Calculates a transition matrix.
 - Use dynamic programming to find the most probable path

Eukaryotic gene prediction

- We must train or know intron-exon and exon-intron junctions.
- The gene finder must find promoter motifs.
- Must have ORF awareness.
- HMMs are the core of several gene prediction algorithms:
 - GenScan
 - Augustus
 - GeneMark
 - GRAIL
 - Twinscan
- Gene prediction accuracy depends partly on transition probabilities calculated based on the training data.

UCSC Genome Browser on Human Dec. 2013 (GRCh38/hg38) Assembly

move <<< << < > >> zoom in 1.5x 3x 10x base zoom out 1.5x 3x 10x 100x

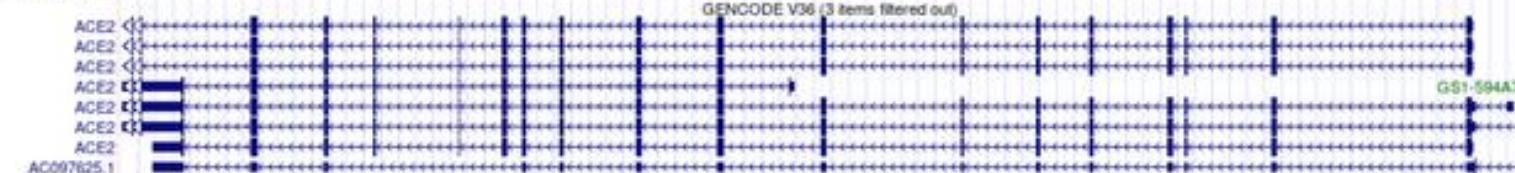
multi-region chrX:15,560,138-15,602,945 42,808 bp. enter position, gene symbol, HGVS or search terms

go

chrX (p22.2) Xp22.2 21.3 p21.1 11.4 p11.3 12.1 13.1 Xq21.1 21.31 22.1 22.3 Xq23 q24 Xq25 26.3 q27.3 Xq28

Scale
chrX: 15,565,000 | 15,570,000 | 15,575,000 | 10 kb | 15,580,000 | 15,585,000 | 15,590,000 | hg38 | 15,595,000 | 15,600,000 |Reference Assembly Fix Patch Sequence Alignments
Reference Assembly Alternate Haplotype Sequence Alignments

All Haplotypes



Gene Expression in 54 tissues from GTEx RNA-seq of 17382 samples, 948 donors (V8, Aug 2019)

ENCODE cCREs
Layered H3K27Ac
H3K27Ac Mark (Often Found Near Regulatory Elements) on 7 cell lines from ENCODE

100 vertebrates Basewise Conservation by PhyloP

Cons 100 Verts

Multiple Alignments of 100 Vertebrates

Rhesus
Mouse
Dog
Elephant
Chicken
X_tropicalis
Zebrafish

Short Genetic Variants from dbSNP release 153

Repeating Elements by RepeatMasker

Common dbSNP(153)

RepeatMasker

move start < 2.0 > Click on a feature for details. Click+shift+drag to zoom in. Click side bars for track options. Drag side bars or labels up or down to reorder tracks. Drag tracks left or right to new position. Press "?" for keyboard shortcuts. move end < 2.0 >

track search default tracks default order hide all add custom tracks track hubs configure reverse resize refresh

collapse all

expand all

Use drop-down controls below and press refresh to alter tracks displayed.
Tracks with lots of items will automatically be displayed in more compact modes.

refresh

Thursday Lab10

Gene prediction in prokaryotes

GeneMarkS-2 (<https://genemark.bme.gatech.edu/>)