

# Bioinformatics Algorithms

## COS-BIOL-530/630

### Lecture 12

| <b>Days &amp; Times</b> | <b>Room</b>                    | <b>Meeting Dates</b>    |
|-------------------------|--------------------------------|-------------------------|
| Tu 2:00PM - 3:50PM      | Thomas Gosnell Hall (GOS)-2178 | 01/13/2025 - 04/28/2025 |
| Th 2:00PM - 3:50PM      | Thomas Gosnell Hall (GOS)-2178 | 01/13/2025 - 04/28/2025 |

Instructor:

Fernando Rodriguez

email: [frvsbi@rit.edu](mailto:frvsbi@rit.edu)

Office: Orange Hall 1311

# De Bruijn Graphs

## - Lecture12-

### Announcements

#### Lecture12

**Lab 12:** Genome Assembly in Oedipus

- Discussion12
- Activity12 (\*Edited due date assignments)

**Quiz10:** Lecture/Lab12 + Lecture13: Opens Friday April 23rd.

**Exam 2:** Thursday , April 17<sup>th</sup> 2pm (GOS)-2178

Lecture/Lab 07 to Lecture/Lab 11

- Gene ontologies
- Algorithms and pattern matching
- RNA secondary structure
- Gene prediction
- Sequencing technologies (HTS Intro)

# De Bruijn Graphs

## - Lecture 12 -

Topics:

- Sanger vs. High Throughput Sequencing Assembly
- Graph Theory / Eulerian cycles
- de Bruijn graphs

# The Human Genome Project Battle

## Public consortium:

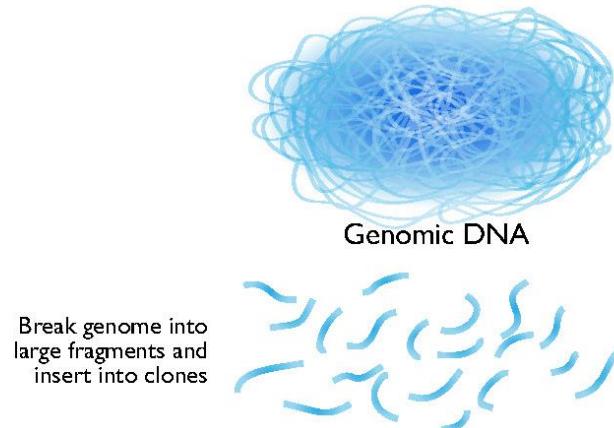
- \$3 billion
- Hierarchical approach (Lander et al. 2001)
- Initial step of clone-based physical mapping to generate a list of overlapping BAC clones to be sequenced (\$\$\$\$)
- Followed by shotgun sequencing of the individually mapped BAC clones (100 kb – 200 kb)

## Celera:

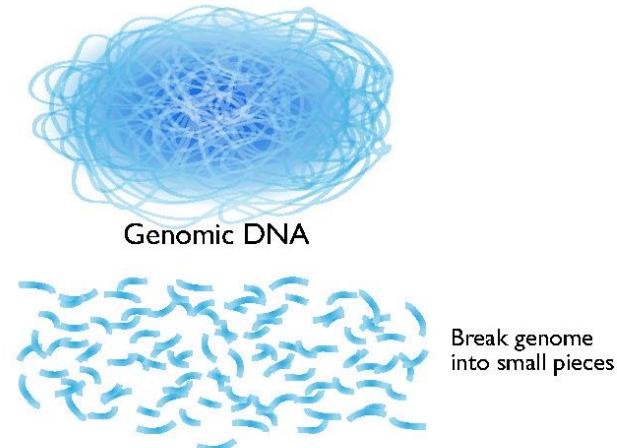
- \$300 million
- It was a proof-of-concept
- Whole genome shotgun approach
- No clones: generates data more quickly
- Direct sequencing of randomly sheared genomic DNA (\$\$).
- Libraries insert size: 2, 10 and 50 Kb
- Computationally expensive

## Human Genome Sequencing

Generating a Reference Genome Sequence (e.g., Human Genome Project)



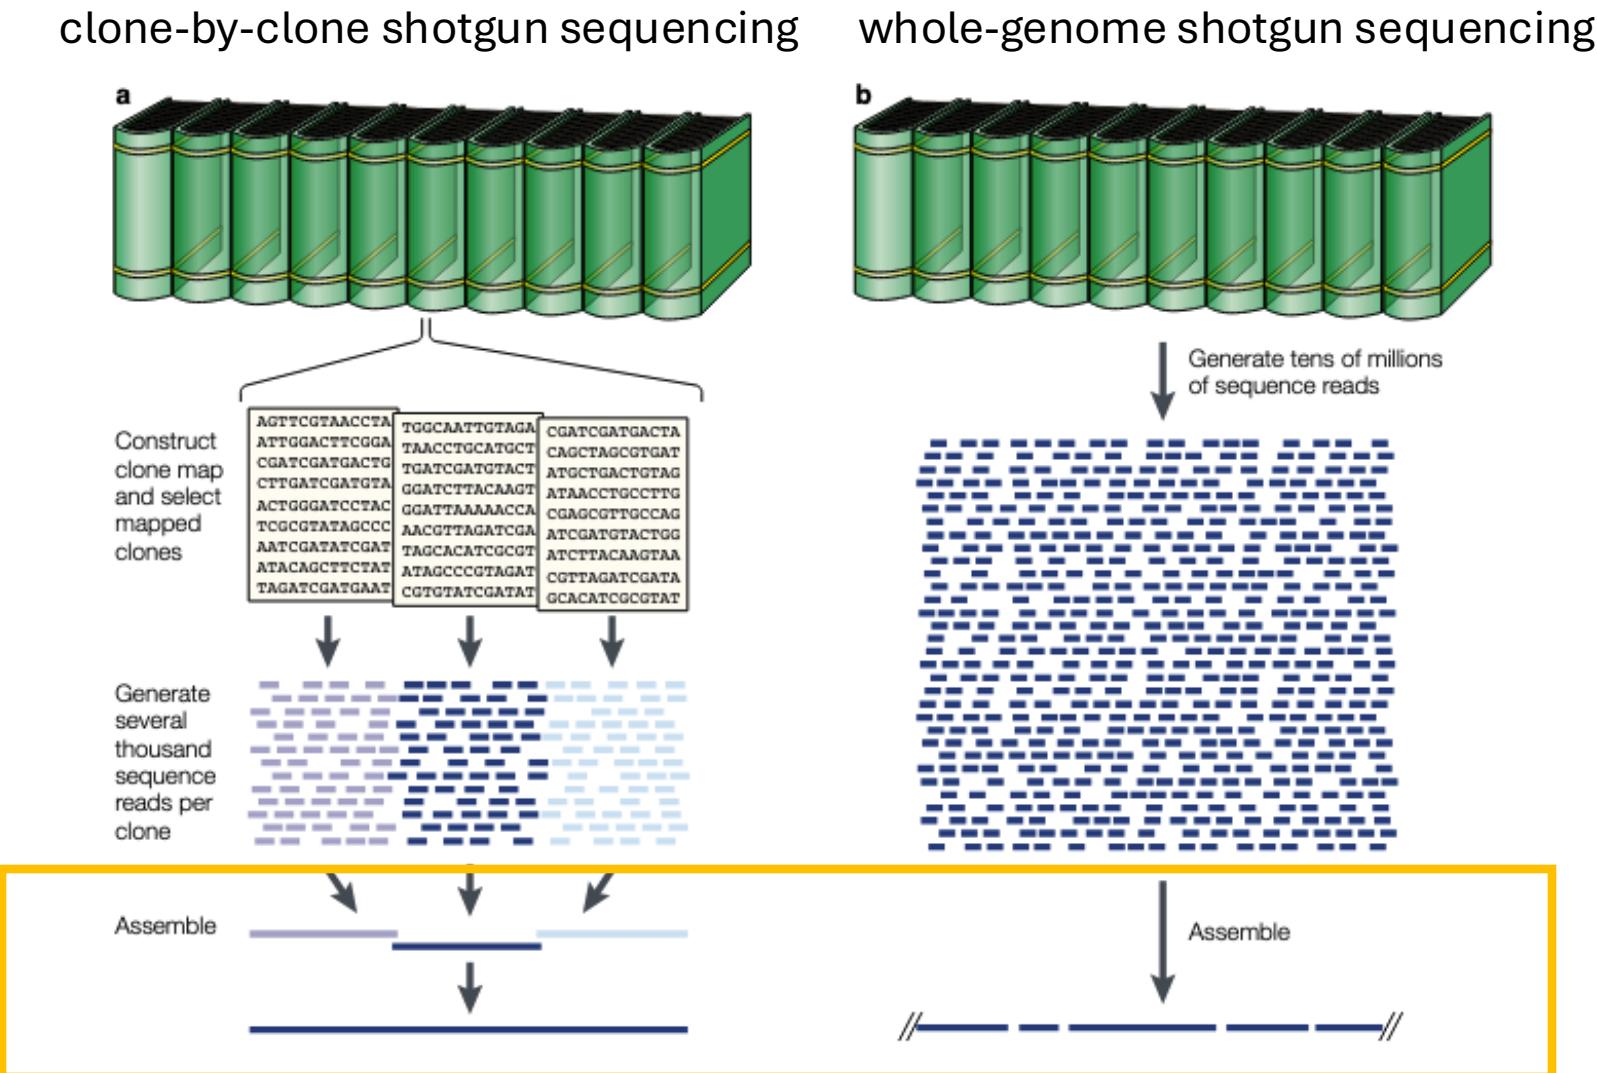
### Celera



Celera used more computational power and more sophisticated (proprietary) algorithms, but it was based on the same algorithm the Human Genome Project was using.

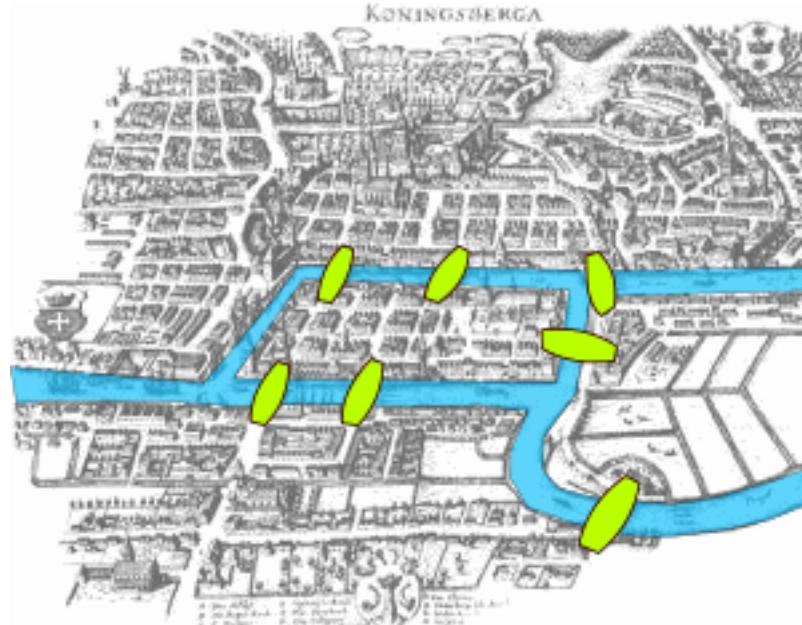
**Computational issues arise from alignment-based assembly.**

# Two main shotgun-sequencing strategies.



# Assembling billions of reads into a contiguous genome: Challenge Accepted

- The development of algorithmic ideas for High Throughput Sequencing can be traced back 300 years to the Prussian city of Königsberg (present-day Kaliningrad, Russia).
- **Bridges of Königsberg problem:**
  - Seven bridges joined the four parts of the city.
  - Residents, who enjoyed strolling through the city, wondered “**is it possible to visit every part of the city by walking across each of the seven bridges exactly once and returning to one’s starting location?**”

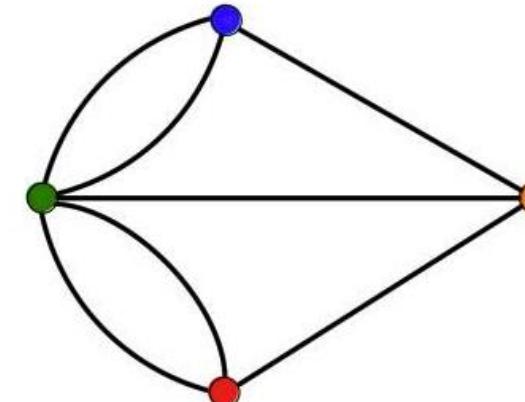
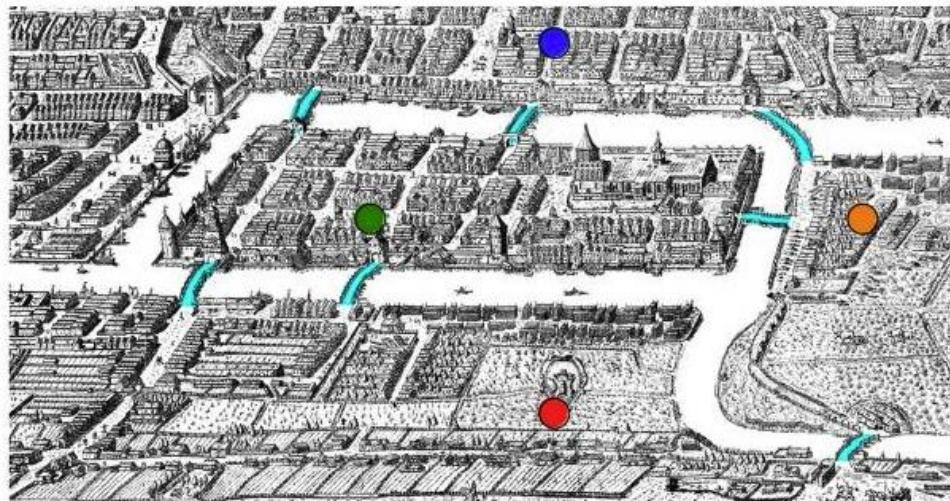


Source: Wikipedia

Remarkably, the conceptual breakthrough used by the mathematician Leonhard Euler in 1735 to solve this problem will enable the assembly of billions of short sequencing reads.

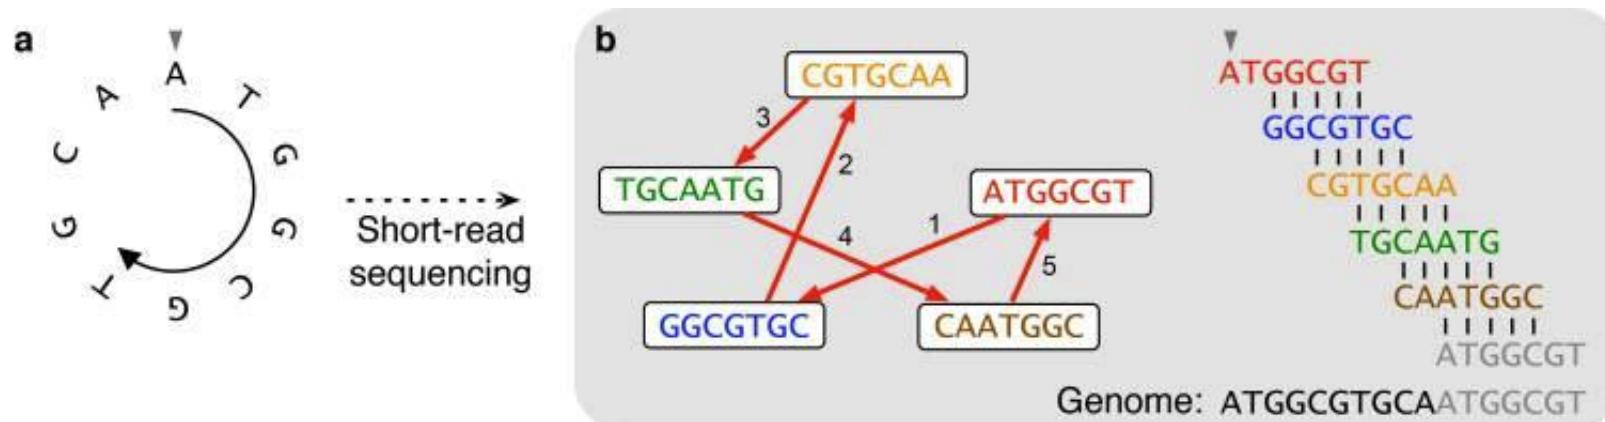
# Assembling billions of reads into a contiguous genome: Challenge Accepted

- Euler represented each landmass as a point (**node**), and each bridge as a line segment (**edge**) connecting the appropriate two points.
- This creates a **graph** – a network of nodes connected by edges.
- Euler described the procedure for determining if an arbitrary graph contains an **Eulerian cycle**: a path through the graph that visits every edge exactly once and returns back where it started.
- Euler not only *resolved* the Bridges of Königsberg Problem but also effectively launched the entire branch of mathematics known today as **graph theory**.



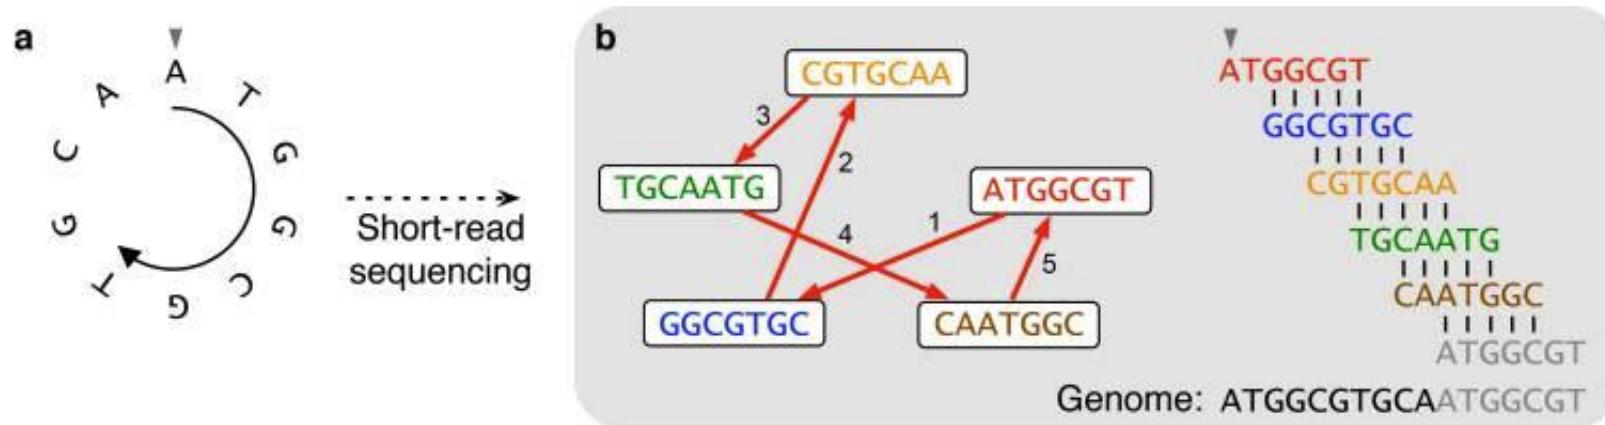
# Why are graphs useful for genome assembly?

- Suppose a really small circular genome.
- With Sanger sequencing, we produce a few short reads.
- A method for assembling reads uses graphs
  - Each read is represented by a **node**
  - Overlap between reads is represented by an arrow (**directed edge**) joining two reads.
  - Eg. two nodes may be connected with a directed edge if the reads overlap 5 bases.



# Why are graphs useful for genome assembly?

- Try to visualize the path (*let's take a stroll*) along the edges of the graph.
- In genome assembly, the path traces a series of overlapping reads -> candidate assembly.
- This path induces a **Hamiltonian cycle** in the graph: a cycle that travels to every **node exactly once** in the assembly.



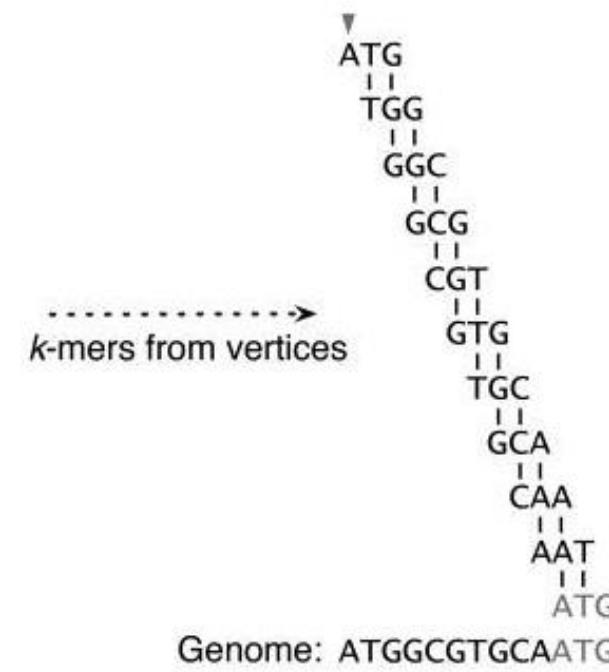
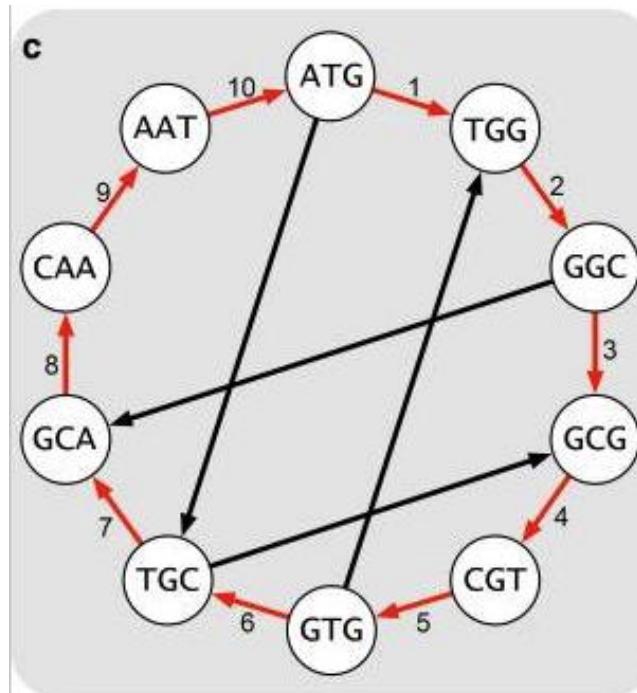
- The circular *genome* resulting from a Hamiltonian cycle contains all five reads

# “Old” assemblers

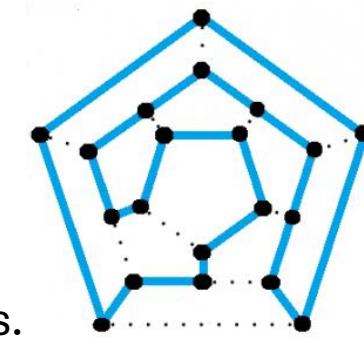
- Assemblers work with strings of length  $k$  ( **$k$ -mers**):
  - Shorter than reads (eg. an Illumina 100bp read can be divided into 46 overlapping 55-mers)
- From a set of reads, form a node for every  $k$ -mer in the reads.
- Connect one  $k$ -mer to another if the two  $k$ -mers completely overlap except for one nucleotide at each end.
- A **Hamiltonian cycle (red edges)** will represent the candidate genome since it travels each  $k$ -mer exactly once.

Hamiltonian cycle  
Visit each node once

ATGGCGT  
 $k = 3$ , comprises  
ATG,  
TGG,  
GGC,  
GCG,  
CGT



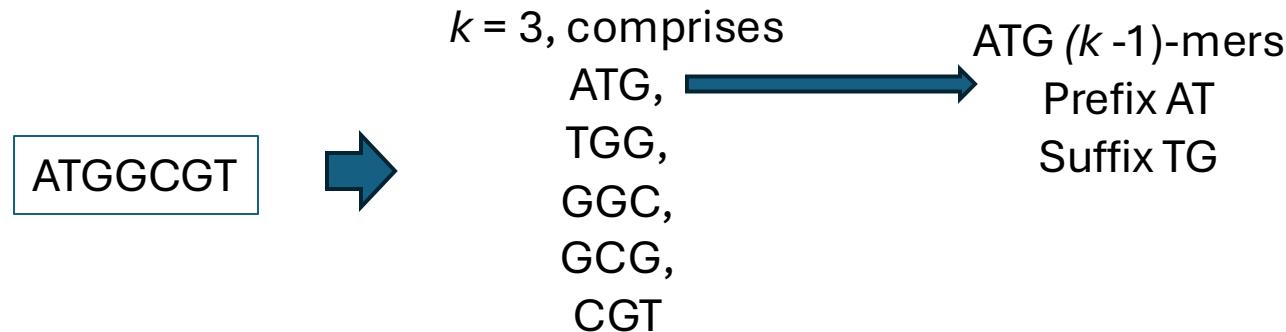
# Hamiltonian cycle



- This method is easy to implement for small genomes/small number of reads.
- The Hamiltonian cycle approach was feasible for sequencing the first microbial genome in 1995 and the human genome in 2001.
- However, the algorithm for finding a Hamiltonian cycle in a large graph with millions of nodes (e.g., Illumina sequencing) would not be efficient.
- The computational burden was so large that most HTS projects have abandoned the Hamiltonian cycle approach.
- Maybe finding a cycle visiting all edges (*bridges*) of a graph once is much easier.

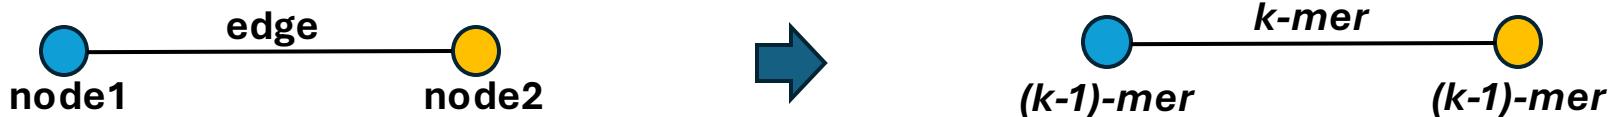
# de Bruijn graphs

- The origin of de Bruijn graphs dated back to 1946, when the mathematician Nicolaas de Bruijn became interested in the *Superstring Problem*: find a shortest circular *superstring* that contains all possible *substrings* of length  $k$  ( $k$ -mers) in a given alphabet.
- In an alphabet containing  $n$  symbols, there are  $n^k$   $k$ -mers
- In DNA: the alphabet {A,T,G,C}, for  $k = 3$  there are  $4^3 = 64$  trinucleotides.
- The de Bruijn graph represents all  $k$ -mer prefixes and suffixes as nodes and draws edges representing  $k$ -mers with a particular prefix and suffix.
- For example,  $k$ -mer edge ATG has prefix AT and suffix TG



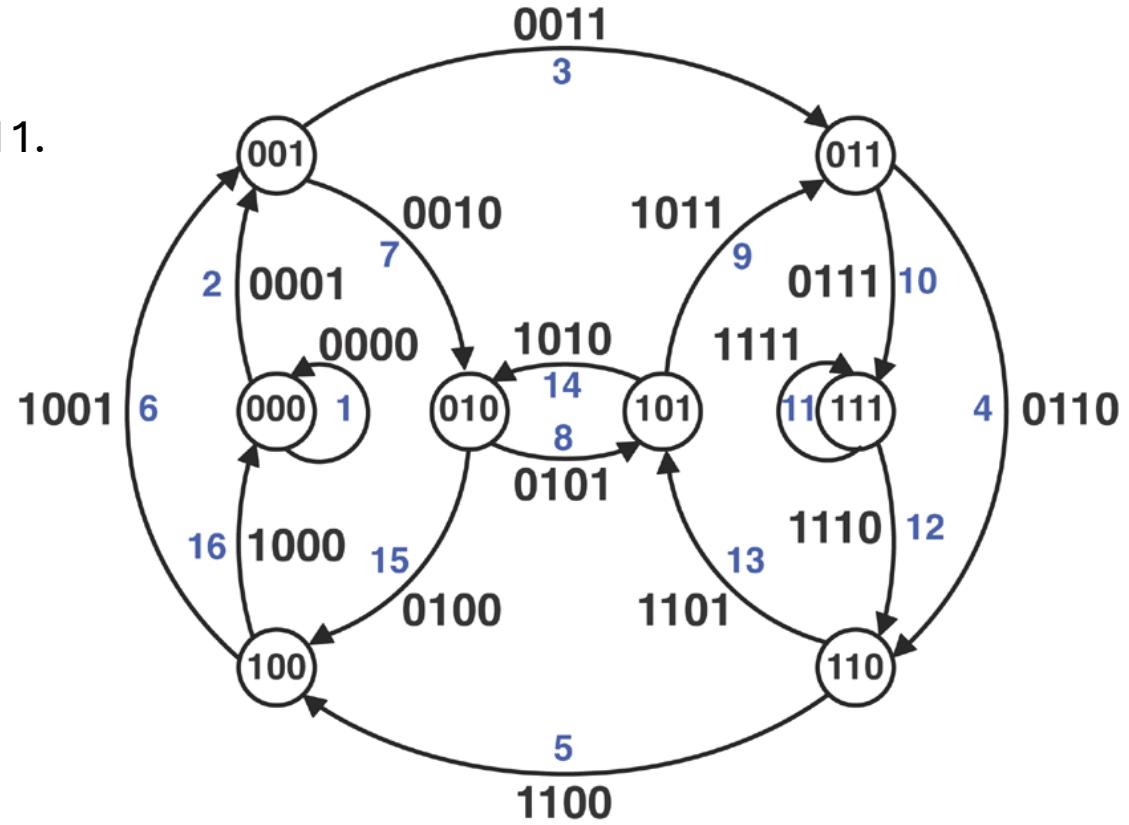
# de Bruijn graphs

- The biologist's decision to fragment the genome into smaller pieces motivated computer scientists to cast fragment assembly as such a problem.
- Instead of assigning each  $k$ -mer to a node, they assign each  $k$ -mer (within a read) to an edge.
- In the construction of a de Bruijn graph, every possible  $(k-1)$ -mer is assigned to a node.
- Connect one  $(k-1)$ -mer by a directed edge to a second  $(k-1)$ -mer if there is some  $k$ -mer whose prefix is the former and whose suffix is the latter.
- A crucial aspect of the de Bruijn graph is that it is an **Eulerian graph**, which implies it has an **Eulerian cycle**.
- An Eulerian cycle implies that there exists an Eulerian path that visits each edge exactly once and returns to the starting point without visiting any vertex twice.



# de Bruijn graphs

- two-character alphabet: **0** and **1**
- All possible 3-mers: 000, 001, 010, 011, 100, 101, 110, 111.
- The circular superstring 0001110100 :
  - contains all 3-mers
  - each 3-mer exactly once (graph short as possible)
- how can one construct such a superstring for all  $k$ -mers
  - if an arbitrary value of  $k$ ?
  - an arbitrary alphabet?
- De Bruijn answered this question by borrowing Euler's solution of the Bridges of Königsberg problem.



Source: <https://doi.org/10.1038/nbt.2023>

# de Bruijn graphs

- For example,  $S=\{A,B\}$  and  $k = 3$
- Take each length-3 input string and split it into two overlapping substrings of length 2 ( $k-1$ ). Call these the left and right 2-mers (prefix and suffix).

For a linear genome: AAABBBA and  $k = 3$

3-mers: AAA, AAB, ABB, BBB, BBA  
( $k-1$ )-mer / 2-mer: AA, AA AA, AB AB, BB BB, BB, BA



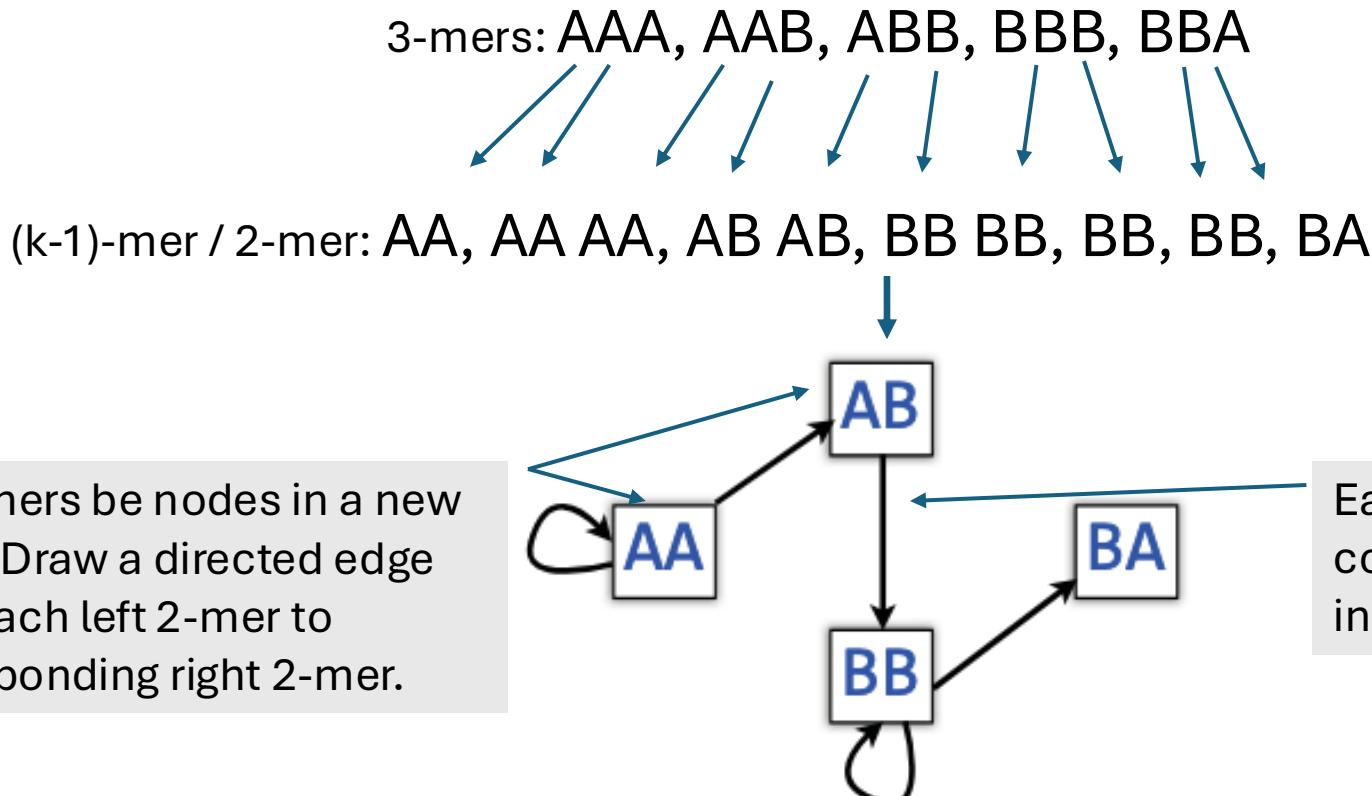
Represent  $k$ -mer in a graph:

- One edge per  $k$ -mer
- One node per distinct  $(k-1)$ -mer

# de Bruijn graphs

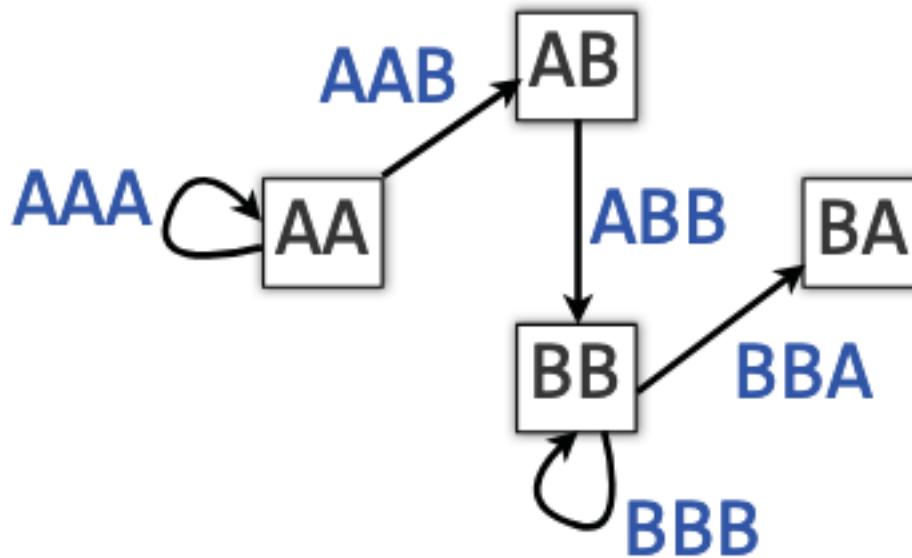
- For example,  $S=\{A,B\}$  and  $k = 3$
- Take each length-3 input string and split it into two overlapping substrings of length 2. Call these the left and right 2-mers (prefix and suffix).

For a genome: AAABBBA



# de Bruijn graphs

An edge corresponds to an overlap (of length  $k-2$ ) between two  $k-1$  mers.  
More precisely, it corresponds to a  $k$ -mer from the input.



This is an Eulerian cycle!

Walk crossing each edge exactly once gives a reconstruction of the genome  
AAABBBBA

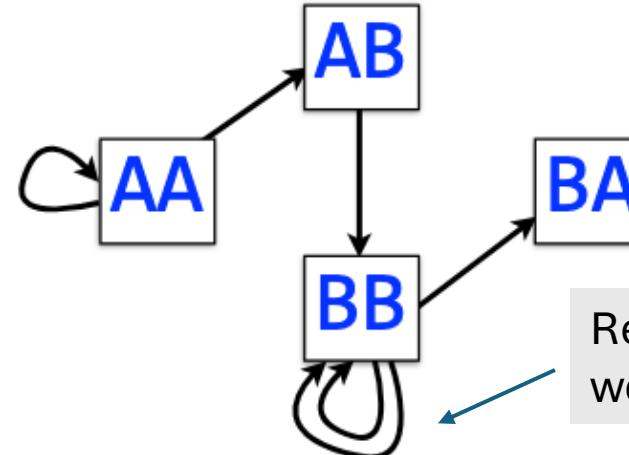
# de Bruijn graphs

- If we add one more B to our input string:

For a genome: AAABBBBA

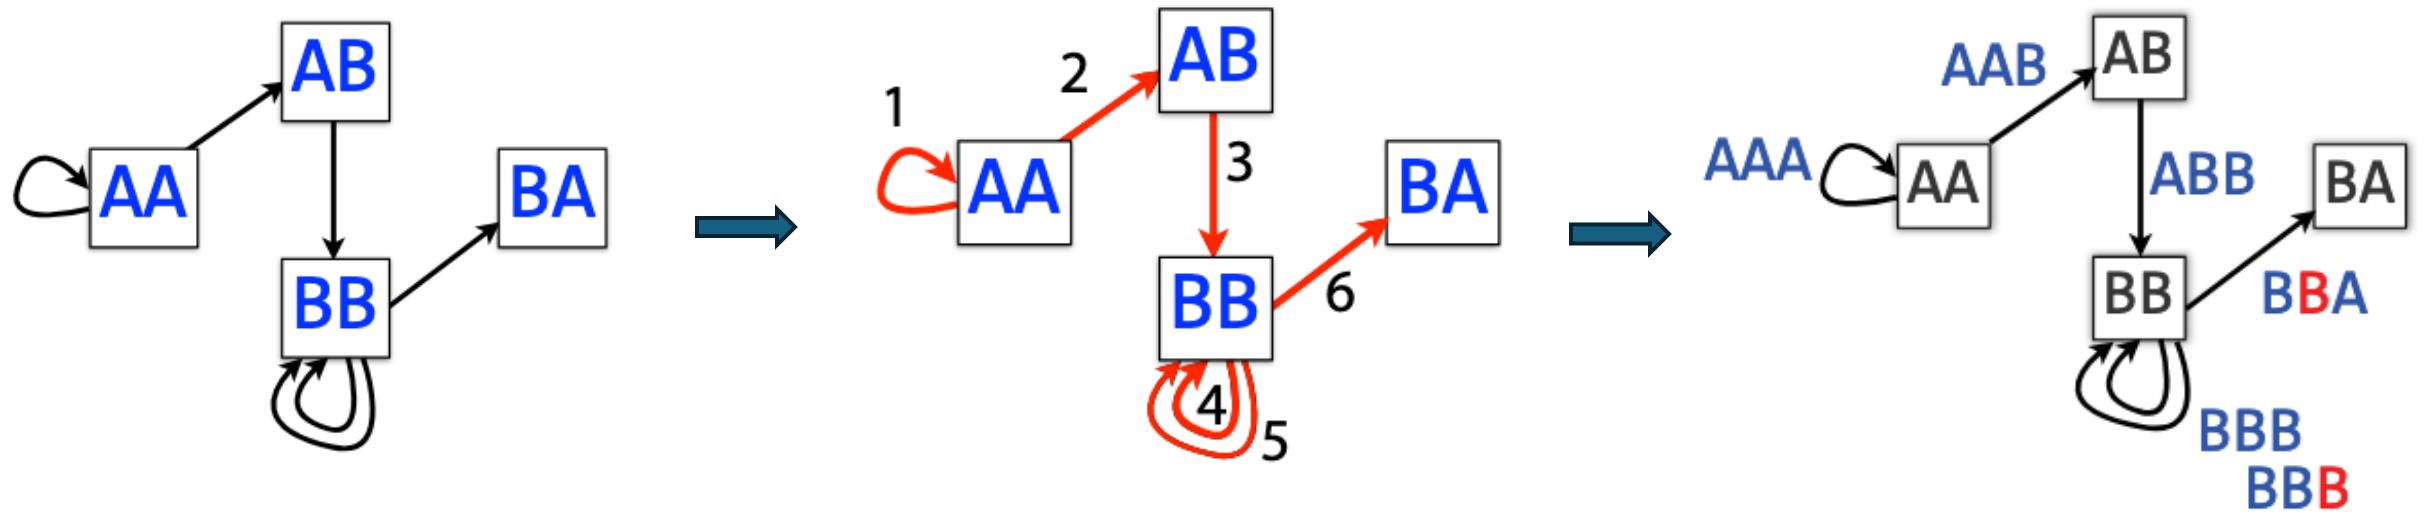
3-mers: AAA, AAB, ABB, BBB, BBB, BBA  
( $k-1$ )-mer / 2-mer: AA, AA AA, AB AB, BB BB, BB BB, BB BB, BA

Nodes (2-mers) in a new graph: no change.



Rebuild the De Bruijn graph accordingly:  
we get a multi-edge.

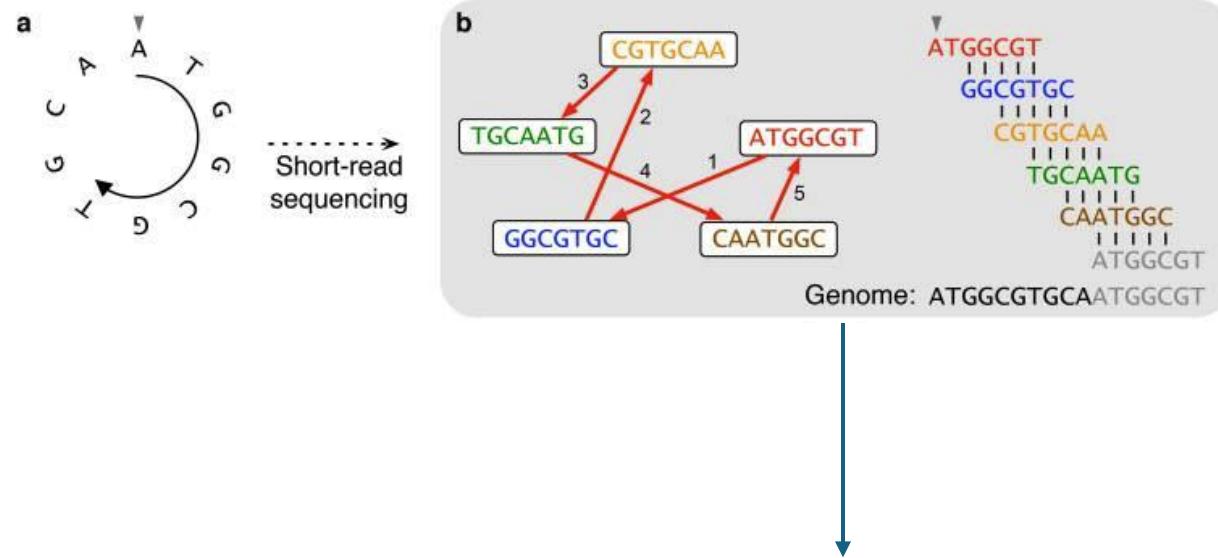
# de Bruijn graphs



This is an Eulerian cycle!

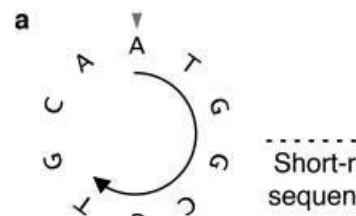
Walk crossing each edge exactly once gives a reconstruction of the genome  
**AAABBBBA**

# de Bruijn graphs

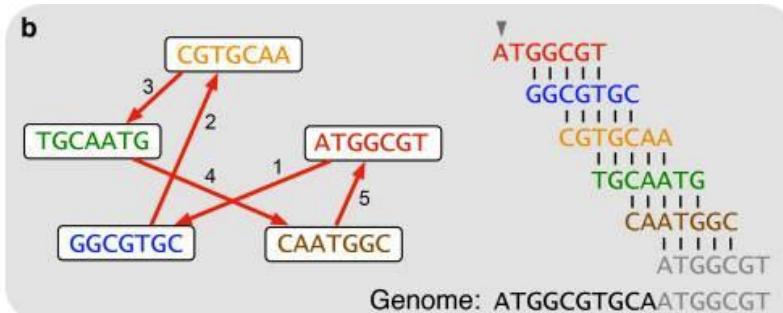


This time, nodes are  $(k-1)$ -mers  
Edges are  $k$ -mers

# de Bruijn graphs

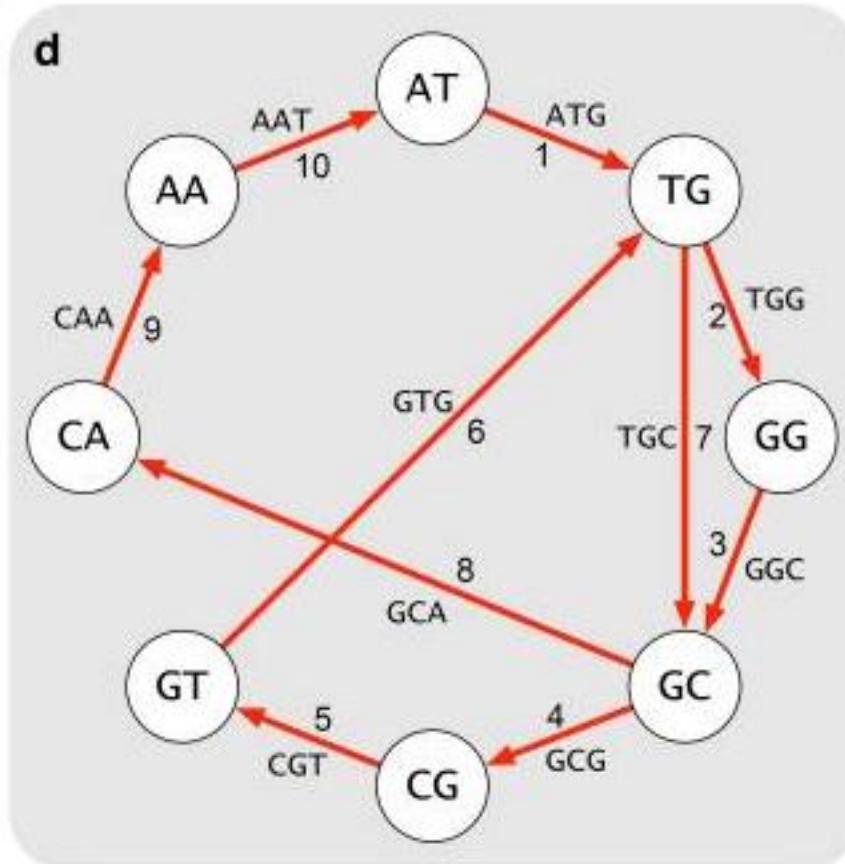


Short-read sequencing

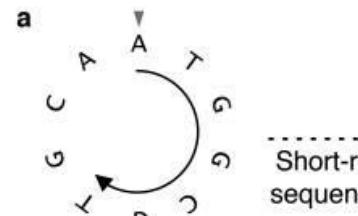


This time, nodes are  $(k-1)$ -mers  
Edges are  $k$ -mers  
 $k = 3$

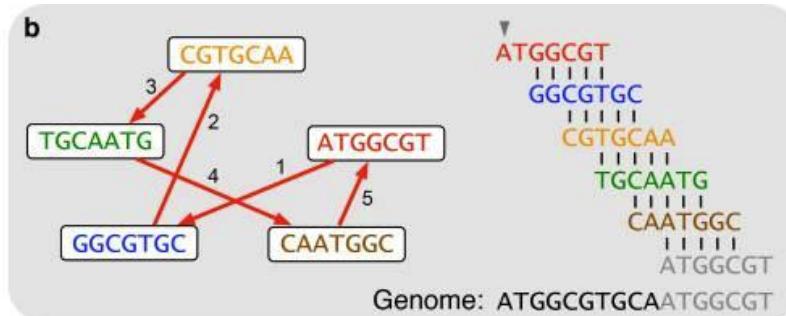
Eulerian cycle  
Visit each edge once



# de Bruijn graphs



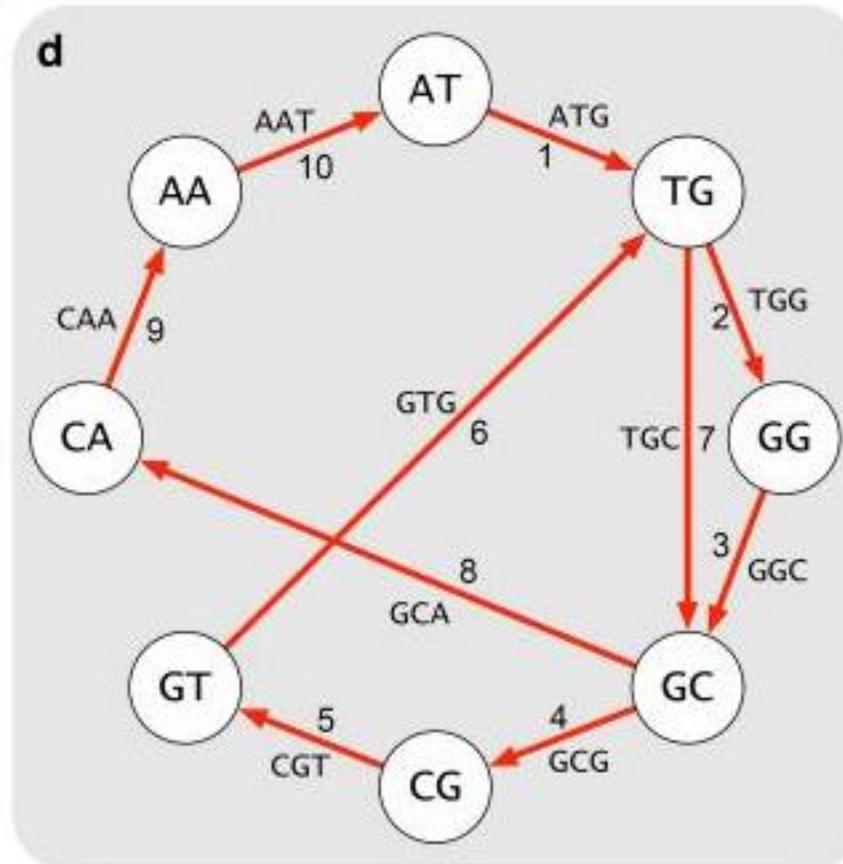
Short-read sequencing



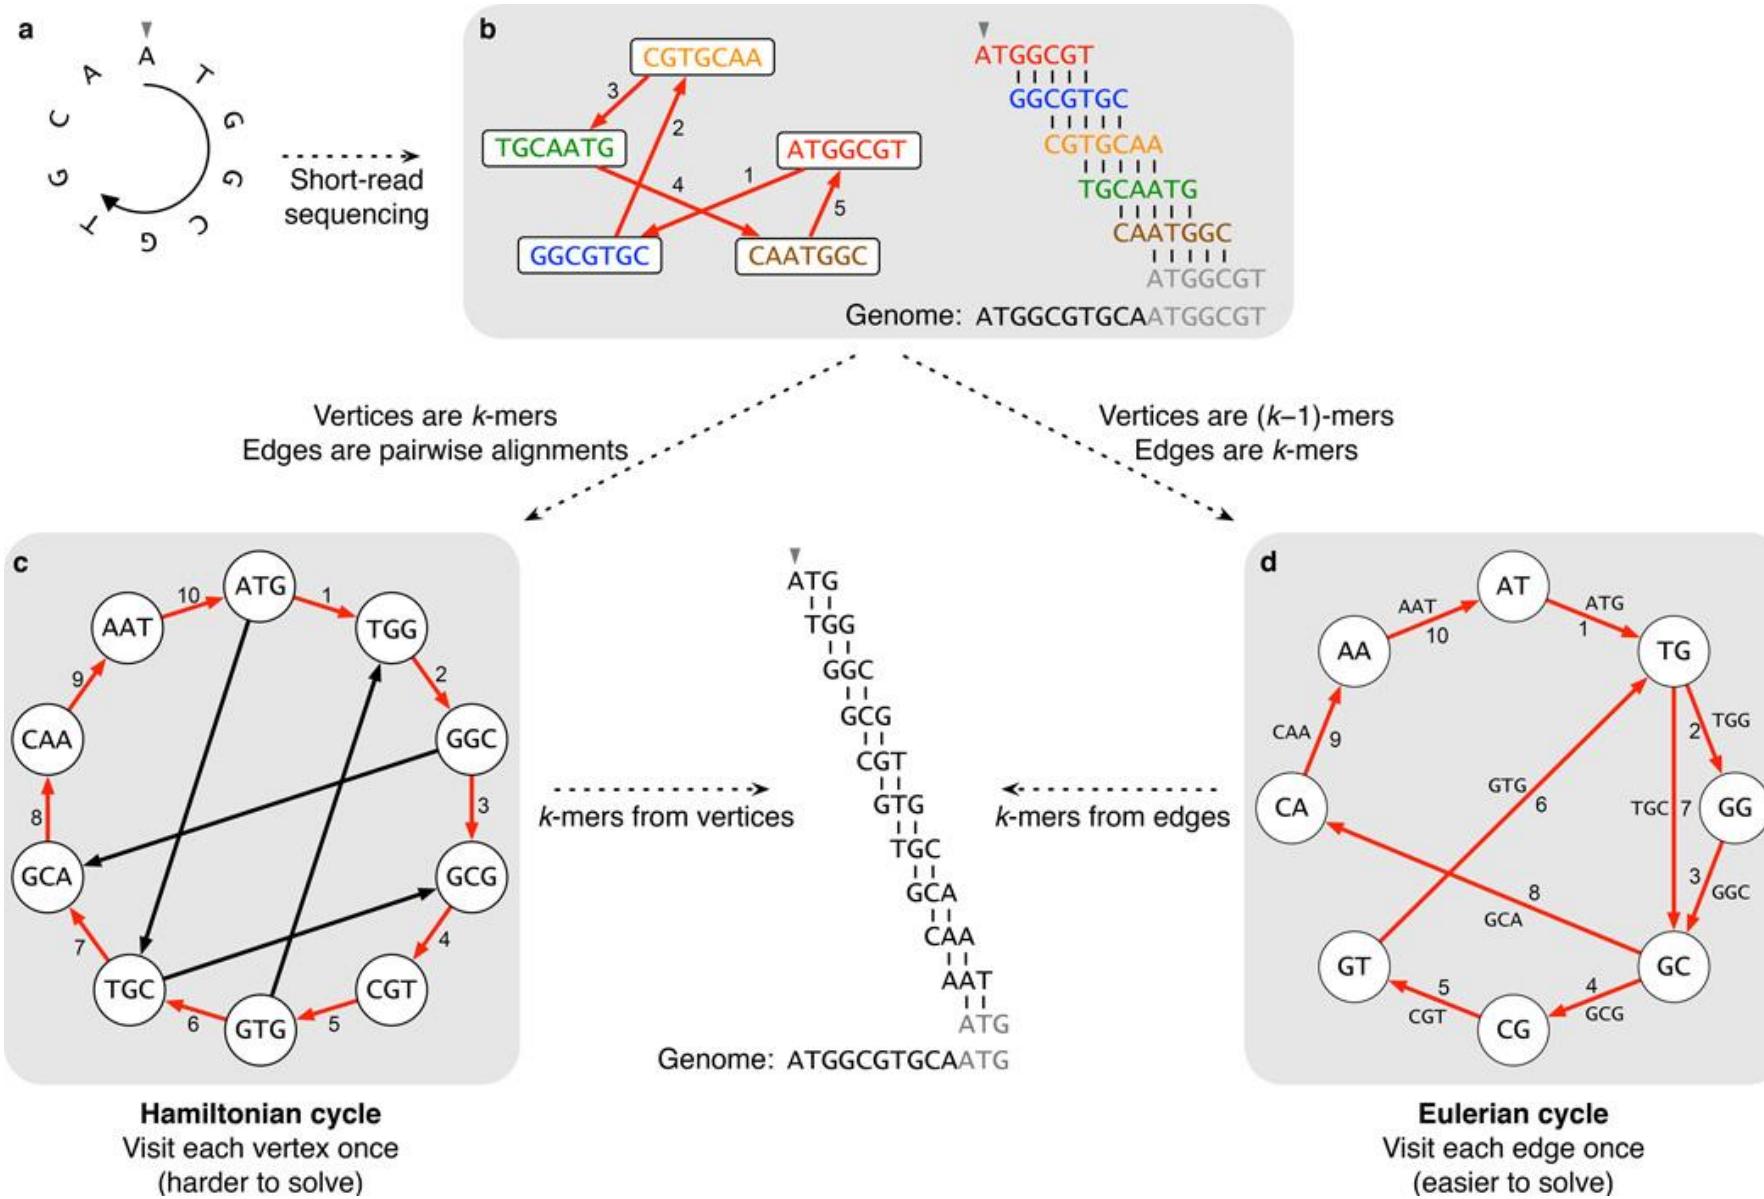
Easier to solve!

Eulerian cycle  
Visit each edge once

This time, nodes are  $(k-1)$ -mers  
Edges are  $k$ -mers



## Two strategies for genome assembly: from Hamiltonian cycles to Eulerian cycles.

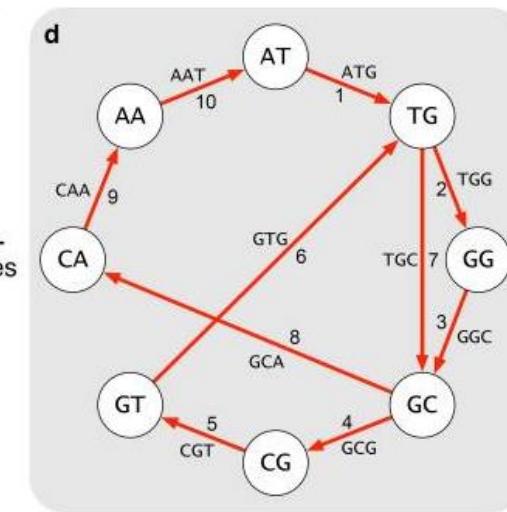
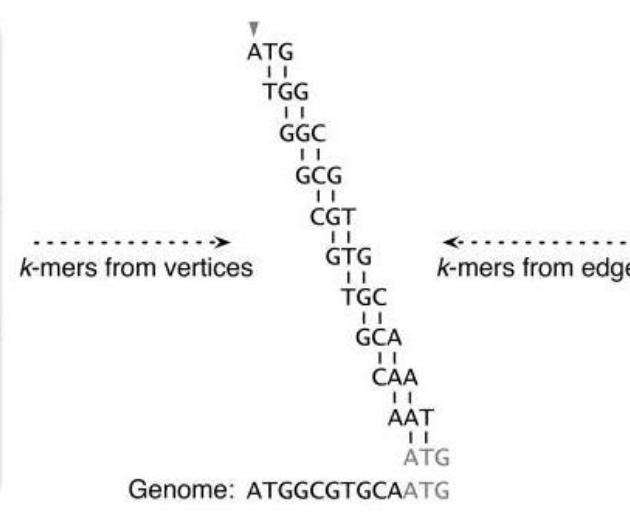
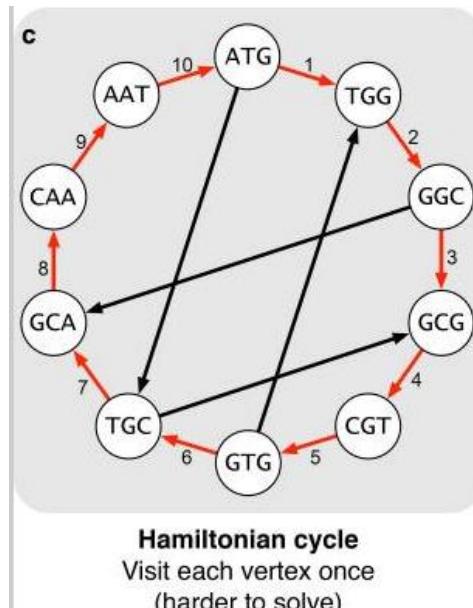


# de Bruijn graphs

- Modern short-read assembly algorithms construct a de Bruijn graph by representing all  $k$ -mer prefixes and suffixes as nodes.
- Then drawing edges that represent  $k$ -mers having a particular prefix and suffix.
- Finding an Eulerian cycle allows one to reconstruct the genome by forming an alignment in which each successive  $k$ -mer (from successive edges) is shifted by one position.
- This generates the same cyclic genome sequence without the computational strain of finding a Hamiltonian cycle.

## Two strategies for genome assembly: from Hamiltonian cycles to Eulerian cycles.

- The run time required by a computer implementation of Euler's algorithm is roughly proportional to the number of edges in the graph.
- In the Hamiltonian approach, the time is potentially a lot larger, due to the large number of pairwise alignments needed to construct the graph, and to the *NP*-Completeness of finding a Hamiltonian cycle.
- The Hamiltonian Path Problem is NP-complete , which means that the efficient algorithms for solving this problem are unknown.

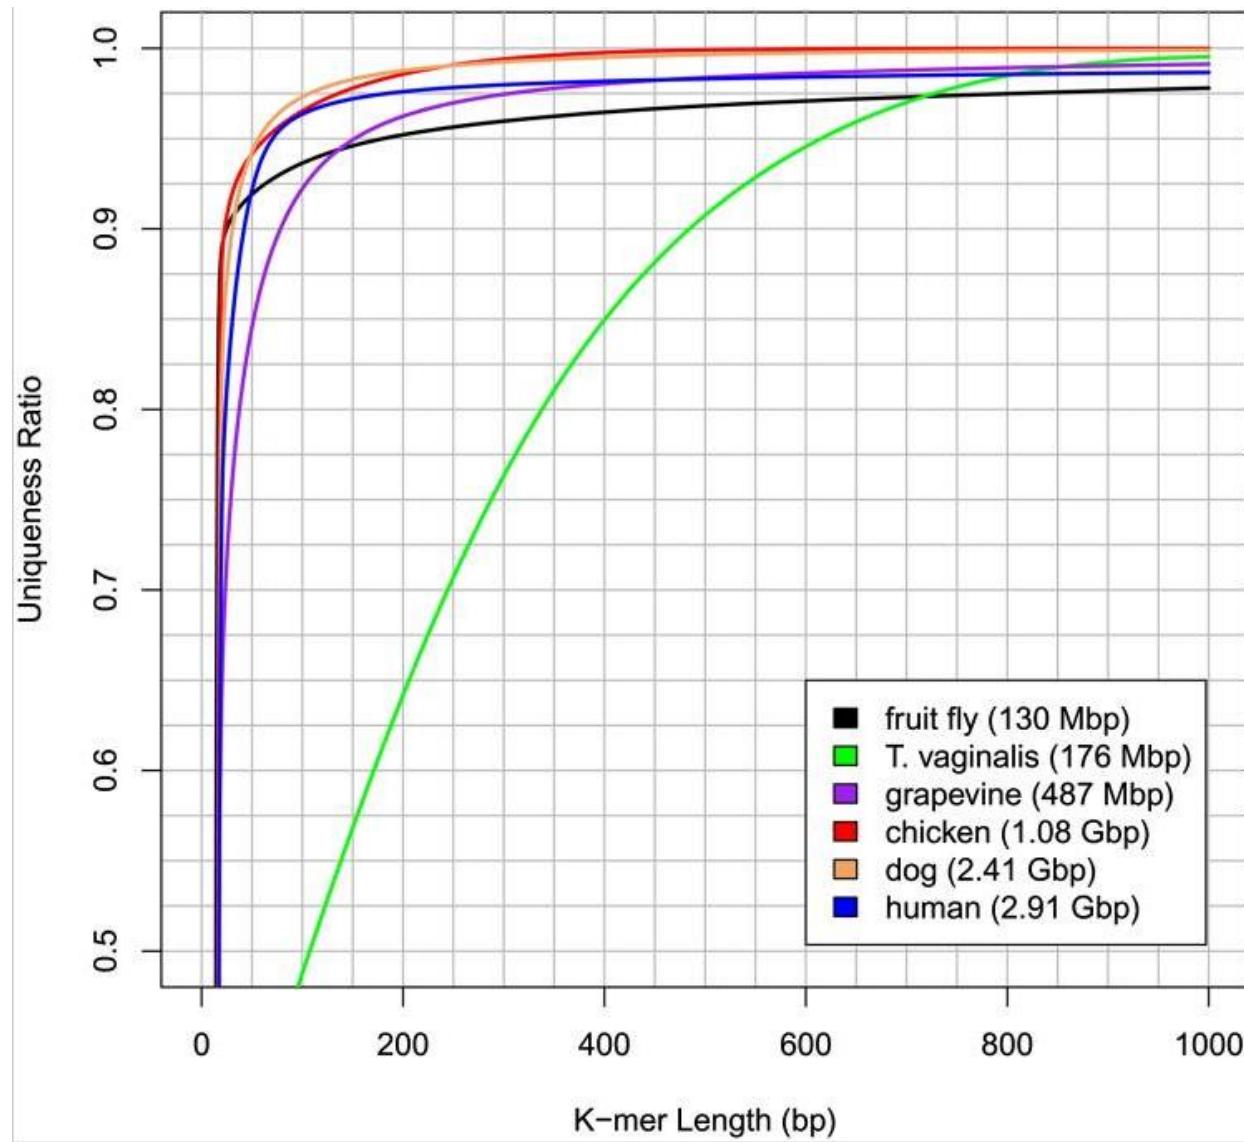


# Assembly methods

- Current genome sequencing technology can only sequence a tiny portion of a genome in a contiguous read.
- DNA sequence reads may fit together in more than one way because of repetitive sequences within the genome.
- Assembly methods aim to create the most complete reconstruction possible without introducing errors.
- The central challenge of genome assembly is resolving repetitive sequences.

How can we measure the repetitiveness among species?

## Uniqueness ratio for varying read lengths in six genomes



The figure shows how much of each genome would be covered by  $k$ -mers (reads) that occur exactly once.

Among the multicellular species, dog and chicken are the least repetitive while fly is the most repetitive.

The percentage of a genome covered uniquely increases rapidly as read length increases to 50 bp and above.

But the rate of increase varies due to the variable repeat lengths in different species.

# Assembly methods

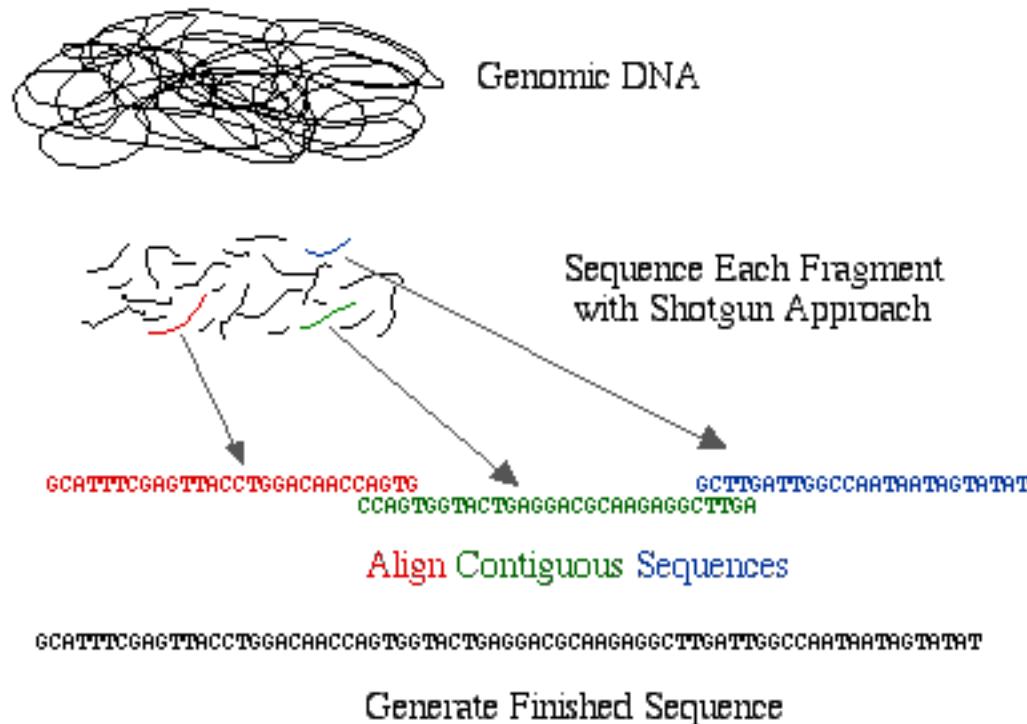
- Early genome assemblers used a simple “greedy” algorithm, in which all pairs of reads are compared with each other, and the ones that **overlap** most are merged first.
- To allow for sequencing errors, assemblers compute these overlaps with a variant of the **Smith-Waterman** algorithm.
- Once all overlaps are computed, the reads with the longest overlap are concatenated to form a contig (contiguous sequence).
- This simple merging process will accurately reconstruct the simplest genomes, but fails for repetitive sequences longer than the read length.

|                  |          |
|------------------|----------|
| R <sub>3</sub> : | CCTACAAG |
| R <sub>4</sub> : | CTACAAGT |
| A:               | TACAAGTT |
| B:               | ACAAGTTA |
| C:               | CAAGTTAG |
| X:               | TACAAGTC |
| Y:               | ACAAGTCC |
| Z:               | CAAGTCCG |

# Large-scale shotgun assembly

- Several assemblers have been developed to assemble large, repetitive genomes from long (Sanger) reads, including the Celera Assembler.

## Whole Genome Shotgun Sequencing Method



# Short read assembly

- In principle, assemblers created for long reads should also function for short reads.
- They need to find overlap and build a consensus.
- They even try to accommodate short reads by fine-tune different parameters (overlaps, seed reads, coverage).
- New generation of genome assemblers has been developed specifically to address the challenges of assembling short reads.
- Rather than using an overlap graph, all assemblers use **de Bruijn graph** algorithm.

# Two approaches to read assembly

## Layout - Overlap - Consensus:

- Construct overlap graph directly from reads, eliminating redundant reads.
- Trace path for assembly.
- String graphs assemblers

## De Bruijn graph-based assemblers:

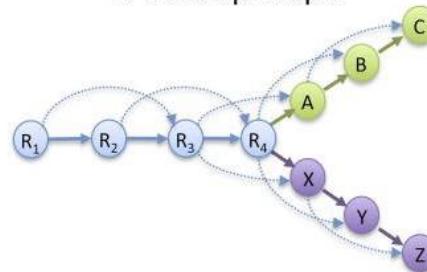
- Construct k-mer graph from reads
- Original reads are discarded
- Trace path in graph for assembly

A Read Layout

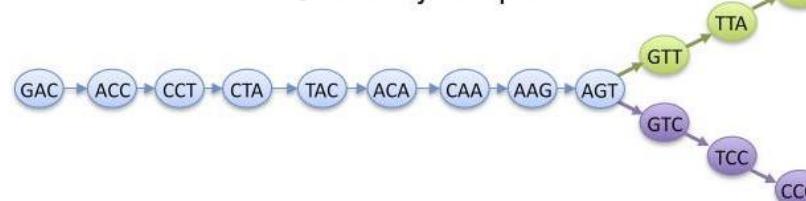
R<sub>1</sub>: GACCTACA  
R<sub>2</sub>: ACCTACAA  
R<sub>3</sub>: CCTACAAAG  
R<sub>4</sub>: CTACAAAGT  
A:  
B:  
C:  
X:  
Y:  
Z:

TACAGTT  
ACAAGTTA  
CAAGTTAG  
TACAAGTC  
ACAAGTCC  
CAAGTCCG

B Overlap Graph



C de Bruijn Graph

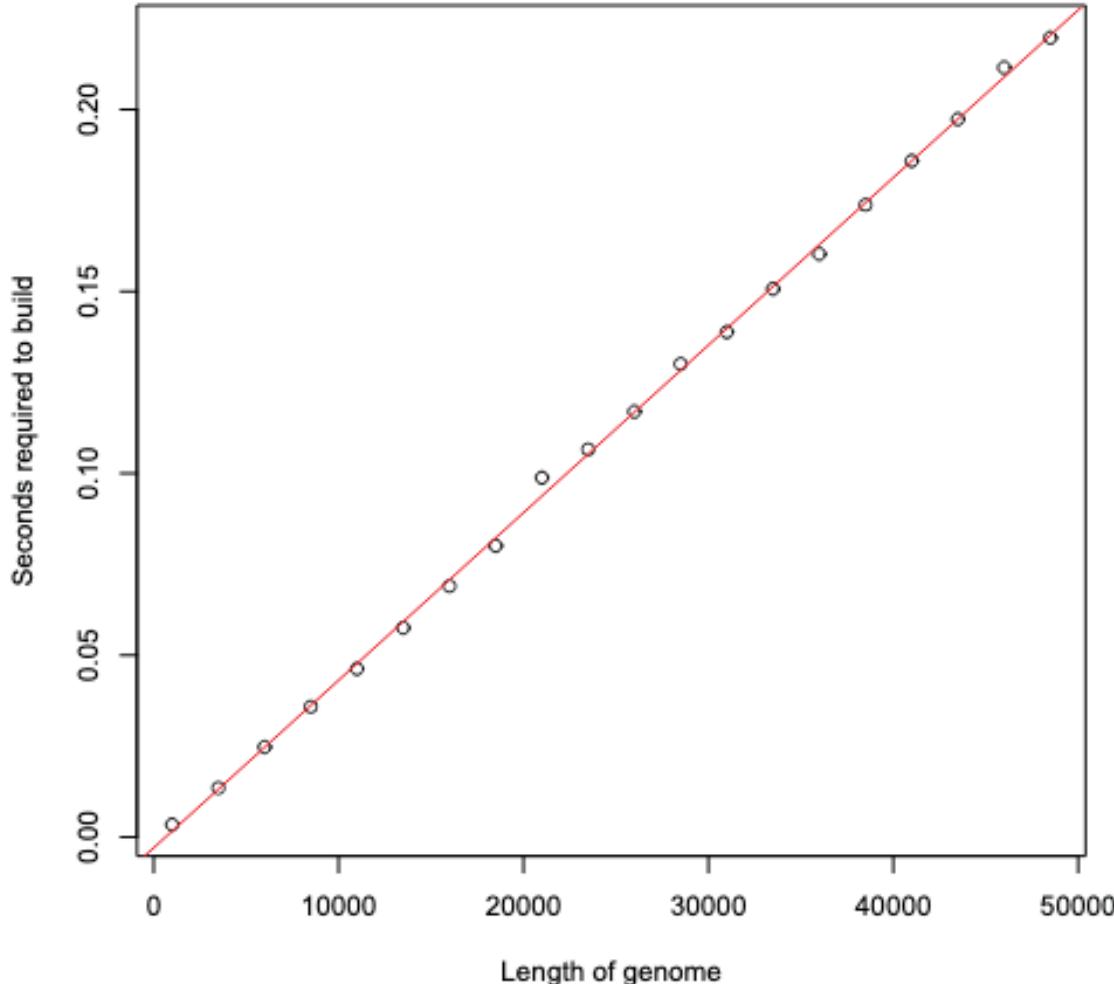


- Differences between an overlap graph and a de Bruijn graph for assembly.
- In both approaches, repeat sequences create a fork in the graph.

# Limitations of de Bruijn graphs

- Sequencing error complicates the de Bruijn graph (errors create their own graph nodes), but many errors are easily recognized by their structure in the graph : **dead-end** “tips” in the graph.
- The main drawback to the de Bruijn approach is the loss of information caused by decomposing a read into a path of  $k$ -mers.
  - can't resolve repeats as well as overlap graph.
- Another potential drawback of the de Bruijn approach is that the de Bruijn graph can require an enormous amount of computer space (RAM).
- In theory, the size of the de Bruijn graph depends only on the size of the genome, including polymorphic alleles.
- Single most important benefit of De Bruijn graph is speed and simplicity.

# Applying de Bruijn graphs



- Timed de Bruijn graph construction applied to progressively longer prefixes of **lambda phage genome** (50 Kb).
- $k = 14$
- $O(N)$  expectation appears to work in practice.

# Lab12: Genome assembly

