

Bioinformatics Algorithms

COS-BIOL-530/630

Lecture 13

Days & Times	Room	Meeting Dates
Tu 2:00PM - 3:50PM	Thomas Gosnell Hall (GOS)-2178	01/13/2025 - 04/28/2025
Th 2:00PM - 3:50PM	Thomas Gosnell Hall (GOS)-2178	01/13/2025 - 04/28/2025

Instructor:

Fernando Rodriguez

email: frvsbi@rit.edu

Office: Orange Hall 1311

RNA-seq - Lecture 13 -

Quiz 10: Lectures 12 and 13. Available April 23rd
(Due on April 27th 5 pm)

Lab 12 Assignments:

- Discussion 12
- Activity 12

Final Project:

- Draft (due on April 24)
- Report (due on May 1st 11:59 pm)

No Class on Thursday April 24th

SRATE evaluation period will be ending soon (4/29)!
rit.smartevals.com



RNA-seq

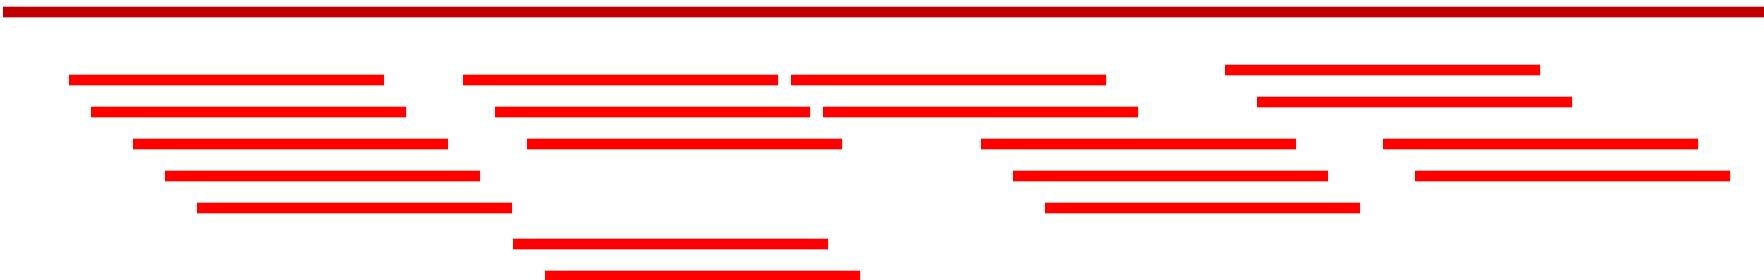
- Lecture 13 -

Topics:

- The short-read mapping (alignment) problem
- The Burrows-Wheeler Transform (BWT)
- RNA sequencing

Short read mapping

Reference genome



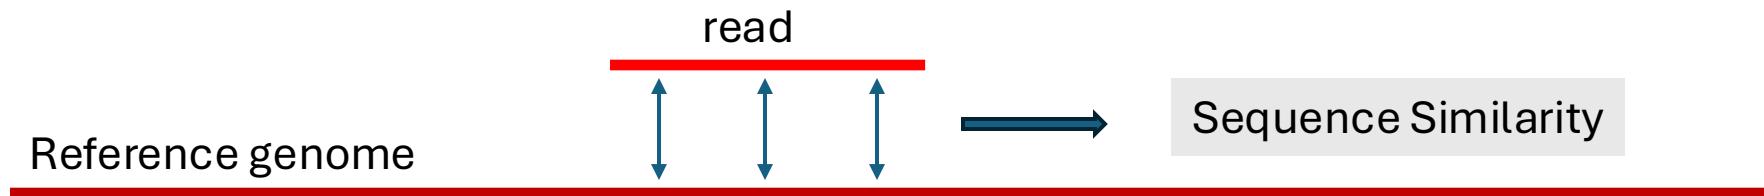
HTS technologies
produce
million of reads
=
a short substring of
the genome

If we already have one reference genome:

- Map genomic read to the reference (**re-sequencing**)
 - Nucleotide Polymorphism
 - Find structural variants
 - Indels
- Map RNA (coding and non-coding) to reference
- Problem: Million of string searches in a long string

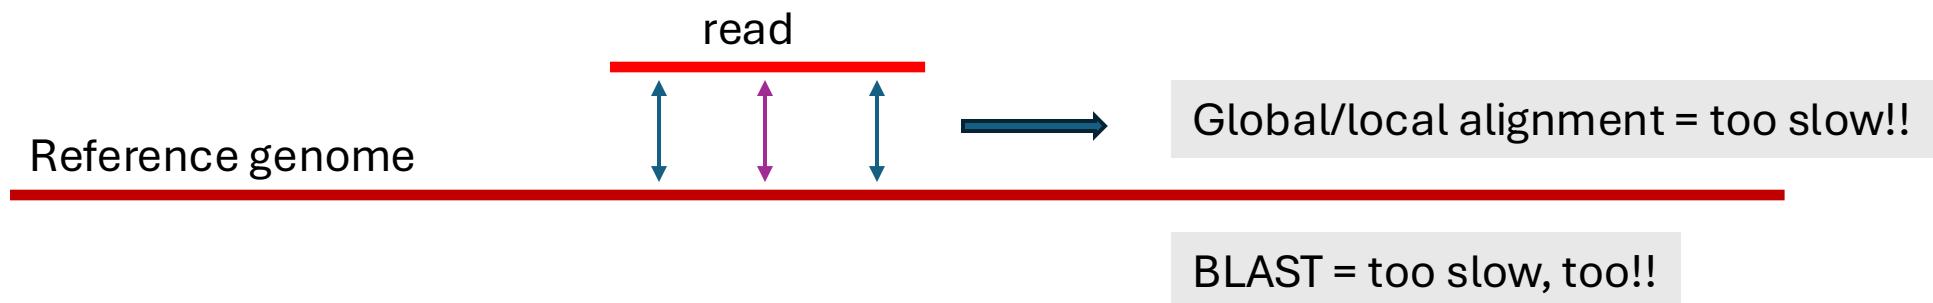
Short read mapping: the problem

- First, genomic DNA is sheared into fragments that can be size-selected to obtain a template library of uniformly short DNA fragments.
- During sequencing, every molecule is cyclically extended by a single base, and all clusters are read in parallel before the cycle is repeated (Illumina).
- The first step in the analysis of NGS data is to determine where in the genome each of the short reads originated.



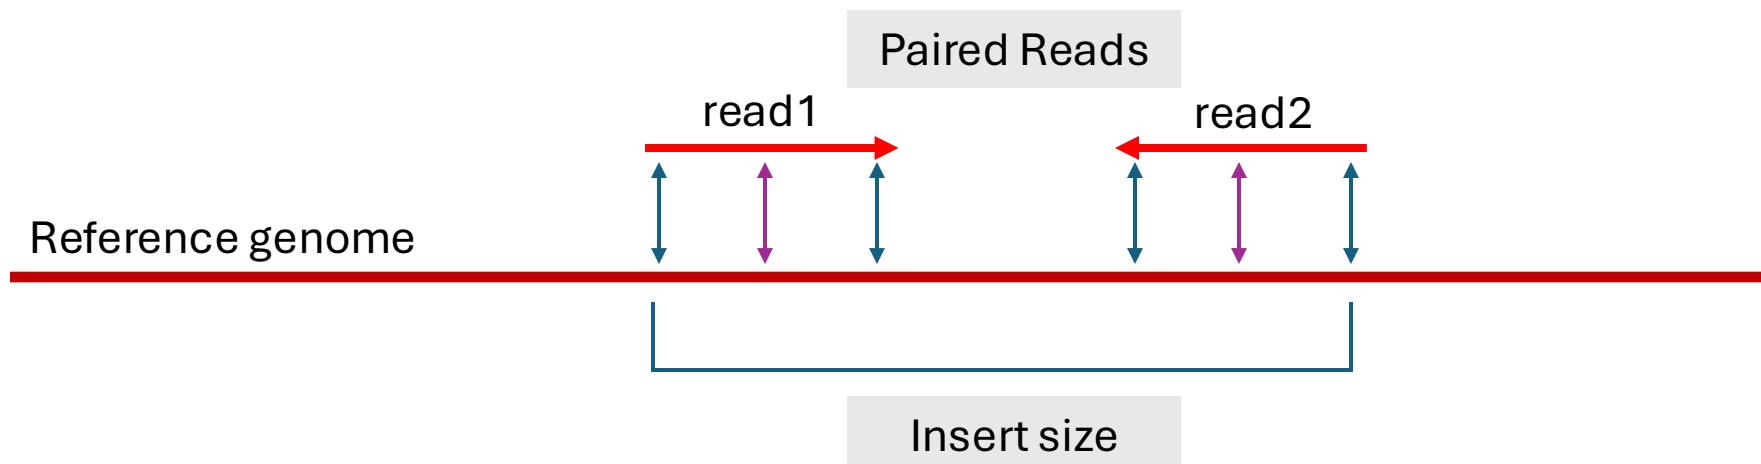
Short read mapping: the problem

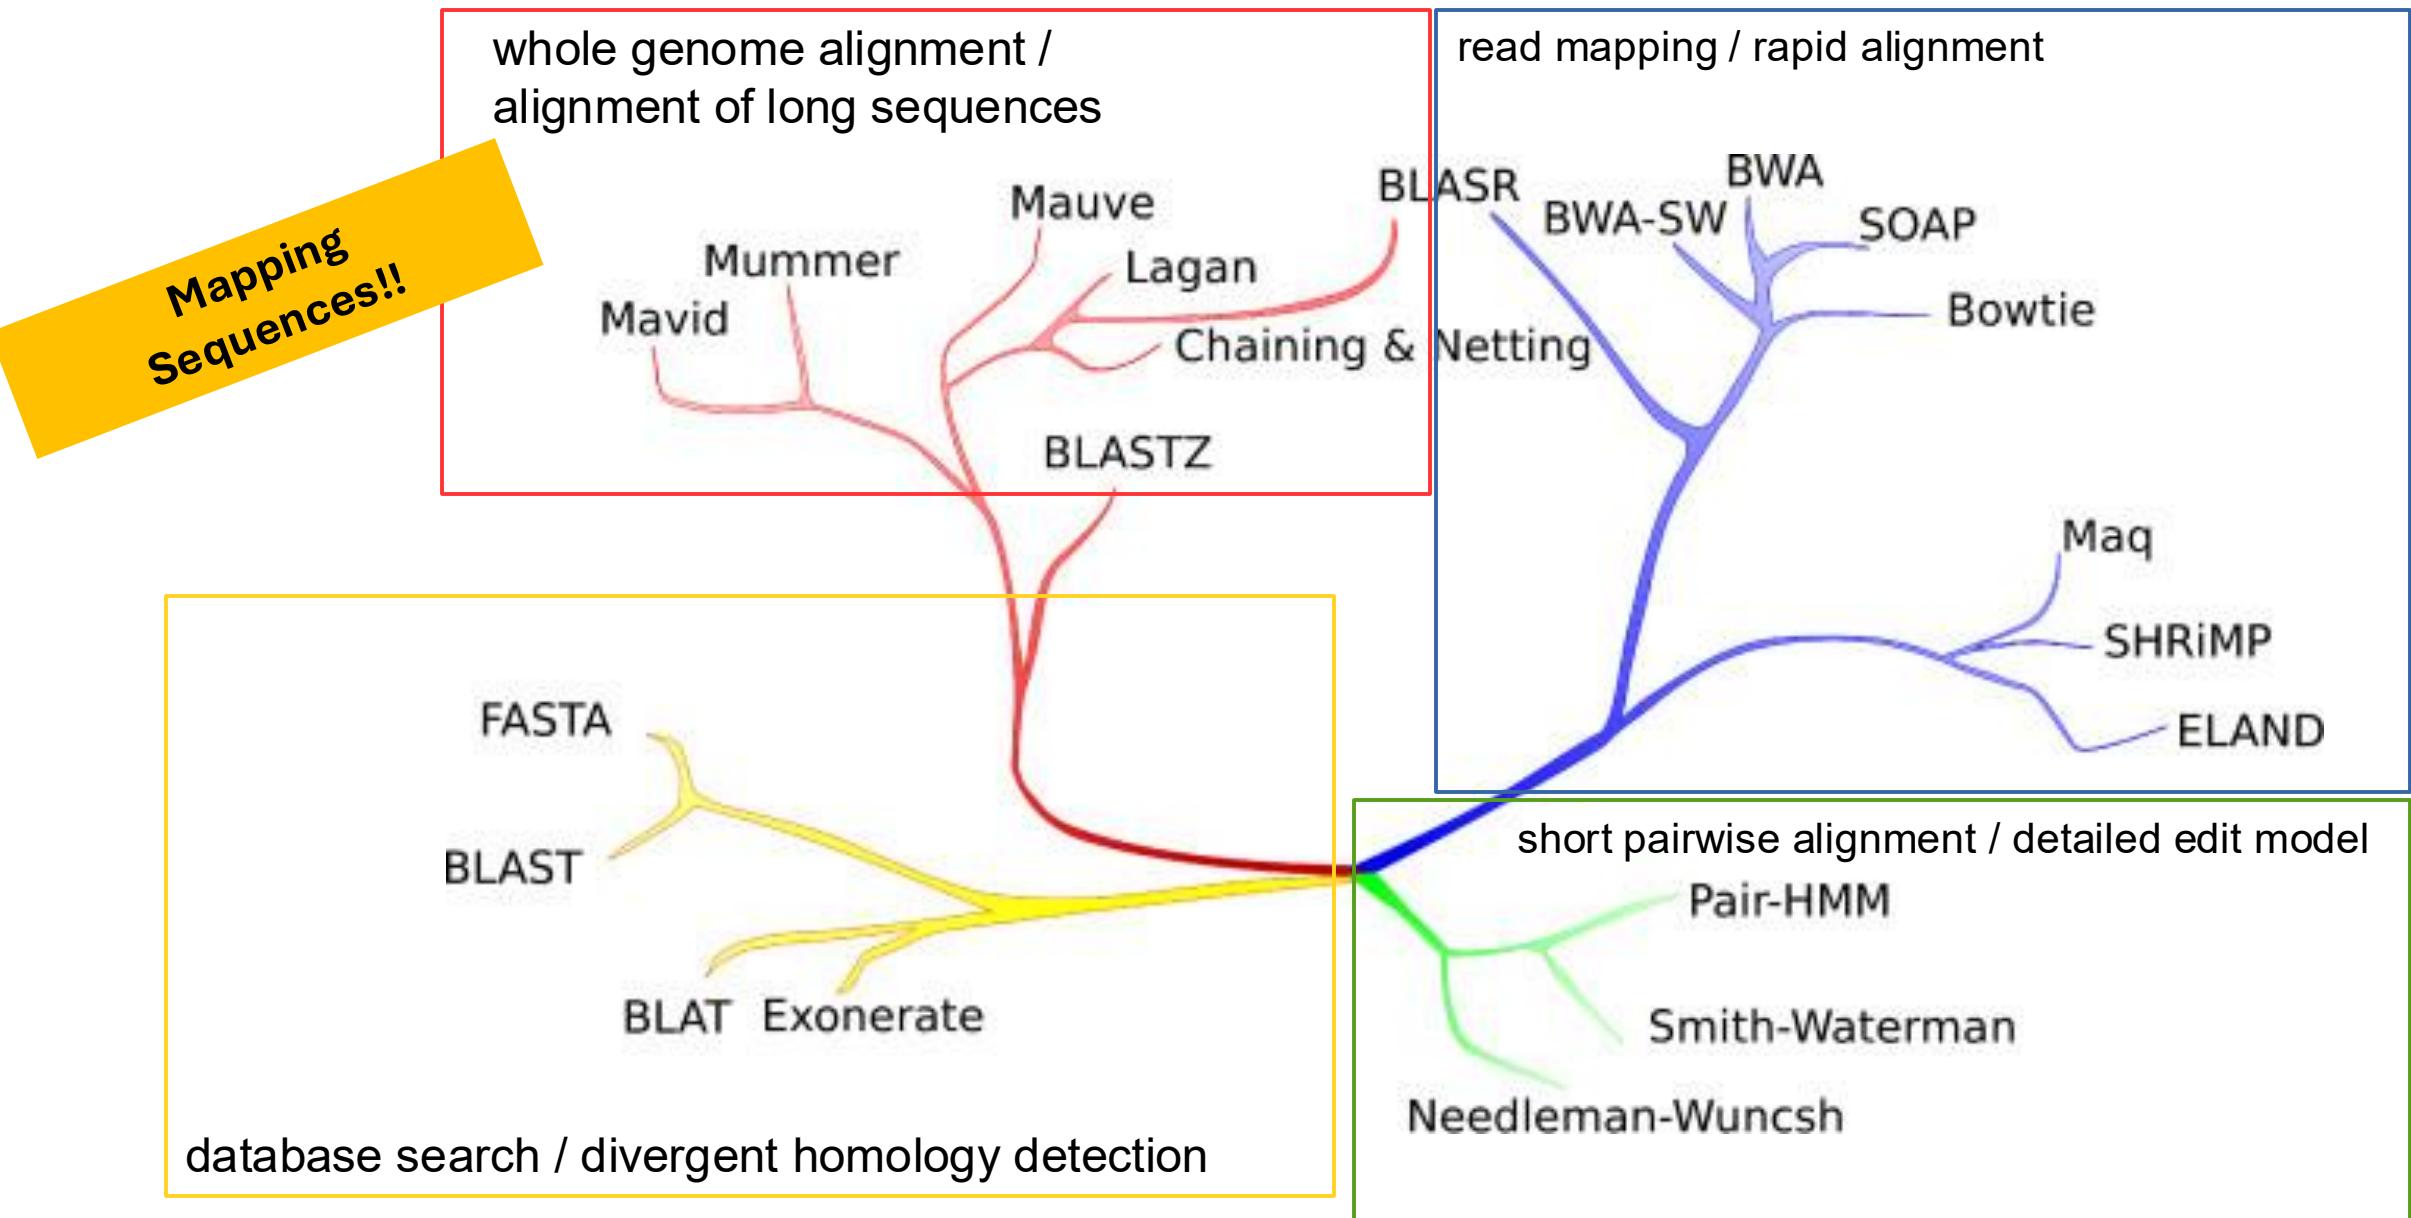
- First, genomic DNA is sheared into fragments that can be size-selected to obtain a template library of uniformly short DNA fragments.
- During sequencing, every molecule is cyclically extended by a single base, and all clusters are read in parallel before the cycle is repeated (Illumina).
- The first step in the analysis of NGS data is to determine where in the genome each of the short reads originated.
- Due to sequencing errors and genuine differences between the reference genome and the sequenced organism, the match is not 100%.



Short read mapping: the problem

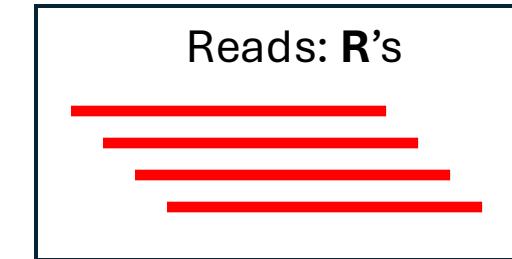
- Big data: sometimes billions of short reads need to be aligned to a large genome. (examples of large genomes?).
- We required new algorithm designs to run orders of magnitude faster.
- Sometimes, we need to incorporate additional information (eg. paired reads)





Read mapping as String Matching

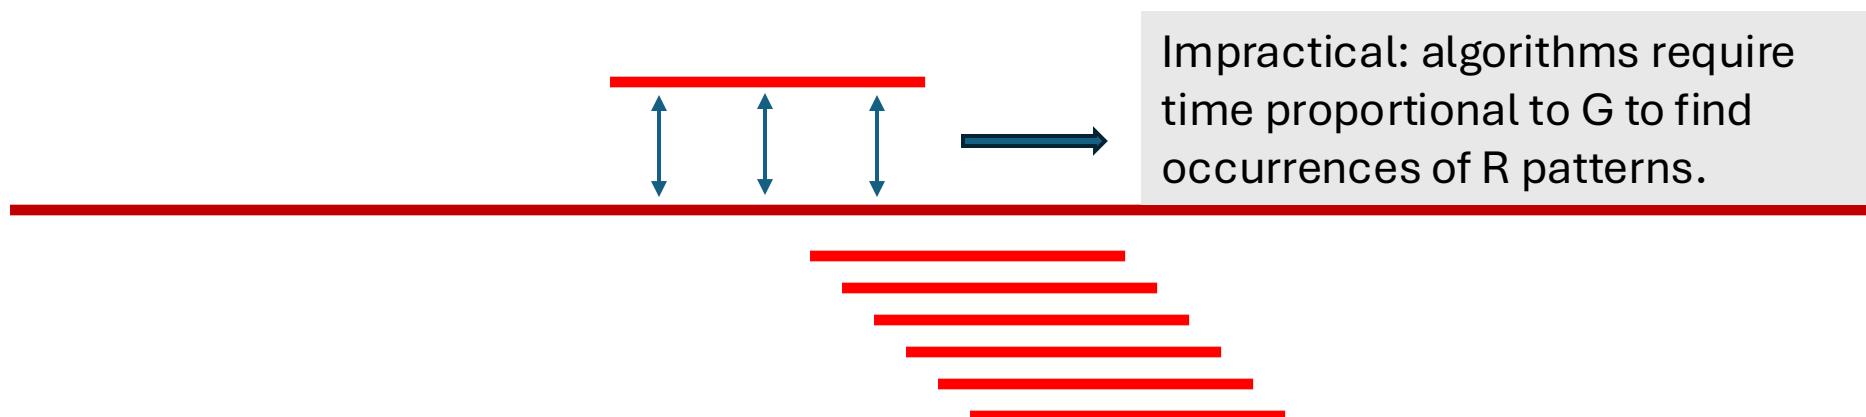
Reference genome: **G**
Substring S: contiguous subsequence



Alphabet: {A, C, G, T}

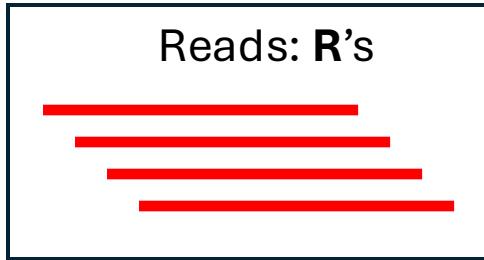
Exact String Matching

- A read R might be detected through sequence identity of R to a substring in G.
- The problem is finding all occurrences of a pattern in R in G.
- But R's are typically hundreds of reads: classical matching algorithms.



Read mapping as String Matching

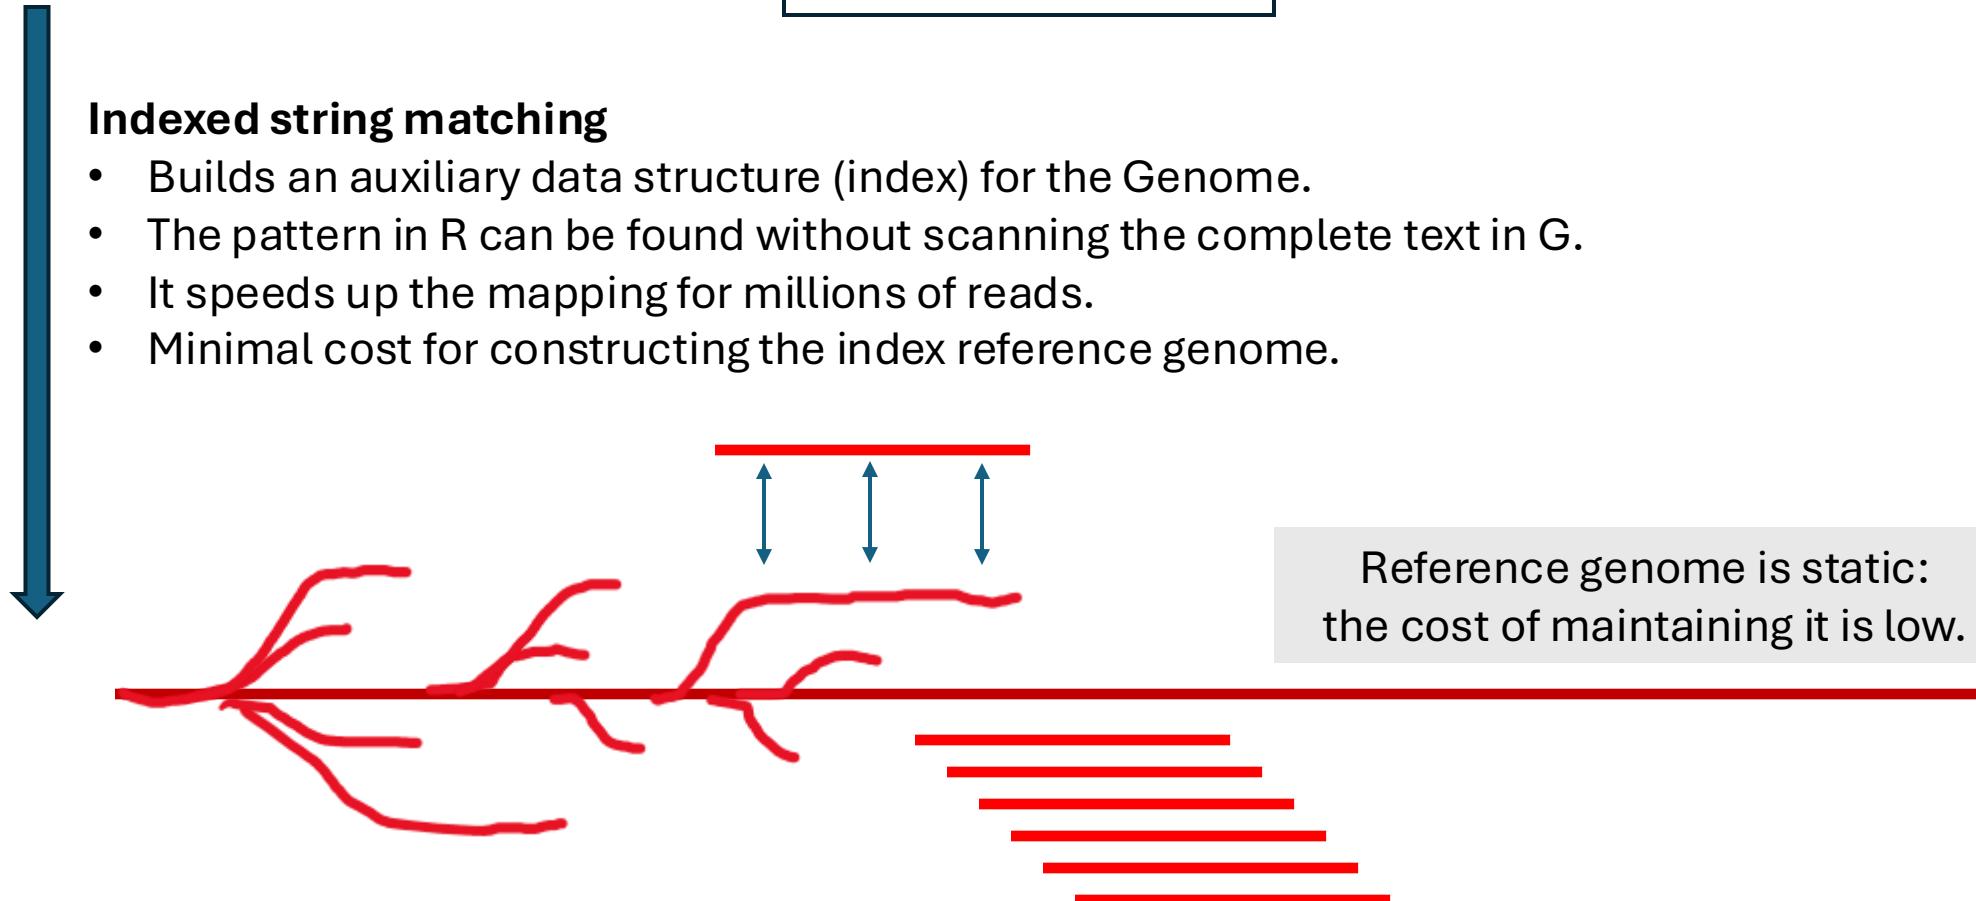
Reference genome: **G**
Substring S: contiguous subsequence



Alphabet: {A, C, G, T}

Indexed string matching

- Builds an auxiliary data structure (index) for the Genome.
- The pattern in R can be found without scanning the complete text in G.
- It speeds up the mapping for millions of reads.
- Minimal cost for constructing the index reference genome.

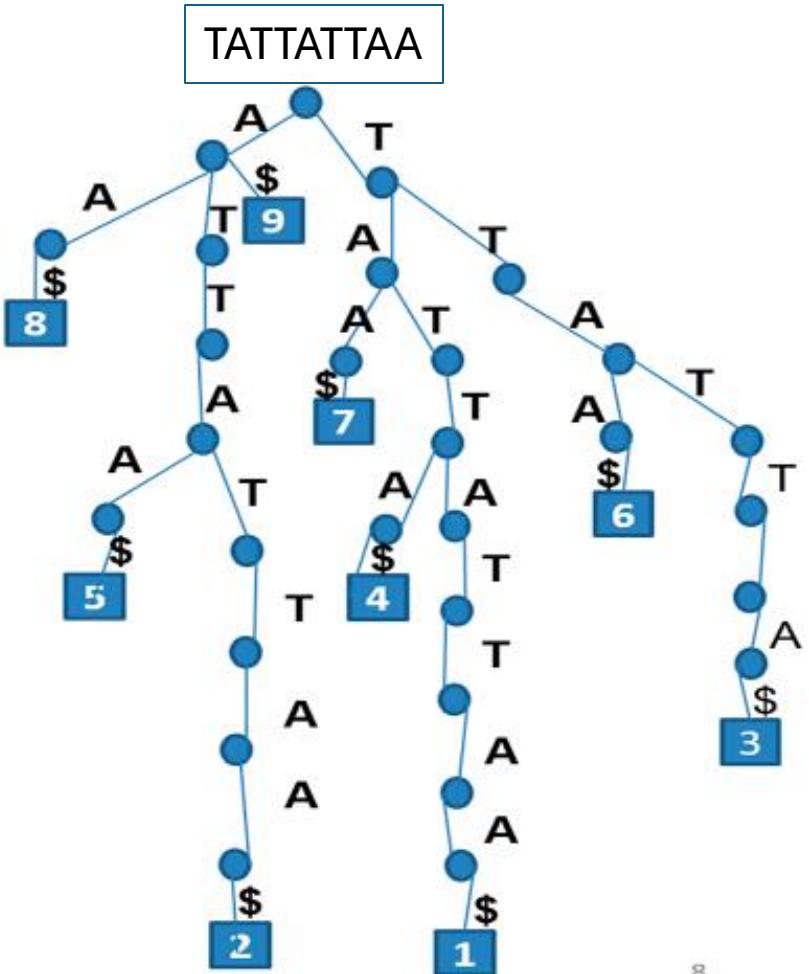


Whole Genome Alignment

- A critical process in comparative genomics, facilitating the detection of genetic variants and aiding our understanding of evolution.
- Classification of Whole Genome Alignment Algorithms:
 - **Suffix Tree-Based** Alignment Methods
 - **MUMmer** Technique
 - Anchor based methods
 - LAGAN
 - Mauve
 - BLASTZ
 - LASTZ
 - Minimap2
 - Hash-based methods
 - Graph-based methods

Whole Genome Alignment

- Suffix Tree-Based Alignment Methods

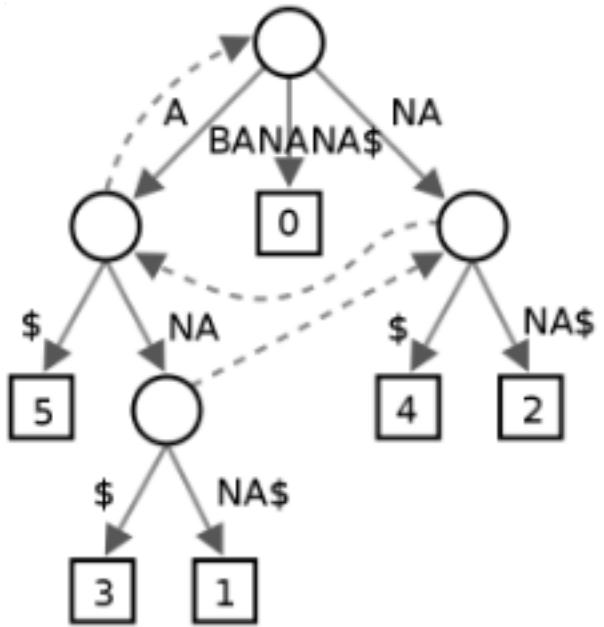


- It is a compressed data tree structure that represents all the suffixes of a given string.
 - The trees save their positions in the text as well as their values.
 - TATTATTAA sequence tree with 9 suffixes and \$ marking the termination of each suffix with a value.
 - Data structured in this way provide fast implementations for string operation.

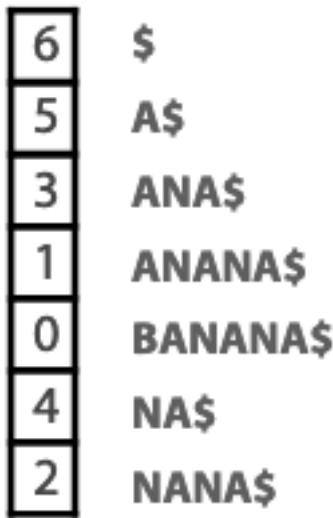
Exact pattern matching with indexing

There are many data structures built on suffixes

String: 'banana'



Suffix Tree



Suffix Array



FM Index

classical *full-text* index structures
Relatively **fast search** functionality
But a large space requirement

Ferragina and Manzini developed their full-text index which takes space proportional to the *compressed* text

Read mapping as String Matching

- **Suffix trees** and **suffix arrays** are classical **full-text index structures**.
- **Suffix arrays** are an array of integers specifying the lexicographical order of the suffixes in the reference genome, and they are more space efficient than suffix trees.
- They permit (almost) optimal $O(n)$ search time, proportional to:
 - number of occurrences of a pattern in the genome.
 - Size of the genome.
 - Length of the read.
 - Number of occurrences
- However, both data structures demand memory resources if the index (Genome) is too large.
 - ie. A suffix tree of the human genome needs ~45 GB of space.
- It wasn't until Ferragina and Manzini (2000) developed the full-text index based on the **Burrows-Wheeler compression algorithm** that the space requirement was reduced while preserving fast functionality.

The Burrows-Wheeler Transform (BWT)

- It is a combination of the Burrows-Wheeler compression algorithm with the the Suffix Array data structure to obtain a sort of compressed suffix array.
- This transformation of a string has been useful in compression: that was its original motivation, and the **b** in bzip2 stands for this.
- BWT also has formed the basis of a number of modern string algorithms that can operate on text using a small amount of space.

Bowtie

Software Highly accessed Open access

Ultrafast and memory-efficient alignment of short DNA sequences to the human genome

Ben Langmead*, Cole Trapnell, Mihai Pop and Steven L Salzberg

* Corresponding author: Ben Langmead langmead@cs.umd.edu

Center for Bioinformatics and Computational Biology, Institute for Advanced Computer Studies, University of Maryland, College Park, MD 20742, USA
For all author emails, please [log on](#).

Genome Biology 2009, **10**:R25 doi:10.1186/gb-2009-10-3-r25

The electronic version of this article is the complete one and can be found online at:
<http://genomebiology.com/2009/10/3/R25>

bwa

Fast and accurate short read alignment with Burrows–Wheeler transform

Heng Li and Richard Durbin*

+ Author Affiliations

* To whom correspondence should be addressed.

Received February 20, 2009.
Revision received May 6, 2009.
Accepted May 12, 2009.

Bioinformatics (2009) **25**(14):1754–1760.

Bowtie alignment performance versus SOAP and Maq

	Platform	CPU time	Wall clock time	Reads mapped per hour (millions)	Peak virtual memory footprint (megabytes)	Bowtie speed-up	Reads aligned (%)
Bowtie -v 2	Server	15 m 7 s	15 m 41 s	33.8	1,149	-	67.4
SOAP		91 h 57 m 35 s	91 h 47 m 46 s	0.10	13,619	351x	67.3
Bowtie	PC	16 m 41 s	17 m 57 s	29.5	1,353	-	71.9
Maq		17 h 46 m 35 s	17 h 53 m 7 s	0.49	804	59.8x	74.7
Bowtie	Server	17 m 58 s	18 m 26 s	28.8	1,353	-	71.9
Maq		32 h 56 m 53 s	32 h 58 m 39 s	0.27	804	107x	74.7

Performance and sensitivity of Bowtie v0.9.6, SOAP v1.10, and Maq v0.6.6 when aligning **8.84 M reads (35 bp)**.

Source: <https://doi.org/10.1186/gb-2009-10-3-r25>

The Burrows-Wheeler Transform (BWT)

- Most modern aligners use a full-text index in Minute space, or **FM-Index** as the genome index structure.
- They perform well in both overall runtime and memory usage.
- The FM index's most important component is its use of the Burrows-Wheeler transform (BWT) of the reference genome.
- The BWT is created by first building an **array of suffix rotations in the genome**.
 - starting with the full genome as the first suffix
 - then the next suffix removes the first letter and appends it onto the end.
 - ie. the suffix rotations of the word “banana” and “\$” match the end of the string:
 - banana\$
 - anana\$b
 - nana\$ba

The Burrows-Wheeler Transform (BWT)

Text transform that is useful for compression & search.

banana

banana\$

anana\$b

nana\$ba

ana\$ban

na\$bana

a\$banan

\$banana

sort →

\$banana

a\$banan

ana\$ban

anana\$b

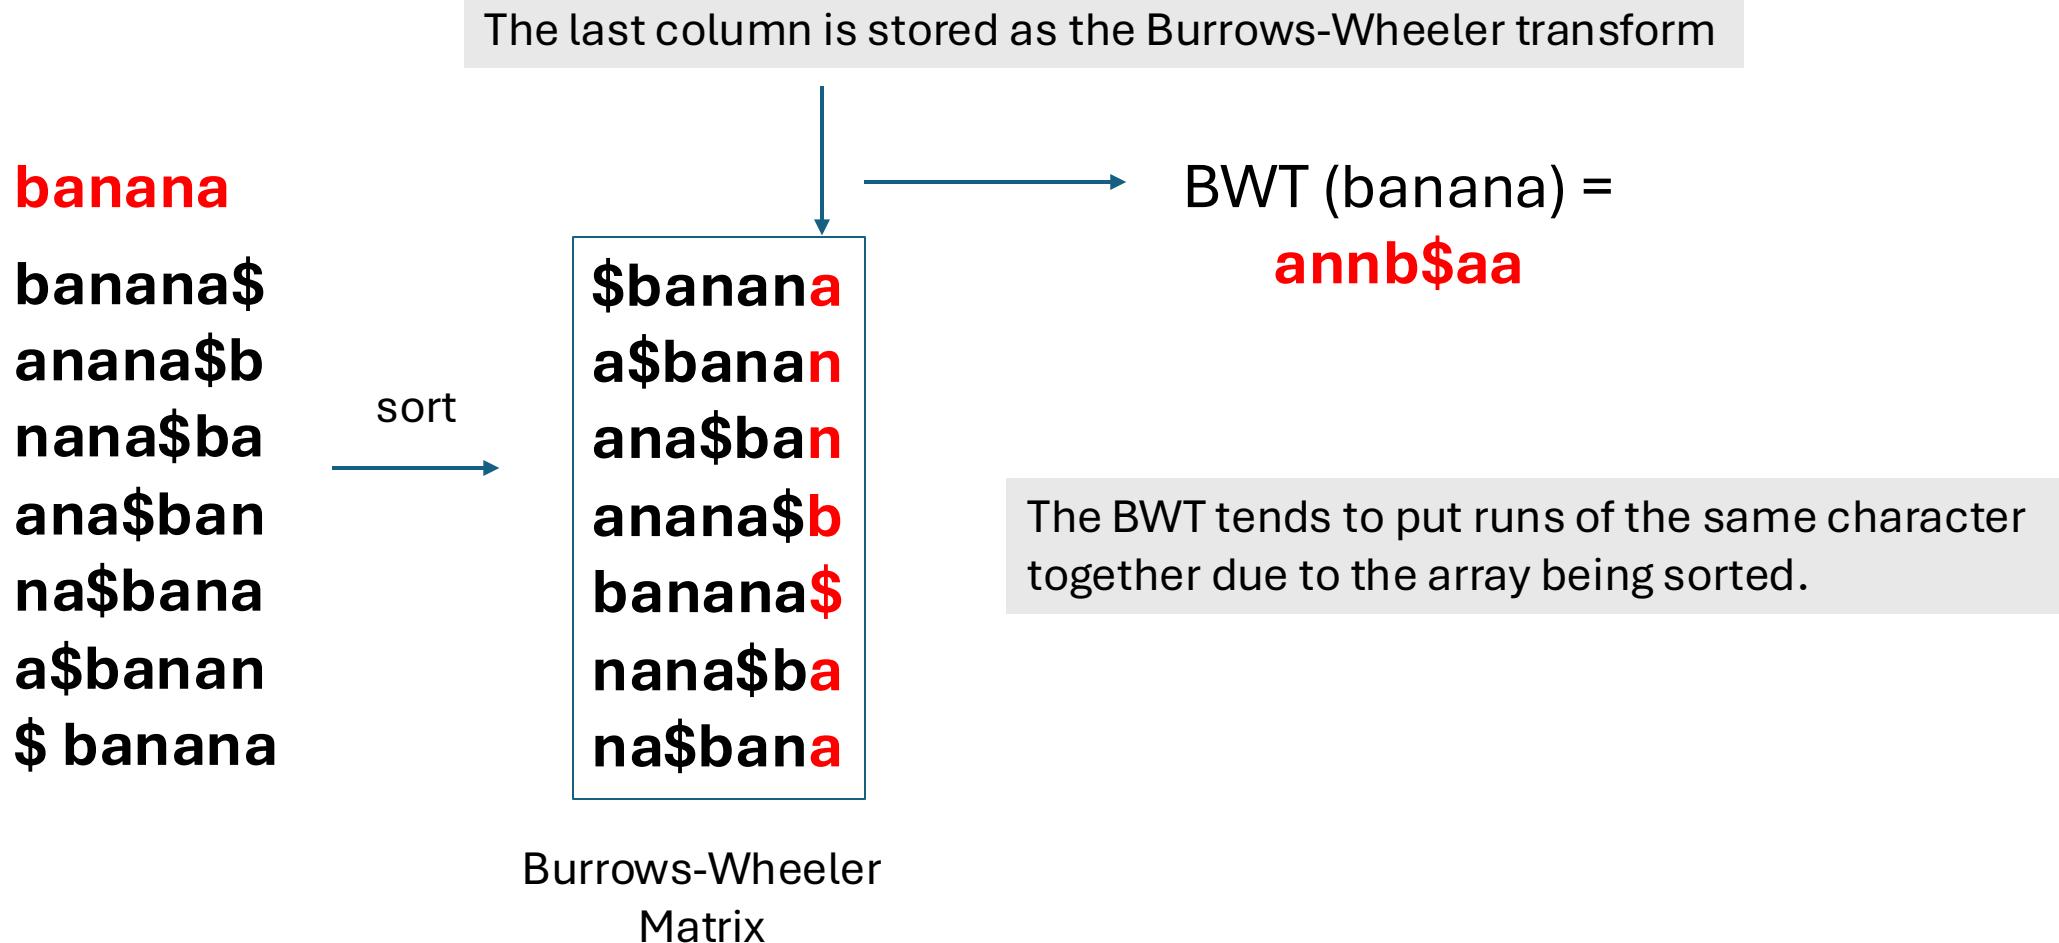
banana\$

nana\$ba

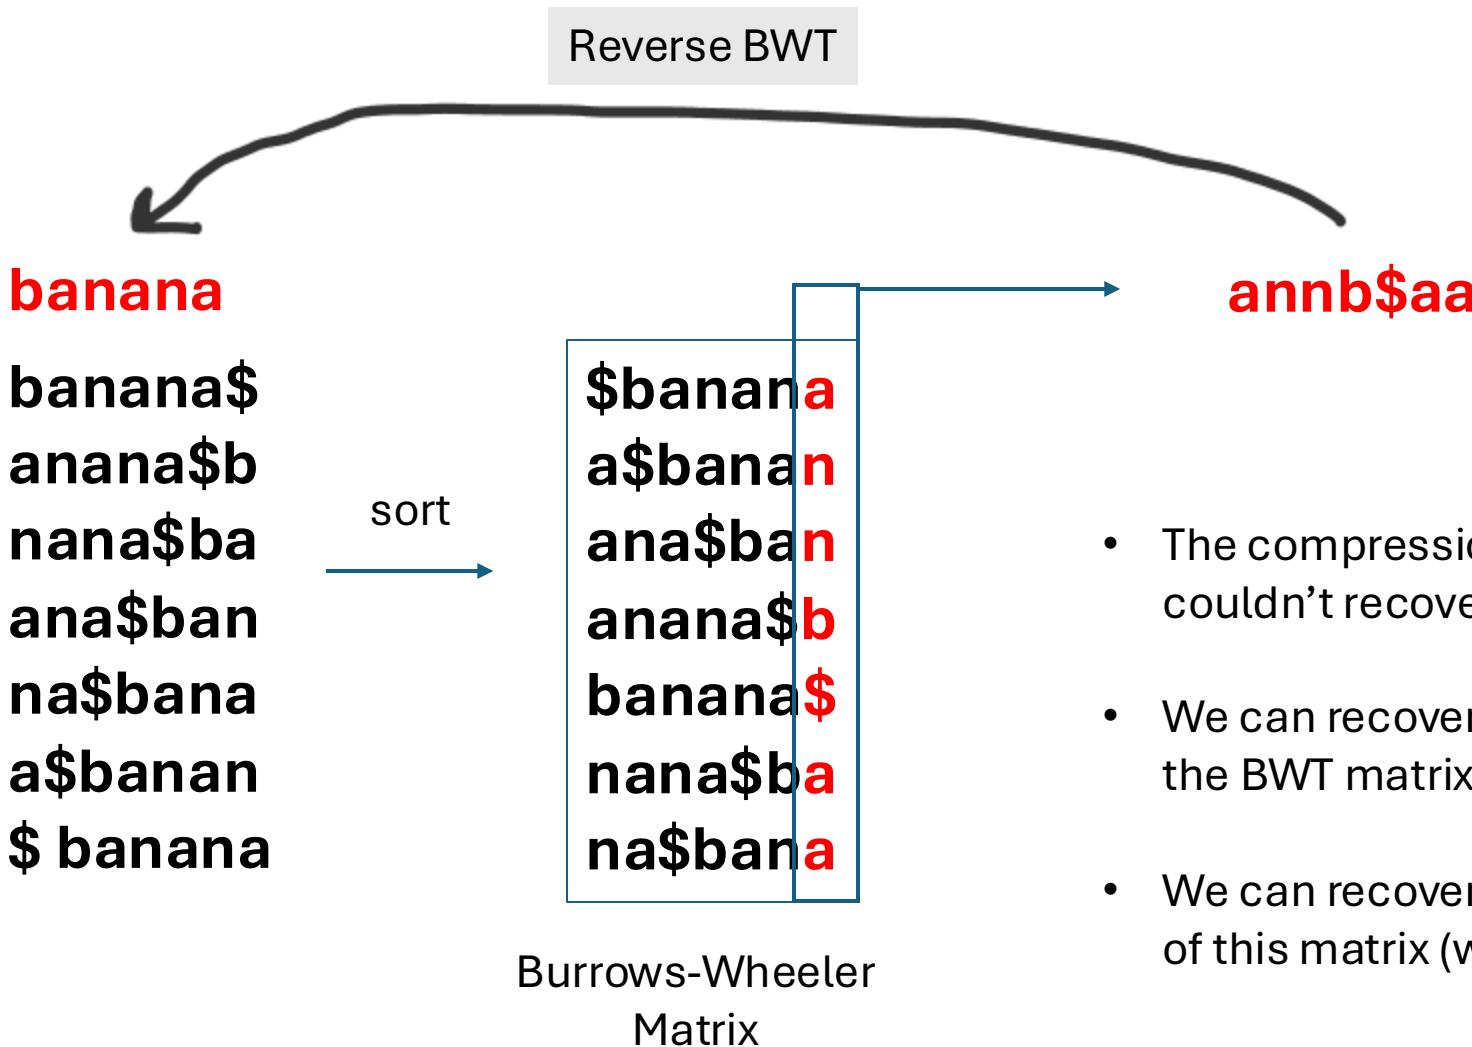
na\$bana

- “\$” match the end of the string
- generate **all rotations** of the reference sequence.
- **sort** them lexicographically.

The Burrows-Wheeler Transform (BWT)



The Burrows-Wheeler Transform (BWT)



- The compression wouldn't be useful if we couldn't recover the original string.
- We can recover the string by recovering the BWT matrix from just the BWT string.
- We can recover the string as the first row of this matrix (which will start with "\$").

The Burrows-Wheeler Transform (BWT)

Now, we pair up F and L (ranked) to establish a mapping from one column to the other. This will be used for LF-mapping.

banana
banana\$
anana\$b
nana\$ba
ana\$ban
na\$bana
a\$banan
\$ banana

sort →

\$banana
a\$banana
ana\$ban
anana\$b
banana\$
nana\$ba
na\$banan

BWT (banana) =
annb\$aa

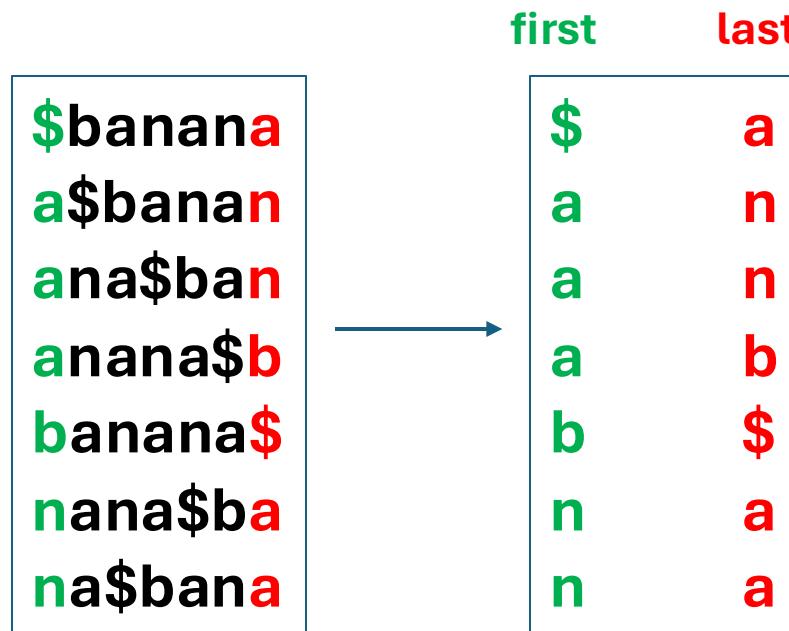
LF Property: The order of occurrence of a letter X in the **last column** corresponds to the same order of occurrence of X in the **first column**.

Burrows-Wheeler
Matrix

The Burrows-Wheeler Transform (BWT)

BWT ("banana\$")
annb\$aa

T-ranking
Give each character a rank, equal to # times the character occurred Previously



Burrows-Wheeler
Matrix

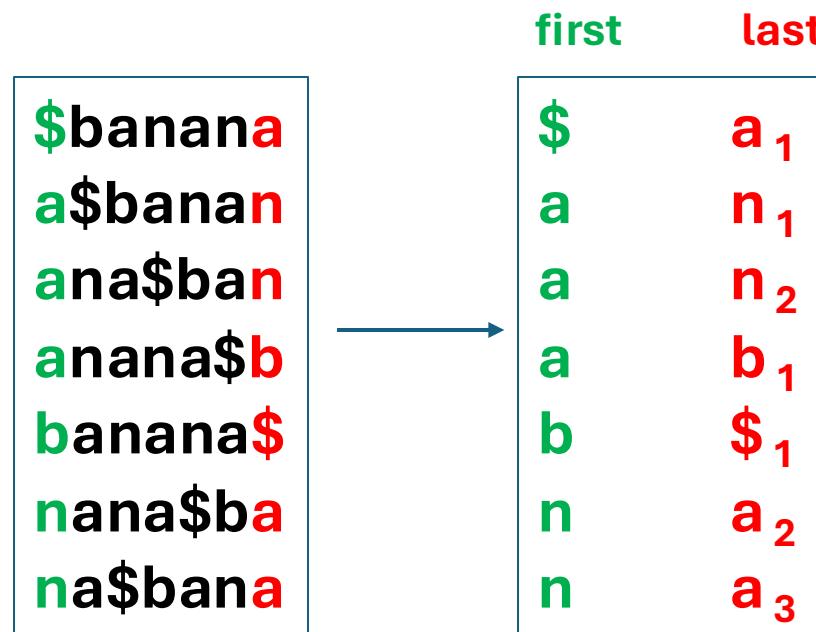
The Burrows-Wheeler Transform (BWT)

BWT ("banana\$")

$a_1 n_1 n_2 b_1 \$_1 a_2 a_3$

T-ranking

Give each character a rank, equal to # times the character occurred previously



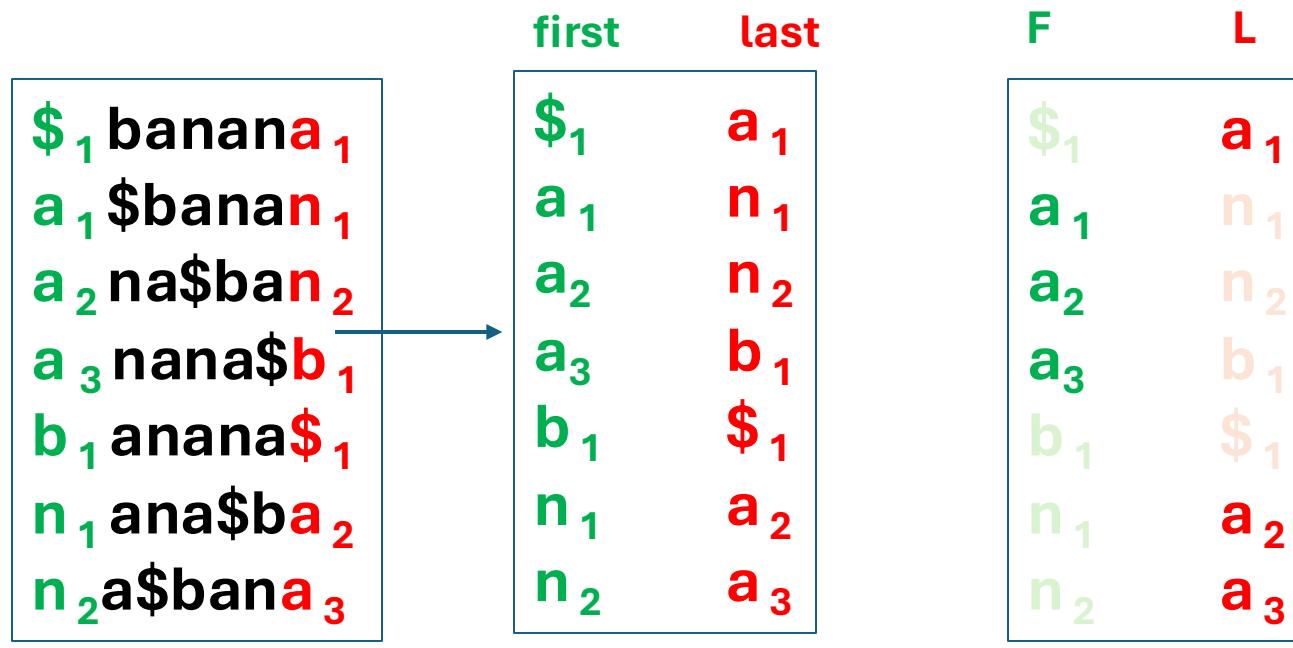
Burrows-Wheeler
Matrix

The Burrows-Wheeler Transform (BWT)

BWT ("banana\$")

$a_1 n_1 n_2 b_1 \$_1 a_2 a_3$

LF Property: The order of occurrence of a letter X in the **last column** corresponds to the same order of occurrence of X in the **first column**.



Burrows-Wheeler
Matrix

a's occur in the same order in F and L.

n's: same order in F and L.

The Burrows-Wheeler Transform (BWT)

Recovering (inverting) original string using BWT

BWT (“banana\$”)

$a_1 n_1 n_2 b_1 \$_1 a_2 a_3$



Reconstruct the first column
sort BWT

first	last
$\$_1$	a_1
a	n_1
a	n_2
a	b_1
b	$\$_1$
n	a_2
n	a_3

The Burrows-Wheeler Transform (BWT)

Recovering (inverting) original string using BWT

BWT (“banana\$”)

$a_1 n_1 n_2 b_1 \$_1 a_2 a_3$



Reconstruct the first column

sort BWT



LF Property: The order of occurrence of a letter X in the **last column** corresponds to the same order of occurrence of X in the **first column**.

first	last
$\$_1$	a_1
a_1	n_1
a_2	n_2
a_3	b_1
b_1	$\$_1$
n_1	a_2
n_2	a_3

Reconstruct the original string by the following letters in the rotations.

The Burrows-Wheeler Transform (BWT)

Recovering (inverting) original string using BWT

BWT (“banana\$”)

a₁n₁n₂b₁\$₁a₂a₃

first	last
\$ ₁ b ₁ a ₃ n ₂	a ₁
a ₁	n ₁
a ₂	n ₂
a ₃	b ₁
b ₁	\$ ₁
n ₁	a ₂
n ₂	a ₃

The Burrows-Wheeler Transform (BWT)

Recovering (inverting) original string using BWT

BWT (“banana\$”)

a₁n₁n₂b₁\$₁a₂a₃



original string

b₁a₃n₂a₂n₁a₁\$₁

first	last
\$ ₁	b ₁
b ₁	a ₃
a ₃	n ₂
n ₂	a ₂
a ₂	n ₁
n ₁	a ₁
a ₁	\$ ₁

First sorted rotation
With T-ranking numbers

The Burrows-Wheeler Transform (BWT)

Pattern matching using BWT

BWT (“banana\$”)

a₁n₁n₂b₁\$₁a₂a₃



Reconstruct the first column

sort BWT



-Keep track of the first and last column

-Provide row numbers

- **First** column
- **Last** to **First** column row numbers

i	first	last	L2F(i)
0	\$ ₁ banan a ₁	a ₁	1
1	a ₁ \$bana n ₁	n ₁	5
2	a ₂ na\$bba n ₂	n ₂	6
3	a ₃ nana\$ b ₁	b ₁	4
4	b ₁ anana \$ ₁	\$ ₁	0
5	n ₁ ana\$b a ₂	a ₂	2
6	n ₂ a\$ban a ₃	a ₃	3

Given a query, I can trace back to the original string

The Burrows-Wheeler Transform (BWT)

Pattern matching using BWT

BWT (“banana\$”)

$a_1 n_1 n_2 b_1 \$_1 a_2 a_3$



Reconstruct the first column



- Keep track of the first and last column
- Provide row numbers
 - First column
 - Last to First column row numbers

i	first	last	<u>L2F(i)</u>
0	$\$_1$	a_1	1
1	a_1	n_1	5
2	a_2	n_2	6
3	a_3	b_1	4
4	b_1	$\$_1$	0
5	n_1	a_2	2
6	n_2	a_3	3

LF mapping: The i^{th} occurrence of a character c in **L** and the i^{th} occurrence of c in **F** correspond to the same occurrence in the original string.

The Burrows-Wheeler Transform (BWT)

Pattern matching using BWT

BWT (“banana\$”)

a₁n₁n₂b₁\$₁a₂a₃

i	first	last	L2F (i)
0	\$ ₁	a ₁	1
1	a ₁	n ₁	5
2	a ₂	n ₂	6
3	a ₃	b ₁	4
4	b ₁	\$ ₁	0
5	n ₁	a ₂	2
6	n ₂	a ₃	3

- Given a query, I can trace back to the original string.
- In the original BWT matrix, all rotations that started with the same prefix were plumbed together.
- If I start with the top first rotation of the start of my query, and bottom the the last rotation of the start of my query, the range between top and bottom is exactly all instances of my query in the original string.

The Burrows-Wheeler Transform (BWT)

Pattern matching using BWT

BWT (“banana\$”)

a₁n₁n₂b₁\$₁a₂a₃

A diagram illustrating the Burrows-Wheeler Transform (BWT) and its inverse mapping. On the left, the original string "banana\$" is shown in red: **a₁n₁n₂b₁\$₁a₂a₃**. To its right is a BWT matrix with 7 rows labeled i from 0 to 6. The matrix columns are labeled **first** and **last**. The **first** column contains the suffixes of the string: \$₁, a₁, a₂, a₃, b₁, n₁, n₂. The **last** column contains the characters of the string: a₁, n₁, n₂, b₁, \$₁, a₂, a₃. An arrow labeled "top" points to row 0, and another arrow labeled "bottom" points to row 6. To the right of the matrix is a vertical column labeled **L2F(i)** with values 1, 5, 6, 4, 0, 2, 3.

i	first	last	L2F(i)
0	\$ ₁	a ₁	1
1	a ₁	n ₁	5
2	a ₂	n ₂	6
3	a ₃	b ₁	4
4	b ₁	\$ ₁	0
5	n ₁	a ₂	2
6	n ₂	a ₃	3

- Given a query, I can trace back to the original string.
- String: ‘an’ (query)
- Search for my word backward.
- First ‘n’: search **rotations** in the **last** column for ‘n’

The Burrows-Wheeler Transform (BWT)

Pattern matching using BWT

BWT (“banana\$”)

a₁n₁n₂b₁\$₁a₂a₃

i	first	last	L2F(i)
0	\$ ₁	a ₁	1
1	a ₁	n ₁	5
2	a ₂	n ₂	6
3	a ₃	b ₁	4
4	b ₁	\$ ₁	0
5	n ₁	a ₂	2
6	n ₂	a ₃	3

top → bottom →

- Given a query, I can trace back to the original string.
- String: ‘an’ (query)
- Search for my word backward.
- First ‘n’
- Look for next ‘a’ between the range (top-bottom).

The Burrows-Wheeler Transform (BWT)

Pattern matching using BWT

BWT (“banana\$”)

a₁n₁n₂b₁\$₁a₂a₃

i	first	last	L2F(i)
0	\$ ₁	a ₁	1
1	a ₁	n ₁	5
2	a ₂	n ₂	6
3	a ₃	b ₁	4
4	b ₁	\$ ₁	0
5	n ₁	a ₂	2
6	n ₂	a ₃	3

top → bottom →

original string
b₁a₃n₂a₂n₁ a₁\$₁

- Given a query, I can trace back to the original string.
- String: ‘an’ (query)
- Search for my word backward.
- First ‘n’
- Look for next ‘a’ between the range (top-bottom).
- a₂ and a₃ are the two instances of my query.
- Map back those rotations to the original string.

The Burrows-Wheeler Transform (BWT)

Pattern matching using BWT

BWT (“banana\$”)

a₁n₁n₂b₁\$₁a₂a₃

i	first	last	L2F(i)
0	\$ ₁	a ₁	1
1	a ₁	n ₁	5
2	a ₂	n ₂	6
3	a ₃	b ₁	4
4	b ₁	\$ ₁	0
5	n ₁	a ₂	2
6	n ₂	a ₃	3

- Given a query, I can trace back to the original string.
- String (query):
 - ‘na’
 - ‘nan’
 - ‘ba’
 - ‘bana’

original string

b₁a₃n₂a₂n₁ a₁\$₁

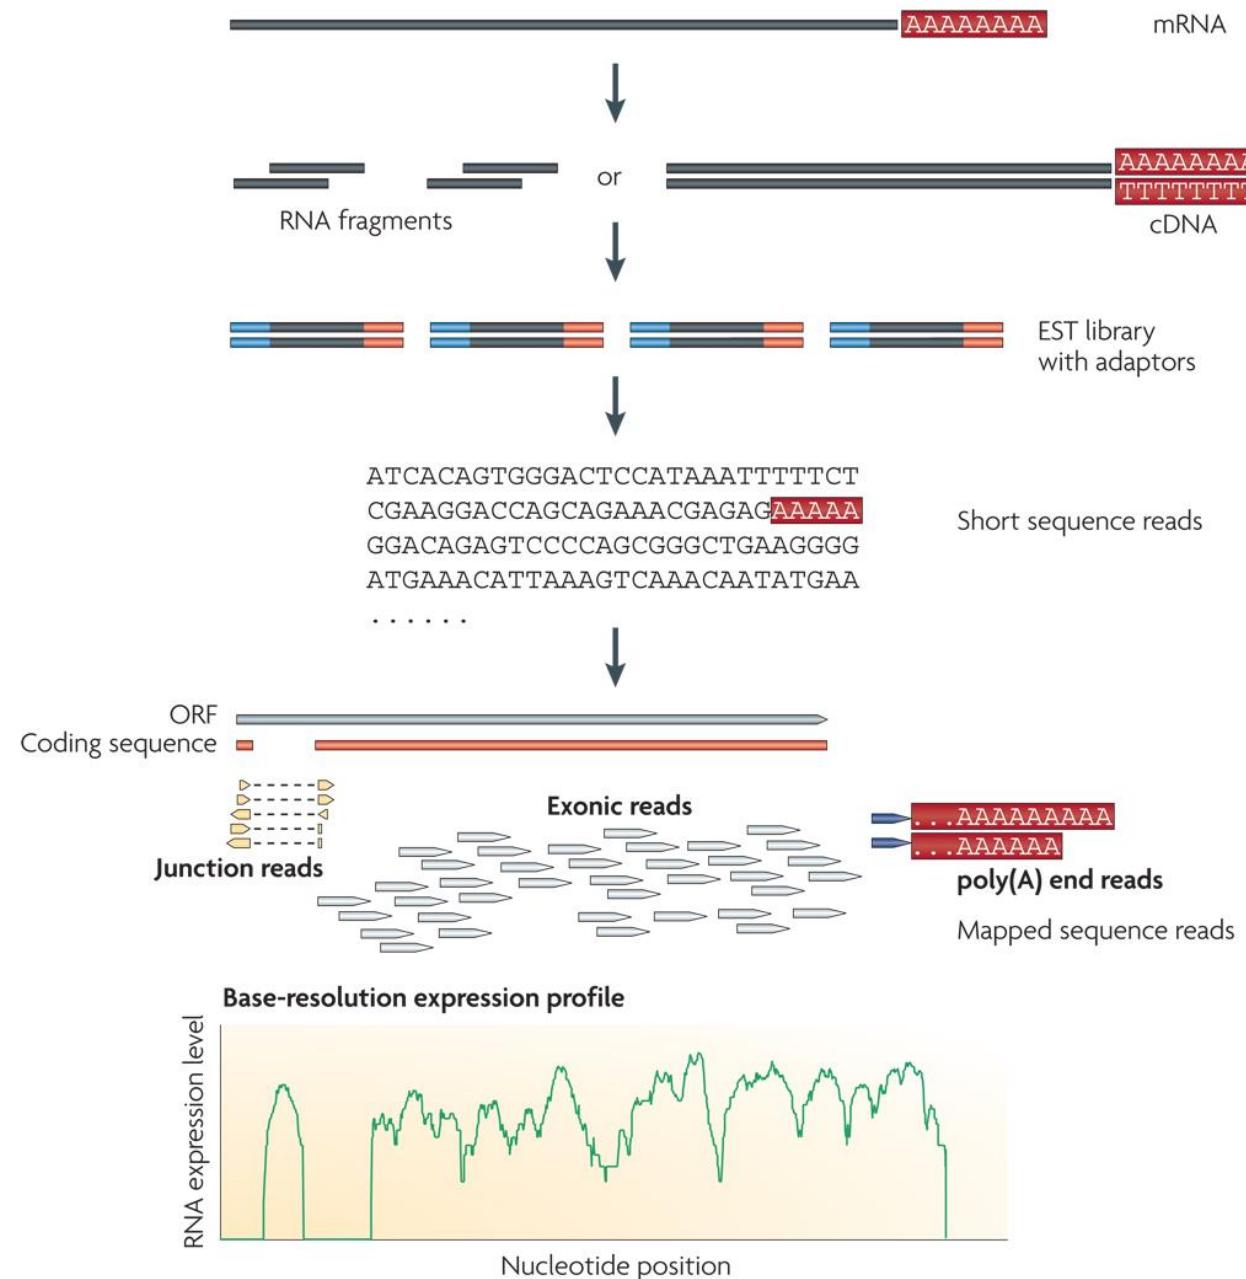
The Burrows-Wheeler Transform (BWT)

- BWT is useful for pattern searching and compression.
- BTW is invertible: given the BTW of a string, the original string can be reconstructed.
- BWT is computable in $O(n)$ time: $O(|pattern|)$
- There is a close relationship between Suffix Trees, Suffix Arrays, and BWT:
 - Suffix Array: order of the suffix numbers of the suffix tree, traversed left to right.
 - BWT: letters at positions given by the suffix array entries.
- After compression, strings can be searched quickly with BWT:
 - Don't have to store the LF mapping
 - Allows computing searches (another algorithm) in compressed data *on the fly* (with some extra storage)

RNA-seq

- RNA-seq is the sequencing of coding and non-coding transcripts.
 - Eukaryotes: polyA tail in mRNAs
 - But also: miRNA, piRNAs, siRNAs, lncRNAs,... is RNA-seq
- Requires reverse transcribe RNA -> DNA -> cDNA library.
- Applications:
 - Gene prediction.
 - *denovo* assembly of transcripts.
 - RNA-seq quantification: gene expression / differential gene expression.
 - Inference of alternating splicing.

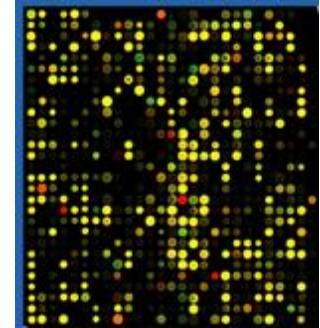
RNA-seq



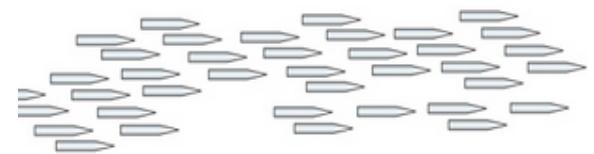
Source: <https://doi.org/10.1038/nrg2484>

RNA-seq vs. Microarrays

- **Microarrays:** measure transcription using probes for specific genes:
 - One probe: one gene
 - Relies on hybridization
 - Requires knowledge of gene sequences
- But microarrays have
 - low reproducibility
 - Limits of detection at extreme high/low expression



- **RNA-seq** has improved reproducibility.
- There are no limitations at high/low expression levels.
 - Low expression is solved by increasing the yield (sequencing depth).
- Now, it is more affordable to perform RNA-seq.



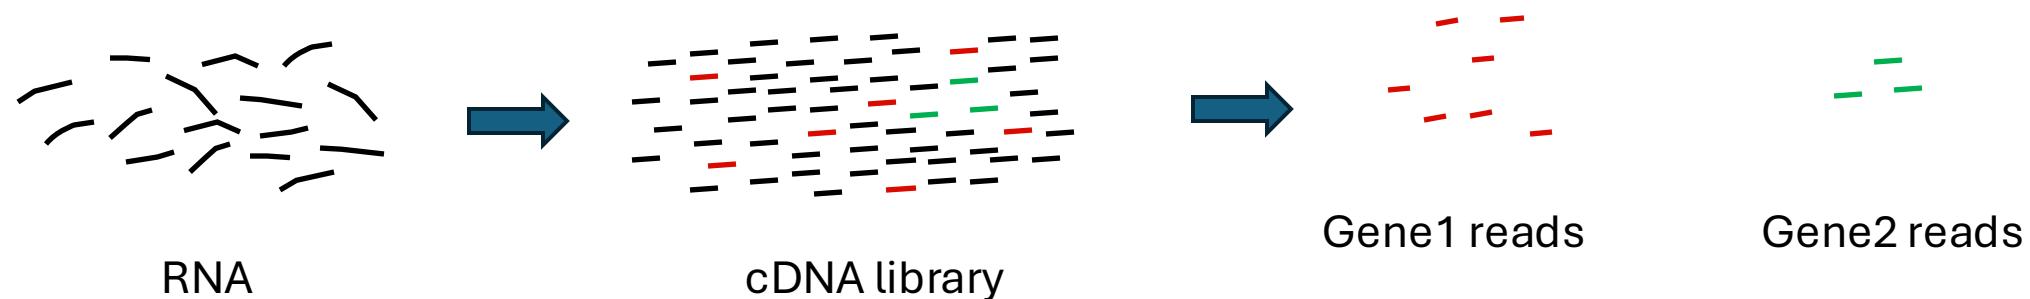
RNA-seq – Gene Prediction

- Gene finding in eukaryotic genomes is notoriously difficult to automate, even with HMMs.
- Most gene prediction tools incorporate RNA-Seq data into unsupervised/supervised training and subsequently generate *ab initio* gene predictions.



RNA-seq quantification

- RNA-seq is a relative abundance measurement technology.
- RNA-seq gives you reads from the ends of a random sample of fragments in your library.
- Without additional data, this only gives information about the **relative** abundances.
- Additional information, such as levels of “spike-in” transcripts, are needed for absolute measurements.



RNA-seq quantification

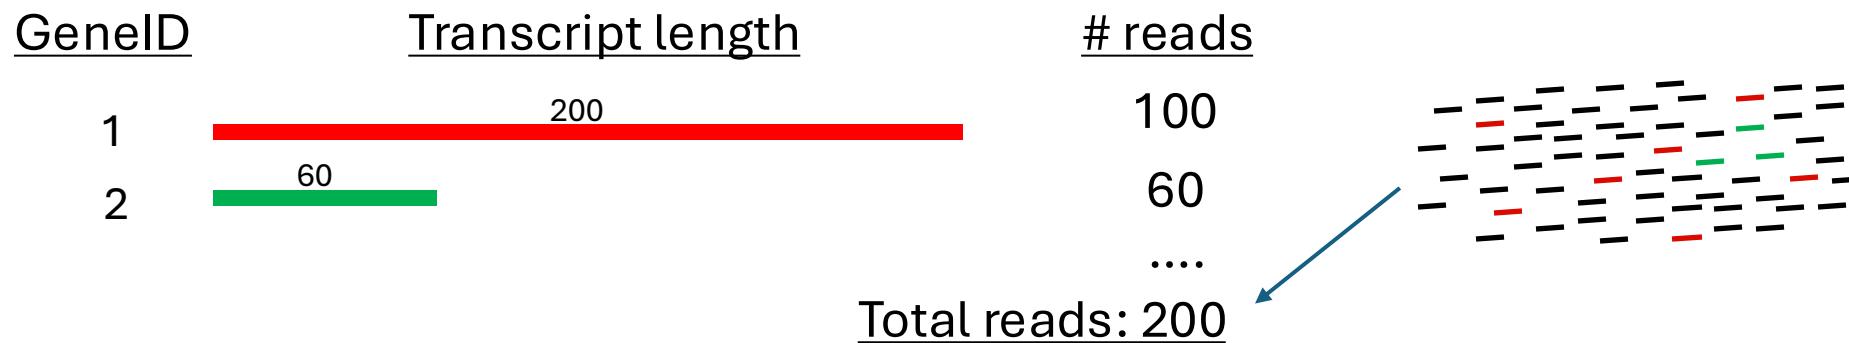
Normalization is required to compare samples in these situations.

Gene	Sample 1 absolute abundance	Sample 1 relative abundance	Sample 2 absolute abundance	Sample 2 relative abundance
1	20	10%	20	5%
2	20	10%	20	5%
3	20	10%	20	5%
4	20	10%	20	5%
5	20	10%	20	5%
6	100	50%	300	75%

RNA-seq quantification

- In RNA-seq read abundance is also dependent of transcript length.
- The probability of a read coming from a transcript is proportional to the relative abundance x length:

$$\text{Relative abundance} = \frac{\text{Probability of read from transcript}}{\text{Transcript length}} = \frac{\frac{\text{reads-trancript}}{\text{total-reads}}}{\text{Transcript length}}$$



$$\text{Relative abundance}_1 = \frac{\frac{100}{200}}{200} = \frac{1}{400}$$



$$\text{Relative abundance}_2 = \frac{\frac{60}{60}}{60} = \frac{2}{400}$$

RNA-seq quantification

-Basic quantification algorithm:

- Align reads (Burrows-Wheeler Transform) against
 - Genome/gene annotations
 - Set of reference transcripts
- Count the number of reads per feature (gene, CDS, exon, etc..)
- Convert read counts into relative expression levels.

-Normalization methods:

- RPKM - Reads Per Kilobase per Million mapped reads
- FPKM (fragments instead of reads, two reads per fragment for paired-end reads)
- TPM - Transcripts Per Million: normalize for gene length first, and then normalize for sequencing dept.

RNA-seq quantification

- RPKM - Reads Per Kilobase per Million mapped reads

$$RPKM = \frac{ExonMappedReads * 10^9}{TotalMappedReads * ExonLength}$$

- TPM - Transcripts Per Million

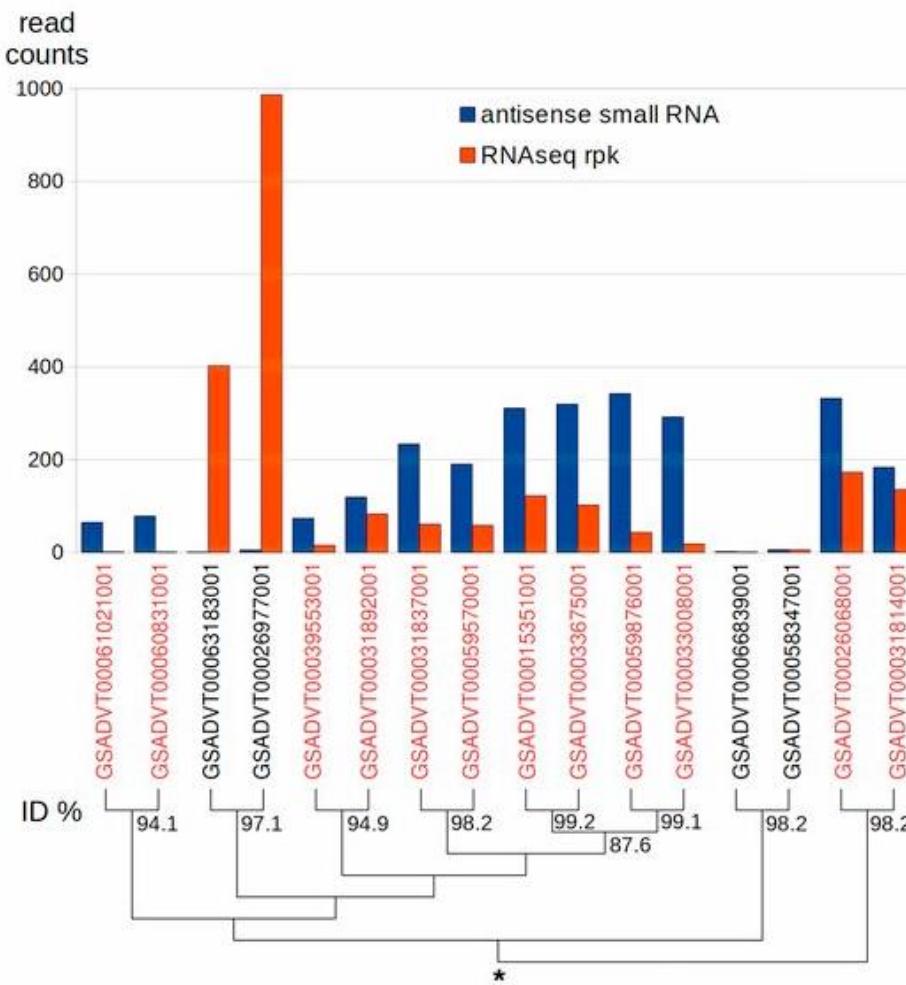
$$TPM = \frac{N_i/L_i * 10^6}{sum(N1/L1 + N2/L2 + \dots + Nn/Ln)}$$

N_i is the number of reads compared to the i -exon

L_i is the length of the i -exon

$Sum(N1/L1 + N2/L2 + \dots + Nn/Ln)$ is the sum of values of all (n) exons after normalization by length.

RNA-seq quantification



RNA profiles from the cytochrome P450 node marked with antisense small RNA counts and RNA-seq rpk counts.

Source: <https://doi.org/10.1534/genetics.116.186734>